



SAS[®] Viya[®] Platform: Deployment Guide

2023.06*

* This document might apply to additional versions of the software. Open this document in [SAS Help Center](#) and click on the version in the banner to see all available versions.

The correct bibliographic citation for this manual is as follows: SAS Institute Inc. 2023. *SAS® Viya® Platform: Deployment Guide*. Cary, NC: SAS Institute Inc.

SAS® Viya® Platform: Deployment Guide

Copyright © 2023, SAS Institute Inc., Cary, NC, USA

All Rights Reserved. Produced in the United States of America.

For a hard copy book: No part of this publication may be reproduced, stored in a retrieval system, or transmitted, in any form or by any means, electronic, mechanical, photocopying, or otherwise, without the prior written permission of the publisher, SAS Institute Inc.

For a web download or e-book: Your use of this publication shall be governed by the terms established by the vendor at the time you acquire this publication.

The scanning, uploading, and distribution of this book via the Internet or any other means without the permission of the publisher is illegal and punishable by law. Please purchase only authorized electronic editions and do not participate in or encourage electronic piracy of copyrighted materials. Your support of others' rights is appreciated.

U.S. Government License Rights; Restricted Rights: The Software and its documentation is commercial computer software developed at private expense and is provided with RESTRICTED RIGHTS to the United States Government. Use, duplication, or disclosure of the Software by the United States Government is subject to the license terms of this Agreement pursuant to, as applicable, FAR 12.212, DFAR 227.7202-1(a), DFAR 227.7202-3(a), and DFAR 227.7202-4, and, to the extent required under U.S. federal law, the minimum restricted rights as set out in FAR 52.227-19 (DEC 2007). If FAR 52.227-19 is applicable, this provision serves as notice under clause (c) thereof and no other notice is required to be affixed to the Software or documentation. The Government's rights in Software and documentation shall be only those set forth in this Agreement.

SAS Institute Inc., SAS Campus Drive, Cary, NC 27513-2414

June 2023

SAS® and all other SAS Institute Inc. product or service names are registered trademarks or trademarks of SAS Institute Inc. in the USA and other countries. ® indicates USA registration.

Other brand and product names are trademarks of their respective companies.

v_039-P1:dplyml0phy0dkr

Contents

Chapter 1 / Pre-installation Tasks	1
ingress-nginx Vulnerability Mitigation	2
Confidential Computing	3
Plan the Workload Placement	4
Deploy the SAS Viya Platform Deployment Operator	10
Retrieve Required Files	18
Preparing for OpenShift	20
Directory Structure	27
 Chapter 2 / Installation	 31
Initial kustomization.yaml File	32
Common Customizations	36
Install the Orchestration Tool	77
Create the SASDeployment Custom Resource	77
Deploy the Software	84
Readiness Service	92
Sign In as the sasboot User	94
Promotion and Migration	95
 Chapter 3 / Post-Installation Tasks	 97
Configure Identities	98
Configure Multi-tenancy	99
Configure the Connection to the Mail Service	100
Configure Files Service	101
Configure Monitoring and Logging	101
Configure Guest Access	101
Obtain and Run Hadoop Tracer Script	102
Configure Cloud Analytic Services (CAS)	106
Configure Cloud Data Exchange	107
Configure Model Access	108
Configure SAS Asset and Liability Management	109
Configure SAS Business Orchestration Services	109
Configure SAS Data Engineering	110
Configure SAS Data Quality	110
Configure SAS Dynamic Actuarial Modeling	111
Deploy SAS Enterprise Session Monitor	111
Configure SAS Event Stream Processing	111
Configure SAS Expected Credit Loss	112
Configure SAS for Microsoft 365	112
Configure SAS Intelligent Decisioning	112
Configure SAS Model Manager	112
Configure SAS Model Risk Management	113
Configure SAS Risk Engine	113
Configure SAS Risk Modeling	113
Configure SAS Stress Testing	113
Configure SAS Visual Analytics	114
Configure SAS Visual Statistics	114
Configure SAS Viya	115

Configure SAS Viya Advanced	115
Configure SAS Viya Enterprise	116
Configure SAS Viya Programming	117
Configure SAS Viya with SingleStore	118
Configure SAS/CONNECT Spawner	118
Chapter 4 / Validating the Deployment	119
The SAS Viya Platform: Deployment Validation	119
The SAS Viya Platform and the SAS Operational Quality Tool	119
Chapter 5 / Modifying the Deployment Configuration	121
Modify Existing Customizations in a Deployment	121
Chapter 6 / Uninstalling	123
Uninstall with the SAS Viya Platform Deployment Operator	123
Remove the SAS Viya Platform Deployment Operator Only	128
Uninstall Deployments That Do Not Use the Deployment Operator	129
Appendix 1 / (Optional) Using a Mirror Registry	133
Create a Mirror Registry	133
Create and Populate a Mirror Registry in Microsoft Azure	141
Create and Populate a Mirror Registry in Amazon ECR	142
Create and Populate a Mirror Registry in Google Cloud Platform	144
Create and Populate a Mirror Registry in Red Hat OpenShift	145
Create and Populate a Mirror Registry in JFrog Artifactory	147
Create a Mirror Registry at a Dark Site	147
Appendix 2 / SAS Viya Deployment Operator Fields and Messages	149
Overview	149
Fields in the SASDeployment Custom Resource	150
Manage Updates	152
Environment Variables for the Operator Pod	153
Communications from the Operator	156
Appendix 3 / Change Deployment Methods	161
Kubernetes Commands to SAS Viya Platform Deployment Operator	161
Kubernetes Commands to sas-orchestration Command	163
SAS Viya Platform Deployment Operator to Any Other Method	163
Appendix 4 / Troubleshooting	165
Deployment Operator Custom Resource Is Too Large	165
Error Checking Database Connection Error Message	166
Certframe Container Fails at init Stage	166
Issues with External PostgreSQL Server Size and Connections	167
SAS Viya Platform Deployment Operator Reconcile Failures	168
Performing Updates Produces Many Error Messages	168
Failure to Create Symbolic Link	169
Long Delay for SAS/CONNECT Server Sign-on from SAS Studio	169
"Operation not permitted" Error in SAS Configurator for Open Source Logs	170
Uninstall Leaves the Deployment Namespace in a Terminating State	170

Pre-installation Tasks

<i>ingress-nginx Vulnerability Mitigation</i>	2
<i>Confidential Computing</i>	3
<i>Plan the Workload Placement</i>	4
Introduction	4
Assign Nodes by Class	4
<i>Deploy the SAS Viya Platform Deployment Operator</i>	10
Overview of the SAS Viya Platform Deployment Operator	10
Retrieve the Files Required by the SAS Viya Platform Deployment Operator	11
Edit the transformer.yaml File	13
Edit the kustomization.yaml for Cluster-Wide Mode	14
Edit the kustomization.yaml for seccomp	14
Configure the Mirror Registry	14
Add an imagePullSecret	16
Configure Proxy Information	17
Apply the SAS Viya Platform Deployment Operator Resources to the Cluster	18
<i>Retrieve Required Files</i>	18
<i>Preparing for OpenShift</i>	20
Workload Node Placement Considerations	20
Security Context Constraints and Service Accounts	21
Enable hostPath Mounts for CAS	25
Networking	26
TLS	26
Additional Security	27
<i>Directory Structure</i>	27
\$deploy/sas-bases Directory	27
\$deploy/site-config Directory	28
Directories for the SAS Viya Platform Deployment Operator	29

ingress-nginx Vulnerability Mitigation

The SAS Viya platform requires configuration of the ingress-nginx controller. By default, ingress-nginx supports custom snippets, which the SAS Viya platform requires. ingress-nginx has discovered that the default setting of "allow-snippets=true" exposes the deployment to known vulnerability CVE-2021-25742. A user with Kubernetes namespace administrator privileges could use custom snippets to exfiltrate the ingress-nginx service account token and gain access to all secrets in the cluster.

Note: For more information about the CVE, see [CVE-2021-25742 Detail](#).

The SAS Viya platform supports a mitigation for CVE-2021-25742, which involves the application of a block list. SAS recommends that you implement it as soon as possible by following the steps described in this section.

Note: Enabling this mitigation affects all applications in the Kubernetes cluster that are using the ingress-nginx controller.

- 1 Ensure that your SAS Viya platform deployment and the cluster it is being deployed in meets all of the following criteria before applying this mitigation:
 - A supported version of the SAS Viya platform must either be deployed, or you must be about to deploy a supported version.
 - The Kubernetes cluster does not have unsupported deployments of the SAS Viya platform that lack the mitigation and that are sharing the ingress-nginx controller.
 - The ingress-nginx controller is a supported version.
 - The Kubernetes cluster does not have any additional applications using the ingress-nginx controller that do not support using the "annotation-value-word-blocklist" annotation.

If the cluster and your SAS Viya platform deployment do not meet these criteria, do not proceed with applying the mitigation.

- 2 Add the "annotation-value-word-blocklist" annotation to the ConfigMap for the ingress-nginx controller in your Kubernetes cluster. Here is an example command:

```
kubectl patch cm ingress-nginx-controller -n ingress-nginx -p '{"data":
{"allow-snippet-annotations":"true", "annotation-value-word-
blocklist":"load_module, lua_package, _by_lua, location, root, proxy_pass, service
account, {, }, \\""}';
```

In this example, the ConfigMap has the default name of ingress-nginx-controller and is deployed in the ingress-nginx namespace. Modify this command as required for your deployment.

The suggested values for the blocklist are supported by the SAS Viya platform:

```
"load_module, lua_package, _by_lua, location, root, proxy_pass, serviceaccount, {, }, \"
```

Confidential Computing

Note: Confidential computing is only supported on Microsoft Azure at this time. If you are not deploying on Microsoft Azure, skip this section.

Confidential computing helps you to protect your data and code from the time it is created until it is destroyed. To support confidential computing SAS uses AMD SEV technology, which supports a new set of security features that protects data-in-use with full VM memory encryption.

To enable confidential computing on Microsoft Azure:

- 1 Ensure that you have selected VM instance types that support confidential computing with the SAS Viya platform. For more information, see [“Requirements for Confidential Computing” in System Requirements for the SAS Viya Platform](#).
- 2 Ensure that the default (system) node pool and the user node pools contain VM instance types that support confidential computing. For details, see [Create confidential VM on AMD in the Azure portal](#).
- 3 If you are using the SAS Viya platform Infrastructure as Code (IaC) project to provision your deployment, make the following changes to the terraform.tfvars file in the IaC script:

- a Update the Default Node pool VM type to one of the AMD SEV confidential VMs. Here is an example:

```
default_nodepool_vm_type = "Standard_DC8as_v5"
```

- b Update the node pool availability zone and location. Here is an example:

```
default_nodepool_availability_zones = ["1", "2", "3"]
location = "East US"
```

If you are unsure of the availability zone, set the value to empty "".

- c Update the storage type. Using `storage_type=standard` creates an NFS server VM. Using `storage_type=ha` creates Azure NetApp files.

Note: If you are deploying using the viya4-deployment GitHub project, do not attempt to configure confidential computing for the NFS server or Jump server. This configuration is not supported. Use NetApp files for secure storage instead.

CAUTION

After you have deployed your software, if you restart your Microsoft Azure cluster, the default `netapp_protocols=["NFSv3"]` might cause file lock issues for SAS Configuration Server (Consul), Redis, and the logging/

monitoring servers. These issues can result in a non-functional deployment.

Plan the Workload Placement

Note: If you have used a GitHub project, SAS Viya Infrastructure as Code (IaC), for one of the cloud providers to provision your infrastructure components, the procedures described in this section have already been performed for you. You might still want to verify the taints and labels with the following command:

```
kubectl get nodes -o=custom-  
columns=NAME:.metadata.name,LABELS:.metadata.labels,TAINTS:.spec.taints
```

For more information about commands to validate the workload node placement in your infrastructure, see [“Place the Workload on Nodes” on page 9](#). For more information about IaC, see [“Help with Cluster Setup” in *System Requirements for the SAS Viya Platform*](#).

Introduction

Before you deploy your SAS Viya platform, you should develop a plan to distribute the workload across your deployment. Labeling and tainting the nodes prepares the environment to use your plan when the manifest is applied during the deployment process. This section explains how to distribute the workload across your deployment according to several factors.

Assign Nodes by Class

Overview

The SAS Viya platform consists of multiple workload classes. Each class of workload has a unique set of attributes that you must consider when planning your deployment. When planning the placement of the workload, it is helpful to think beyond the initial deployment and to also consider the Kubernetes maintenance life cycle.

IMPORTANT SAS strongly recommends labeling and tainting all your nodes, especially the CAS nodes. Note, however, that the connect workload class is only required if you are not using dynamically launched pods. For

more information about dynamically launched pods, see [“connect” on page 6](#).

Workload Classes

Properties associated with the pods in your deployment describe the type of work that they perform. By using the commands described in [“Place the Workload on Nodes” on page 9](#), you can configure where you want each class of pods to run, thereby enabling you to manage the associated workload.

Here are the SAS Viya platform workload classes:

stateless

These workloads include web applications and microservices. Stateless applications scale horizontally in order to handle increased workload and to provide higher availability. By default, one replica of each stateless workload is deployed. A high-availability (HA) scaling overlay is provided with the deployment artifacts for each SAS Viya platform software order.

stateful

These workloads include open-source components that are used by the SAS Viya platform in order to store critical information that is required to maintain operational integrity. These workloads are deployed in an HA configuration by default.

cas

SAS Cloud Analytic Services (CAS) is the primary compute engine in the SAS Viya platform. Although CAS is a stateful process, it has unique attributes and requirements within the Kubernetes infrastructure. CAS has been engineered to use all available resources to complete each request in the shortest amount of time. Therefore, CAS can quickly drive every available CPU core to 100% utilization. CAS also aggressively consumes RAM to process large volumes of data as quickly as possible.

Because CAS was designed with performance as a primary objective, it must be hosted on dedicated resources. Therefore, each CAS pod must be placed on its own dedicated node. In addition, if you are not running with auto-resourcing (which is enabled by default), SAS recommends that you configure the CAS server with guaranteed quality of service (QoS) in order to prevent CAS pods from being evicted by Kubernetes.

Note: The requirement to place each CAS pod on its own dedicated node does not apply to the personal CAS server. For more information about the personal CAS server, see [“Create a Personal CAS Server” on page 64](#).

For increased granularity, you can also use the cascontroller and casworker classes. The same restrictions described for the cas class also apply to cascontroller and casworker.

compute

Compute processes are stateful processes that run to completion (that is, they are not long-running services). They deliver multiple functions within a SAS Viya platform deployment, including but not limited to batch jobs, interactive SAS

Studio sessions, and parallel compute processes that are executed as part of a SAS Model Studio workflow.

connect

SAS/CONNECT provides the ability for clients to launch SAS compute processes, the same processes described in “[compute](#)” on page 5, within a local or remote SAS environment. The SAS/CONNECT Spawner is a SAS Viya platform service that launches these processes on behalf of SAS/CONNECT clients. The processes can be launched in their own pods (referred to as “dynamically launched pods”) or in the SAS/CONNECT Spawner pod. SAS 9 clients from releases before SAS 9.4M7 and SAS Viya 3 clients from versions before SAS Viya 3.5 do not support dynamically launched pods, so by default their processes are launched in the SAS/CONNECT Spawner pod. For other clients, dynamically launched pods are enabled.

IMPORTANT If you are using dynamically launched pods, you do not need to create a connect workload class and should therefore skip the rest of this description of the connect workload class.

If you are using external SAS/CONNECT clients with your SAS Viya platform deployment, the SAS/CONNECT Spawner requires more resources in order to avoid being evicted by Kubernetes. Like CAS, the SAS/CONNECT Spawner pod requires a dedicated Kubernetes node, and SAS recommends that it be configured with Kubernetes guaranteed QoS.

If you do use the connect workload class, you must add a transformer to the base kustomization.yaml file. For more details, see the README file at `$deploy/sas-bases/examples/sas-connect-spawner/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configure_sasconnect_spawner_in_the_sas_viya_platform.htm` (for HTML format).

Default Node Pool Configuration

SAS recommends that you use a node pool that does not run SAS workloads in the cluster. This node pool is referred to as the *default node pool*. The default node pool is useful for ingress controllers and for hosting LDAP and monitoring stacks. SAS also recommends that the default node pool be labeled with `kubernetes.azure.com/mode=system` so that SAS pods are not scheduled there.

This label is automatically applied to the default Microsoft Azure Kubernetes Service system node pool. If you are provisioning the cluster with the SAS-provided IaC tooling, the label is automatically applied to the equivalent default node pool for other infrastructures. If you are deploying on an infrastructure other than Microsoft Azure Kubernetes Service and you are not provisioning using IaC, you must apply the label yourself with the following command:

```
kubectl label nodes node-name kubernetes.azure.com/mode=system
```

Note: The *node-name* can be the name of a single node or it can be a list of node names separated by a space. The term **azure** is applied to all of the supported cloud platforms because it represents a default SAS Viya platform deployment setting that is not vendor-specific.

Considerations and Default Settings for Workload Placement

The Kubernetes cluster administrator will recognize that each of the SAS Viya platform workload classes should be managed differently. In situations where the deployment has rigorous availability requirements and the Kubernetes cluster software is being upgraded, the different requirements might resemble the following:

- Nodes that host the stateless and stateful workload class can be individually drained and updated with no additional management concerns as long as the HA transformer was used for the deployment.
- Nodes that host the compute workload class should be cordoned. The workloads that run on those nodes should be allowed to run to completion in order to avoid interrupting work in progress. Examples are a SAS Studio session, a batch job, or a SAS Model Studio session. When all compute processes have completed, the node can be drained so that maintenance tasks can be performed.

Each node that hosts compute class workloads can be cordoned, drained, or updated in a manner that does not disrupt end-user activity as long as each compute session is allowed to run to completion.

Note: If your deployment includes SAS Workload Management, you must have at least one node labeled for the compute workload class.

- Nodes that host workloads for the cas and connect classes must be terminated and started on different nodes before the underlying nodes can be maintained as required.

Some aspects of workload placement are set by default during the SAS Viya platform deployment, including both nodeAffinity and tolerations. In addition, SAS Viya platform workloads are not scheduled on the system node pool in Microsoft Azure, the default node group in Amazon Web Services (AWS), or the default node pool in Google Cloud Platform (GCP).

The following table summarizes additional implications of the SAS Viya platform default deployment for each workload class:

Table 1.1 Default Taints and nodeAffinity Settings per Class

Workload Class	Default Settings
Stateless workloads	<p>Prefer to schedule on nodes that are labeled workload.sas.com/class=stateless</p> <p>Tolerate the following taints:</p> <ul style="list-style-type: none"> ■ workload.sas.com/class=stateless:NoSchedule ■ workload.sas.com/class=stateful:NoSchedule
Stateful workloads	<p>Prefer to schedule on nodes that are labeled workload.sas.com/class=stateful</p> <p>Tolerate the following taints:</p>

Workload Class	Default Settings
	<ul style="list-style-type: none"> ■ workload.sas.com/class=stateful:NoSchedule ■ workload.sas.com/class=stateless:NoSchedule
Compute workloads	<p>Prefer to schedule on nodes that are labeled workload.sas.com/class=compute</p> <p>Tolerate the taint workload.sas.com/ class=compute:NoSchedule</p>
CAS workloads	<p>Prefer to schedule on nodes that are labeled workload.sas.com/class=cas</p> <p>Tolerate the taint workload.sas.com/ class=cas:NoSchedule</p>
Connect workloads Note: This workload class is not required if you are using dynamically launched pods.	<p>Prefer to schedule on nodes that are labeled workload.sas.com/class=connect</p> <p>Tolerate the taint workload.sas.com/ class=connect:NoSchedule</p>

To obtain further granularity for scheduling your CAS controllers and CAS workers, divide your CAS workloads between cascontroller and casworker workload groups.

Table 1.2 Default Taints and nodeAffinity Settings for Cascontroller and Casworker Classes

Workload Class	Default Settings
Cascontroller workloads	<p>Prefer to schedule on nodes that are labeled workload.sas.com/class=cascontroller</p> <p>Tolerate the following taints:</p> <ul style="list-style-type: none"> ■ workload.sas.com/class=cas:NoSchedule ■ workload.sas.com/ class=cascontroller:NoSchedule
Casworker workloads	<p>Prefer to schedule on nodes that are labeled workload.sas.com/class=casworker</p> <p>Tolerate the following taints:</p> <ul style="list-style-type: none"> ■ workload.sas.com/class=cas:NoSchedule ■ workload.sas.com/class=casworker:NoSchedule

By default, the CAS operator uses auto-resourcing to utilize the full resources of a node. When a CAS workload is scheduled to run on a specific node, no other application pods can be scheduled to run on that node. If you want to disable auto-resourcing and modify resource requests manually, consider configuring guaranteed QoS on these pods. For more information, see the README file located at `sas-bases/overlays/cas-server/auto-resources/README.md` (for Markdown format) or `$deploy/sas-bases/docs/`

`auto_resources_for_cas_server_for_the_sas_viya_platform.htm` (for HTML format).

If you deploy the SAS Viya platform into a cluster with untainted nodes, the Kubernetes scheduler may attempt to schedule CAS onto those untainted nodes instead of the nodes that have been labeled with `workload.sas.com/class=cas`. An overlay that ensures that CAS servers are only scheduled onto nodes that have the dedicated label is available. For more information, see the “Optional CAS Server Placement Configuration” section of the README file located at `$deploy/sas-bases/overlays/cas-server/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/mpp_cas_server_for_sas_viya.htm` (for HTML format).

Multi-tenancy

If you are enabling multi-tenancy in your deployment, the tenants will share most of the nodes with the provider tenant. However, because each tenant has its own CAS server, the total number of nodes required for CAS for the full deployment is greater than that for a non-multi-tenant deployment. The number of additional CAS nodes required per tenant depends on whether the tenant is deploying SMP or MPP CAS.

Place the Workload on Nodes

SAS requires that you identify the node or nodes on which CAS pods should be scheduled. SAS further recommends that you identify the nodes on which all classes should be scheduled. The following commands place the workload on the nodes in your deployment:

To list the names of the nodes:

```
kubectl get nodes
```

To assign a class of pods to a specific node:

```
kubectl taint nodes node-name workload.sas.com/class=class-name:NoSchedule --overwrite
kubectl label nodes node-name workload.sas.com/class=class-name --overwrite
```

Note: The *node-name* can be the name of a single node, or it can be a list of node names separated by spaces.

Here is an example that assigns only CAS pods to a node named node1.

```
kubectl taint nodes node1 workload.sas.com/class=cas:NoSchedule --overwrite
kubectl label nodes node1 workload.sas.com/class=cas --overwrite
```

To verify the nodes that have been tainted:

```
kubectl get nodes -o=custom-columns=NAME:.metadata.name,TAINTS:.spec.taints
```

To view node labels and taints at the same time:

Note: Because of the length of the command and the margin of the page, this command appears as more than one line. The command should be entered as a single line.

```
kubectl get nodes -o=custom-  
columns=NAME:.metadata.name,LABELS:.metadata.labels,TAINTS:.spec.taints
```

After you have assigned the pods to nodes, to check the assignments:

Note: Because of the length of the command and the margin of the page, this command appears as more than one line. The command should be entered as a single line.

```
kubectl -n name-of-namespace get pod -o=custom-  
columns=NAME:.metadata.name,NODE:.spec.nodeName | sort
```

Deploy the SAS Viya Platform Deployment Operator

Overview of the SAS Viya Platform Deployment Operator

The SAS Viya Platform Deployment Operator provides an automated method for deploying and updating your SAS Viya platform deployment. It runs in the Kubernetes cluster and watches for declarative representations of SAS Viya platform deployments in the form of custom resources (CRs) of the type SASDeployment. The operator can watch for SASDeployments in the namespace where it is deployed (namespace mode) or in all the namespaces in a cluster (cluster-wide mode). Using namespace mode means that the operator and the SAS Viya platform deployment are located in the same namespace. Operators in cluster-wide mode have their own namespaces.

Note: SAS recommends using only one mode of operator in a cluster.

When a new SASDeployment is created or an existing SASDeployment is updated, the operator updates the SAS Viya platform deployment to match the state that is described in the CR. The operator determines the release of the SAS Viya platform that it is working with. It also uses the appropriately versioned tools for that release while deploying and updating the SAS Viya platform software. Thus, a single instance of the operator can manage all SAS Viya platform deployments in the cluster.

The SAS Viya Platform Deployment Operator requires updates occasionally but not at the same frequency as the SAS Viya platform.

If you want to consider using the SAS Viya Platform Deployment Operator to manage your SAS Viya platform deployments, continue reading this section for its deployment steps. If you want to deploy and update your software using Kubernetes commands or the sas-orchestration tool, skip this section and go to [“Retrieve Required Files” on page 18](#).

Retrieve the Files Required by the SAS Viya Platform Deployment Operator

Note: If you plan to use a mirror registry for the deployment of your SAS Viya platform software, you should download and install SAS Mirror Manager before you begin your deployment. For more information, see [“Create a Mirror Registry” on page 133](#).

The required deployment assets are delivered in a TGZ file that is accessible from the my.sas.com portal.

- 1 If you plan to use a cluster-wide mode SAS Viya Platform Deployment Operator and you have not already created a namespace that is dedicated to the deployment of the SAS Viya Platform Deployment Operator, create that namespace now. If you plan to use a namespace mode SAS Viya Platform Deployment Operator and you have not created a namespace for your SAS Viya platform deployment, create that namespace now.

Note: This namespace is referred to as `sasoperator` in the examples that follow.

- 2 Create a directory on the kubectl machine or on a machine that can be reached by your kubectl machine:

```
mkdir directory-name
```

Note: If you created a mirror registry, the new directory should be parallel to `$deploy`.

SAS recommends that you name the directory `operator-deploy`, but you should use a name that is meaningful to you. The directory is referred to as `$operator-deploy` in this guide. Replace `$operator-deploy` with the directory name that you prefer.

- 3 Click the **Get Started** link that is provided in your Software Order Email (SOE). The my.sas.com page opens.
- 4 Examine the order information. Ensure that you are working with the order you expect.
- 5 Select the **Downloads** tab. The line above the **Asset Type** table indicates the release cadence and the version of SAS Viya platform software that you will deploy. If you want to deploy a different version, find the listing for that version on the page.

Note: If you are changing your deployment method to use the deployment operator prior to performing an update, be sure to download the deployment

assets for the version you intend to update to. The target version should be used for deploying the deployment operator.

- 6 Select **Deployment Assets** from the **Asset Type** list, then click **Download** to download the files that are required to deploy your software.
- 7 Save the TGZ file from the my.sas.com page to the directory that you created in step 2.
- 8 Extract the files from the TGZ file in the same directory:

```
tar xvfz file-name.tgz
```

The result is a directory structure that looks like this:

```
$operator-deploy/
└─ sas-bases/
   │ base/
   │ components/
   │ docs/
   │ examples/
   │ extras/
   └─ overlays/
```

- 9 Copy the operator files to the top level of the \$operator-deploy directory and make them writable:

```
cd $operator-deploy
cp -r sas-bases/examples/deployment-operator/deploy/* .
chmod +w site-config/transformer.yaml
```

If you are deploying the deployment operator in cluster-wide mode or into a cluster that does not support seccomp, the operator kustomization.yaml file should also be made writable. In such a scenario, this would be the last line in the previous command:

```
chmod +w site-config/transformer.yaml kustomization.yaml
```

The result is a directory structure that looks like this:

```
$operator-deploy/
├─ kustomization.yaml
├─ operator-base/
├─ sas-bases/
└─ site-config/
   │ cluster-wide-transformer.yaml
   │ remove-seccomp-transformer.yaml
   └─ transformer.yaml
```

The kustomization.yaml file in the \$operator-deploy directory is referred to as the operator kustomization.yaml file throughout the documentation.

Note: Verify that you are using the correct directory for your changes: `$operator-deploy/site-config` or `$deploy/site-config`.

Edit the transformer.yaml File

Replace the variables in the `$operator-deploy/site-config/transformer.yaml` file as indicated in that file. Two patches must be updated.

- The first patch sets the name of the ClusterRoleBinding resource that is needed by the operator. The SAS Viya Platform Deployment Operator requires cluster-wide permissions even when the operator is deployed in namespace mode because the SAS Viya platform deployment requires the deployment of cluster resources. Because ClusterRoleBinding is a cluster-wide resource, it must be unique within the cluster. Each SAS Viya Platform Deployment Operator instance must have a unique ClusterRoleBinding.
- The second patch sets the name of the namespace in which the operator is being deployed, and therefore where the RBAC resources, such as ServiceAccount, are deployed

Here is an example of the transformer.yaml file. It contains the name *sasoperator* for the ClusterRoleBinding and the namespace *sasoperator*.

Note: If you are deploying the operator in namespace mode, the namespace value is also used in the base kustomization.yaml file since you are also deploying the SAS Viya platform into it.

```
---
apiVersion: builtin
kind: PatchTransformer
metadata:
  name: patch-transformer-sas-deployment-operator
patch: |-
  # name: MUST BE PROVIDED BY USER AND UNIQUE IN CLUSTER.
  - op: replace
    path: /metadata/name
    value:
      "sasoperator"
  # namespace: MUST BE PROVIDED BY USER. THIS MUST BE THE NAMESPACE INTO WHICH
  # THE SAS DEPLOYMENT OPERATOR IS DEPLOYED.
  - op: add
    path: /subjects/0/namespace
    value:
      "sasoperator"
target:
  annotationSelector: sas.com/component-name=sas-deployment-operator
  # Some required components for SAS Viya are scoped to the cluster. In order to
  # deploy your software, the operator requires the ClusterRoleBinding, regardless
  # of whether the intended scope is to the namespace or the cluster.
  kind: ClusterRoleBinding
```

Edit the kustomization.yaml for Cluster-Wide Mode

If the operator is being deployed in namespace mode, skip this section. However, if the operator is being deployed in cluster-wide mode, edit the operator kustomization.yaml to uncomment the reference to `site-config/cluster-wide-transformer.yaml`. Here is an example of the file when deploying in cluster-wide mode.

```
resources:
- operator-base
transformers:
- site-config/transformer.yaml
## Uncomment the following inclusion if you are deploying the
## operator in cluster-wide mode.
- site-config/cluster-wide-transformer.yaml
```

Edit the kustomization.yaml for seccomp

If you are using the operator in a cluster that does not support secure computing mode (seccomp), uncomment the relevant line in the operator kustomization.yaml file. Red Hat OpenShift does not support seccomp.

Here is an example:

```
transformers:
...
## Uncomment the following inclusion if you are deploying the
## operator in a cluster that does not support seccomp.
- site-config/remove-seccomp-transformer.yaml
```

Configure the Mirror Registry

Note: If you not using a mirror registry for your deployment, skip this section.

- 1 Copy the mirror configuration yaml file from `$operator-deploy/sas-bases/examples/mirror/` to the `$operator-deploy/site-config` directory. If you are using any cloud provider except Red Hat OpenShift, copy the `mirror.yaml` file. If you are using Red Hat OpenShift, copy the `mirror-flattened.yaml` file.
- 2 Revise the mirror configuration file based on the infrastructure you are using for your deployment.
 - If you are deploying on any cloud provider other than Red Hat OpenShift, open the `mirror.yaml` file and replace each instance of `{{ MIRROR-HOST }}` with the fully qualified domain name (FQDN) of the mirror registry in this format:

`FQDN/registry-namespace/sasdeployment`

Here is an example:

`example.company.com/registry-namespace/sasdeployment`

After replacing all instances of `{{ MIRROR-HOST }}`, save and close the file.

- If you are deploying on Red Hat OpenShift, modify the `mirror-flattened.yaml` file.

1 Find the image registry name and port:

```
oc get svc -n name-of-registry-namespace
```

The output looks like this:

Name	Type	Cluster-IP	External-IP	Ports	Age
image-registry	ClusterIP	172.30.146.74	<none>	5000/TCP	174m

- 2** Open the `mirror-flattened.yaml` file and replace each instance of `{{ MIRROR-HOST }}` with the internal route to the registry from inside the cluster followed by the name of the SAS Viya platform namespace. Use this format:

```
service-name.name-of-registry-namespace.svc:port/platform-namespace
```

The *service-name* and *port* values are in the output from the command in the previous step. The *name-of-registry-namespace* is the same one that you used in the command itself. The *platform-namespace* is the namespace where you will deploy the SAS Viya platform. Here is an example using the sample output from the previous step:

```
image-registry.openshift-image-registry.svc:5000/myviya
```

- 3** After replacing all instances of `{{ MIRROR-HOST }}`, save and close the `mirror-flattened.yaml` file.
- 3** In the operator `kustomization.yaml` file, add the path to mirror configuration `yaml` file as the last entry in the `transformers` section. For example, if you are deploying on any cloud provider except Red Hat OpenShift and you copied the `mirror.yaml` file to the top level of the `$operator-deploy/site-config` directory, the entry would look like this:

```
transformers:
- site-config/transformer.yaml
- site-config/mirror.yaml
## Uncomment the following inclusion if you are deploying the
...
```

If you are deploying on Red Hat OpenShift and you copied the `mirror-flattened.yaml` file to the top level of the `$operator-deploy/site-config` directory, the entry would look like this:

```
transformers:
- site-config/transformer.yaml
- site-config/mirror-flattened.yaml
## Uncomment the following inclusion if you are deploying the
...
```

Add an imagePullSecret

Note: If you are not using a mirror registry or if you are using a mirror registry but authentication is not required, skip this section.

If SAS content has been mirrored, and that mirror requires authentication, you must configure an imagePullSecret and patch that secret into the SAS Viya Platform Deployment Operator's Deployment resource.

All Cloud Providers

For information about creating an imagePullSecret, see [Pull an Image from a Private Registry](#). The secret should be added to the SAS Viya Platform Deployment Operator's Deployment resource in the operator kustomization.yaml file. Here is an example of the resulting kustomization.yaml file:

```
resources:
- operator-base
transformers:
- site-config/transformer.yaml
## Uncomment the following inclusion if you are deploying the
## operator in clusterwide mode.
- site-config/cluster-wide-transformer.yaml
secretGenerator:
- name: site-image-pull-secret
  type: kubernetes.io/dockerconfigjson
  literals:
  - .dockerconfigjson={{ MY-IMAGEPULLSECRET }}
patches:
- patch: |-
  - op: add
    path: /spec/template/spec/imagePullSecrets
    value:
      - name: site-image-pull-secret
target:
  kind: Deployment
  name: sas-deployment-operator
```

Red Hat OpenShift Alternative

If you are deploying on Red Hat OpenShift, you can use an existing secret that OpenShift includes instead of creating a new one.

1 Find the name of the secret:

```
kubectl -n name-of-namespace get secret | grep default-dockercfg
```

The output looks like this:

```
default-dockercfg-#####
```

- 2 Edit the operator kustomization.yaml file by adding a patch that contains a reference to the secret:

```
patches:
- patch: |-
  - op: add
    path: /spec/template/spec/imagePullSecrets
    value:
      - name: output-of-step-1
target:
  kind: Deployment
  name: sas-deployment-operator
```

Configure Proxy Information

Note: If you are not using a forward proxy for your cluster, skip this section. If you would like to use a proxy but limit it to the deployment you are working on, add proxy variables to your SASDeployment Custom Resource. For details, see [“Revise the Custom Resource for Proxy Information” on page 83](#).

To define the same proxy information for all operator-based SAS Viya platform deployments, you should modify the deployment operator manifest file. The manifest file is located at `$operator-deploy/operator-base/deployment.yaml`.

The deployment.yaml file is initially organized like this:

```
containers:
- args:
  - reconcile
  command: []
  env:
```

In the `env` section, add the following environment variables:

```
containers:
- args:
  - reconcile
  command: []
  env:
    - name: HTTP_PROXY
      value: "proxy-URL-for-HTTP-requests"
    - name: HTTPS_PROXY
      value: "proxy-URL-for-HTTPS-requests"
    - name: NO_PROXY
      value: "do-not-proxy-list"
  ...
```

The `do-not-proxy-list` is a comma-separated list of host names, fully qualified host names, and IP addresses that the proxy should ignore.

Here is an example of the deployment.yaml file with values:

```

containers:
  - args:
    - reconcile
    command: []
    env:
    - name: HTTP_PROXY
      value: "http://webproxy.example.com:5000"
    - name: HTTPS_PROXY
      value: "http://webproxy.example.com:5000"
    - name: NO_PROXY
      value: "localhost,noproxy.example.com,apiserver.example.com,
kubernetes.default.svc,10.96.0.1"
...

```

Note: The following values must be included in the list of values for the NO_PROXY variable:

- kubernetes.default.svc (the Kubernetes API server)
 - the value of the KUBERNETES_SERVICE_HOST environment variable for the cluster
-

Apply the SAS Viya Platform Deployment Operator Resources to the Cluster

Run the following command from the \$operator-deploy directory to deploy the SAS Viya Platform Deployment Operator into Kubernetes:

```
kustomize build . | kubectl -n name-of-deployment-operator-namespace apply -f -
```

Retrieve Required Files

Note: If you plan to use a mirror registry for the deployment of your SAS Viya platform and you have not already done so, you should download and install SAS Mirror Manager before beginning your deployment. For more information, see [“Create a Mirror Registry” on page 133](#).

- 1 Deploying the SAS Viya platform requires a separate namespace for each deployment of your software. If you are using the SAS Viya Platform Deployment Operator in namespace mode, the SAS Viya platform deployment must be performed in the same namespace that was used for the operator.

If you have not already done so, create a namespace for your SAS Viya platform deployment.

- 2 Create the directory for your deployment assets.

- If you have created a mirror registry, you have already created the `$deploy` directory and should skip to step 3.
- Create a directory on the `kubectl` machine or on a machine that can be reached by your `kubectl` machine:

```
mkdir directory-name
```

Note: If you have deployed the SAS Viya Platform Deployment Operator, the new directory should be parallel to `$operator-deploy`.

SAS recommends that you name the directory `deploy`, but you should use a name that is meaningful to you. The directory is referred to as `$deploy` in this guide. Replace `$deploy` with the directory name that you prefer.

If you are performing multiple deployments of the SAS Viya platform in a cluster, each `$deploy` directory should be named in a way that allows you to associate a directory with a specific deployment.

Note: If you are deploying with the SAS Viya Platform Deployment Operator, you can use a GitOps repository to store your deployment assets rather than create a directory. If you choose to use GitOps, the instructions for `$deploy` throughout the documentation apply to the GitOps directory as well.

- 3 Click the **Get Started** link that is provided in your Software Order Email (SOE). The `my.sas.com` page opens.
- 4 Examine the order information. Ensure that you are working with the order you expect.
- 5 Select the **Downloads** tab. The line above the **Asset Type** table indicates the release cadence and the version of SAS Viya platform software that you will deploy. If you want to deploy a different version, find the listing for that version on the page.
- 6 Select **Deployment Assets** from the **Asset Type** list, then click **Download** to download the files that are required to deploy your software.
Save the TGZ file to the directory that you created in step 2.

- 7 Extract the files from the TGZ file in the same directory:

```
tar xvfz file-name.tgz
```

The result is a directory structure that looks like this:

```
$deploy/
├─ sas-bases/
│   ├── base/
│   ├── components/
│   ├── docs/
│   ├── examples/
│   └─ overlays/
```

For a description of each subdirectory, see [“\\$deploy/sas-bases Directory” on page 27](#).

- 8 If you are performing a deployment with Kubernetes commands, skip this step.

- a Create a new directory parallel to the `$deploy` directory in which to store your license and `*-certs.zip` files.

```
mkdir directory-name
```

You should use a name that is meaningful to you. The directory is referred to as `$license` in this guide. Replace `$license` with the directory name that you prefer.

- b Copy the `*-certs.zip` file and the license file, identified by the `.jwt` extension, to the `$license` directory.

Note: If you plan to have multiple SAS Viya platform deployments in your cluster, you should organize the `$license` directory as you see fit. Whatever strategy you use to organize, ensure that you can easily differentiate the license and `certs.zip` files by order.

- 9 Parallel to the `sas-bases` subdirectory, create a new directory named `site-config`.

For a description of the `$deploy/site-config` directory, see “[\\$deploy/site-config Directory](#)” on page 28.

Preparing for OpenShift

Note: If you are not deploying the SAS Viya platform in a Red Hat OpenShift environment, skip this section.

Deployment steps for OpenShift are different from those required for other infrastructures. This section describes those differences.

Workload Node Placement Considerations

Although the workload placement documentation recommends labelling and tainting all worker nodes, you must have at least two untainted default nodes to allow for the default ingress controller replica count. For more information, see [Ingress controller configuration parameters](#).

Security Context Constraints and Service Accounts

Overview

A deployment in OpenShift requires multiple custom security context constraints (SCCs) to provide permissions to SAS Viya platform services. SCCs are required in order to enable the Pods to run. (Pods provide essential SAS Viya platform components.) In addition, some SCCs can be customized to meet your unique requirements.

A security context acts like a request for privileges from the OpenShift API. In an OpenShift environment, each Kubernetes Pod starts up with an association with a specific SCC, which limits the privileges that Pod can request. An administrator configures each Pod to run with a certain SCC by granting the corresponding service account for that pod access to the SCC. For example, if Pod A requires its own SCC, an administrator must grant access to that SCC for the service account under which Pod A is launched. Use the OpenShift OC administrative command-line tool to grant or remove these permissions.

Note: For additional details about SCC types, see [“SCCs and Pod Service Accounts” in *System Requirements for the SAS Viya Platform*](#).

Apply and Bind the Security Context Constraints

After you have downloaded the deployment assets as described in [“Retrieve Required Files” on page 18](#), apply the following SCCs that pertain to your deployment and then bind the SCC to the appropriate service account.

sas-cas-server

Every deployment on OpenShift must apply one of the following SCCs for the CAS server.

Why the SCC is needed: CAS relies on SETUID, SETGID, and CHOWN capabilities. CAS is launched by a SETUID root executable called caslaunch. By default, caslaunch starts the CAS controller running under the runAsUser/runAsGroup values and a root process named launchsvcs. Caslaunch connects these processes with a pipe. The SETUID and SETGID capabilities are required by launchsvcs in order to launch session processes under the user's operating-system (host) identity instead of launching them using runAsUser/runAsGroup values. The CHOWN capability is necessary to support Kerberos execution, which requires modification of the ownership of the cache file that is created by a direct Kerberos connection. By default, the cache file is owned by runAsUser/runAsGroup identities, but in order to support Kerberos, it must be owned by the user's host identity.

Therefore, at a minimum, the `cas-server-scc` SCC must be applied. If the custom group "CAS Host Account Required" is used, apply the `cas-server-scc-host-launch` SCC. If the CAS server is configured to use a custom SSSD, apply the `cas-server-scc-sssd` SCC. For more information on enabling SSSD, see ["Configure SSSD" on page 74](#).

- 1 Apply the SCC with one of the following commands:

```
kubectl apply -f sas-bases/examples/cas/configure/cas-server-scc.yaml
kubectl apply -f sas-bases/examples/cas/configure/cas-server-scc-host-launch.yaml
kubectl apply -f sas-bases/examples/cas/configure/cas-server-scc-sssd.yaml
```

- 2 Bind the SCC to the service account with the command that includes the name of the SCC that you applied:

```
oc -n name-of-namespace adm policy add-scc-to-user sas-cas-server -z sas-cas-server

oc -n name-of-namespace adm policy add-scc-to-user sas-cas-server-host -z sas-cas-server

oc -n name-of-namespace adm policy add-scc-to-user sas-cas-server-sssd -z sas-cas-server
```

sas-connect-spawner

By default, no SCC is required for SAS/CONNECT and you can skip this item. The SCC is required only if you intend to launch your SAS/CONNECT servers in the Spawner pod, rather than in their own pods. For more information, see the README file at `$deploy/sas-bases/examples/sas-connect-spawner/openshift/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/granting_security_context_constraints_on_an_openshift_cluster_for_sasconnect.htm` (for HTML format).

Why the SCC is needed: The SAS/CONNECT Launcher must be able to launch the SAS/CONNECT Server under end user identity.

- 1 Apply the SCC with the following command:

```
kubectl apply -f sas-bases/examples/sas-connect-spawner/openshift/sas-connect-spawner-scc.yaml
```

- 2 Bind the SCC to the service account with the following command:

```
oc -n name-of-namespace adm policy add-scc-to-user sas-connect-spawner -z sas-connect-spawner
```

sas-esp-project

To determine if your deployment includes SAS Event Stream Processing, look for it in the "License Information" section of your Software Order Email (SOE) for the list of products that are included in your order. If your SOE is unavailable, look for `$deploy/sas-bases/examples/sas-esp-operator` in your deployment assets. If that directory exists, then your deployment includes SAS Event Stream Processing. If it does not exist, skip this SCC.

To run SAS Event Stream Processing projects with a user other than "sas", you must bind the `sas-esp-project` service account to the nonroot SCC.

The nonroot SCC is a standard SCC defined by OpenShift. For more information about the nonroot SCC, see [Managing SCCs in OpenShift](#).

- 1 There is no SCC to apply.
- 2 Bind the SCC to the service account with the following command:

```
oc -n name-of-namespace adm policy add-scc-to-user nonroot -z sas-esp-project
```

sas-microanalytic-score

To determine if the `sas-microanalytic-score` SCC is needed for your deployment, check for a README file in your deployment assets at `$deploy/sas-bases/overlays/sas-microanalytic-score/service-account/README.md`. If the README file is present, then the SCC is available for deployments on OpenShift.

Why the SCC is needed: Security context constraints are required in an OpenShift cluster if the `sas-micro-analytic-score` pod needs to mount an NFS volume. If the Python environment is made available through an NFS mount, the service account requires NFS volume mounting privileges.

For more information, see the README file at `$deploy/sas-bases/overlays/sas-microanalytic-score/service-account/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configure_sas_micro_analytic_service_to_add_service_account.htm` (for HTML format).

- 1 Apply the SCC with the following command:

```
kubectl apply -f sas-bases/overlays/sas-microanalytic-score/service-account/sas-microanalytic-score-scc.yaml
```

- 2 Bind the SCC to the service account with the following command:

```
oc -n name-of-namespace adm policy add-scc-to-user sas-microanalytic-score -z sas-microanalytic-score
```

sas-model-publish-kaniko

To determine if the `sas-model-publish-kaniko` service account exists in your deployment, check for a README file in your deployment assets at `$deploy/sas-bases/examples/sas-model-publish/kaniko/README.md`. If the README file is present and you plan to publish models with SAS Model Manager or SAS Intelligent Decisioning to containers using kaniko, you must bind the `sas-model-publish-kaniko` service account to the anyuid SCC. The anyuid SCC is a standard SCC defined by OpenShift. For more information about the anyuid SCC, see [Managing SCCs in OpenShift](#). For more information about publishing models with kaniko, see the README at `$deploy/sas-bases/examples/sas-model-publish/kaniko/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configure_kaniko_for_sas_model_publish_service.htm` (for HTML format).

- 1 There is no SCC to apply.
- 2 Bind the service account with the following command:

```
oc -n name-of-namespace adm policy add-scc-to-user anyuid -z sas-model-publish-kaniko
```

sas-model-repository

To determine if the `sas-model-repository` SCC is needed for your deployment, check for a README file in your deployment assets at `$deploy/sas-bases/overlays/sas-model-repository/service-account/README.md`. If the README file is present, then the SCC is available and might be required for deployments on OpenShift.

Why the SCC is needed: The `sas-model-repository` pod requires a service account with privileges if the Python environment is made available through an NFS mount. NFS volumes are not permitted in the restricted SCC, so an SCC that has NFS in the allowed volumes section is required.

For more information, see the README file at `$deploy/sas-bases/overlays/sas-model-repository/service-account/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configure_sas_model_repository_service_to_add_service_account.htm` (for HTML format).

- 1 Apply the SCC with the following command:

```
kubectl apply -f sas-bases/overlays/sas-model-repository/service-account/
sas-model-repository-scc.yaml
```

- 2 Bind the SCC to the service account with the following command:

```
oc -n name-of-namespace adm policy add-scc-to-user sas-model-repository -z
sas-model-repository
```

sas-opendistro

Every SAS Viya platform deployment contains `sas-opendistro`. If your deployment uses an internal instance of OpenSearch, the `sas-opendistro` is required.

Deploying OpenSearch on OpenShift requires changes to a few kernel settings. The `sysctl-transformer.yaml` file can apply the necessary `sysctl` parameters to configure the kernel, but it requires the following special privileges:

`allowPrivilegeEscalation` option enabled, `allowPrivilegedContainer` option enabled, and `runAsUser` set to `RunAsAny`.

Therefore, before you apply the `sas-opendistro-scc.yaml` file, you must modify it to enable these privileges. For the instructions to modify them, see the “Modify `sas-opendistro-scc.yaml` for `sysctl-transformer.yaml`” section of the README file at `$deploy/sas-bases/examples/configure-elasticsearch/internal/openshift/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/openshift_on_red_hat_openshift.htm` (for HTML format).

Why the SCC is needed: For optimal performance, deploying OpenSearch software requires the `vm.max_map_count` parameter to be set on the Kubernetes nodes running the stateful workloads to ensure there is adequate virtual memory available for use with `mmap` for accessing the search indices. To provide a method to set this argument automatically, the SAS Viya platform includes an optional transformer as part of the internal-elasticsearch overlay that adds an init container to automatically set this parameter. The init container must be run at a privileged level (using both the `privileged = true` and `allowPrivilegeEscalation = true` security context options) since it modifies the kernel parameters of the host. The container terminates after it sets the kernel parameter, it terminates, and the OpenSearch software will then proceed to start as a non-privileged container.

- 1 Apply the SCC with the following command:

```
kubectl apply -f sas-bases/examples/configure-elasticsearch/internal/
openshift/sas-opendistro-scc.yaml
```

- 2 Bind the SCC to the service account with the following command:

```
oc -n name-of-namespace adm policy add-scc-to-user sas-opendistro -z sas-
opendistro
```

SAS Watchdog

SAS Watchdog is included in every SAS Viya platform deployment. If you choose to deploy SAS Watchdog, the SCC is required. For more information, see the README file at `$deploy/sas-bases/overlays/sas-programming-environment/watchdog/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configuring_sas_compute_server_to_use_sas_watchdog.htm` (for HTML format).

Why the SCC is needed: SAS Watchdog monitors processes to ensure that they comply with the terms of LOCKDOWN system option. It emulates the restrictions imposed by LOCKDOWN by restricting access only to files that exist in folders that are allowed by LOCKDOWN. It therefore requires elevated privileges provided by the custom SCC.

- 1 Apply the SCC with the following command:

```
kubectl apply -f sas-bases/examples/sas-programming-environment/watchdog/sas-watchdog-scc.yaml
```

- 2 Bind the SCC to the service account with the following command:

```
oc -n name-of-namespace adm policy add-scc-to-user sas-watchdog -z sas-programming-environment
```

sas-programming-environment

The SAS Watchdog service account includes the same permissions as the sas-programming-environment service account described here. If you have already bound the SAS Watchdog service account as described above, you should skip this step.

- 1 If your deployment does not require hostPath volume mounts, bind the nonroot SCC to the sas-programming-environment service account using this command:

```
oc -n name-of-namespace adm policy add-scc-to-user nonroot -z sas-programming-environment
```

- 2 If your deployment requires hostPath volume mounts, bind the hostmount-anyuid SCC to the sas-programming-environment service account using this command:

```
oc -n name-of-namespace adm policy add-scc-to-user hostmount-anyuid -z sas-programming-environment
```

Note: The nonroot and hostmount-anyuid SCCs are standard SCCs defined by OpenShift. For more information about them, see [Managing SCCs in OpenShift](#).

Enable hostPath Mounts for CAS

If you want to use hostPath mounts for CAS in your OpenShift deployment, perform the following steps:

- 1 Modify the CAS SCC you chose to use at [“Security Context Constraints and Service Accounts” on page 21](#).
 - a Set `allowHostDirVolumePlugin` to `true`.
 - b Add `hostPath` to the list of volume types. Here is an example:

```
volumes:
- configMap
- downwardAPI
- emptyDir
- persistentVolumeClaim
- projected
- secret
- nfs
- hostPath
```

- 2 Run the following command on the CAS worker node with the directory to be used as a hostPath:

```
sudo chcon -t container_file_t -R directory-path
```

Replace *directory-path* with the directory path to be used as a hostPath mount. This directory must not be located in the file system's root directory (/).

Note: If you prefer, you can use a MachineConfig to apply the same chcon change instead. Consult your cluster administrator to determine which option is best for your deployment.

Networking

OpenShift uses *routes* for its ingress controllers. The initial kustomization.yaml file does not account for routes, so you must change the resources block. Here is the original kustomization.yaml content:

```
resources:
- sas-bases/base
- sas-bases/overlays/network/networking.k8s.io
```

Here is the revision for OpenShift:

```
resources:
- sas-bases/base
- sas-bases/overlays/network/route.openshift.io
```

For details, see the README file at `$deploy/sas-bases/examples/security/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configure_network_security_and_encryption_using_sas_security_certificate_framework.htm` (for HTML format).

TLS

Because OpenShift uses routes instead of Ingress, the security component listed in the initial kustomization.yaml file must be revised. Here is the original kustomization.yaml content:

```
components:
- sas-bases/components/security/core/base/full-stack-tls
```

```
- sas-bases/components/security/network/networking.k8s.io/ingress/
  nginx.ingress.kubernetes.io/full-stack-tls
```

Here is the revision for OpenShift:

```
components:
- sas-bases/components/security/core/base/full-stack-tls
- sas-bases/components/security/network/route.openshift.io/route/full-
  stack-tls
```

For details, see the README file at `$deploy/sas-bases/examples/security/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configure_network_security_and_encryption_using_sas_security_certificate_framework.htm` (for HTML format).

Additional Security

You must perform two additional tasks to apply security for OpenShift:

- remove the Secure Computing (seccomp) profile
- update the fsGroup field

For the instructions to perform both tasks, see the README file at `$deploy/sas-bases/examples/security/container-security/README.md` (for Markdown format) or `$deploy/sas-bases/docs/modify_container_security_settings.htm` (for HTML format).

Directory Structure

\$deploy/sas-bases Directory

The `$deploy/sas-bases` directory contains the files that are provided by SAS to deploy your software. Here is the structure:

```
sas-bases/
├─ base/
├─ components/
├─ docs/
├─ examples/
├─ extras/
└─ overlays/
```

The examples subdirectory contains files that are organized by categories, such as CAS and security, that control the configuration of your deployment. README.md files describe specific configuration settings and the files and modifications that are required to enable those settings. Example files contain the variables that control configuration settings, such as the number of workers in an MPP deployment of CAS. If you decide to enable the configuration setting that is described in a README.md file, you should copy the example file that is listed in the README.md

file to the site-config directory, and replace the variables with the values that you want for your deployment.

The overlays subdirectory also contains files, organized by category, that affect the configuration of your deployment. However, unlike the examples directory, the files in the overlays directory do not need to be modified before they are used. Therefore, they also do not have to be copied to the site-config directory. To use an overlay, you add a reference to the specific file to your base kustomization.yaml file.

The docs subdirectory contains HTML versions of the README files in your software order with an HTML index for them.

Depending on the software in your order, you might have an extras subdirectory. It contains configuration files for special deployment options. It is also subdivided into examples and overlays subdirectories that act in the same way as the higher level examples and overlays directories.

The base and components subdirectories do not contain any files that you should revise.

\$deploy/site-config Directory

Overview

If you are following the instructions in this guide, you have created a directory structure for your SAS Viya platform deployment that looks like this:

```
$deploy/
├─ sas-bases/
└─ site-config/
```

The `$deploy/sas-bases` directory contains the files that SAS provides to deploy your software. During the process of deploying your software, you will be instructed to edit files in the `sas-bases` directory. You should not edit these files directly. Instead, copy the file to the `$deploy/site-config` directory, ensure that the copied file is writeable, and make your modifications to the copied file.

You are being directed to modify copied files because future updates of your software involving replacing the `sas-bases` directory with an updated one. If you store modified files that control the installation and configuration of your software in `sas-bases`, those files are lost during the update. Updates do not change the files in `$deploy/site-config`, although you might be asked to modify a file that exists there.

Organization

Because the `$deploy/site-config` directory contains files that you modify, you should organize the directory according to your needs. However, it is important that you plan the organization before you continue the deployment process. You should organize your files in the `site-config` directory so that you and others who work on the deployment can locate them easily.

Because the structure that you choose for your site-config directory depends on your choices, explicit instructions on where to save specific files cannot be provided. Instead, only general instructions can be offered. An example of a general instruction is “Save the file in your site-config directory.” Suggestions for the organization, however, can be provided.

By Function

You can organize the files according to the functional area of the deployment that they are associated with. When you are working with the files, it should be clear from this guide or a README what aspect of the deployment they relate to. Here is an example:

```
$deploy/
├─ sas-bases/
└─ site-config/
    ├─ authentication/
    └─ cpu/
        └─ storage/
```

By Type

Your deployment is based on a file that you create named kustomization.yaml, which contains five main blocks. When you modify an example file or use an overlay file, you will also be instructed to add a reference to that file in the kustomization.yaml file. Therefore, you can consider organizing the site-config directory according to the blocks in kustomization.yaml. Here is an example:

```
$deploy/
├─ sas-bases/
└─ site-config/
    ├─ configMapGenerator/
    ├─ patches/
    ├─ resources/
    ├─ secretGenerator/
    └─ transformers/
```

Directories for the SAS Viya Platform Deployment Operator

Note: If you are deploying your SAS Viya platform using Kubernetes commands or the sas-orchestration command, skip this section.

Deploying with the SAS Viya Platform Deployment Operator requires a more extensive directory structure that is based on the structure for the \$deploy directory that is described in a preceding section.

```
├─ $license/
|   └─ *-certs.zip
|   └─ *.jwt
```

```

├─ $operator-deploy/
|   ├─ kustomization.yaml
|   ├─ operator-base/
|   ├─ sas-bases/
|   └─ site-config/
|       ├─ cluster-wide-transformer.yaml
|       ├─ remove-seccomp-transformer.yaml
|       └─ transformer.yaml
└─ $deploy/
    ├─ kustomization.yaml
    ├─ sas-bases/
    └─ site-config/

```

The \$license directory contains the .jwt file that licenses your software and the *-certs.zip file that contains the entitlement and CA certificates. If you have more than one order of the SAS Viya platform, there will be multiple license and *-certs.zip files.

The \$operator-deploy directory contents, including the operator kustomization.yaml file, are used only for configuring and deploying the SAS Viya Platform Deployment Operator.

The \$deploy directory is organized as described in a preceding section. Configurations and customizations for your SAS Viya platform deployment are made in the \$deploy/sas-bases and \$deploy/site-config subdirectories. You should have a uniquely named \$deploy directory for each SAS Viya platform deployment.

Note: Do not put additional files in the \$deploy directory unless directed to do so by SAS documentation or SAS Technical Support.

Installation

<i>Initial kustomization.yaml File</i>	32
Overview	32
Create the File	33
<i>Common Customizations</i>	36
Overview	36
Using a Mirror Registry	36
Add the Update Checker to Your Deployment	41
Add a sitedefault File to Your Deployment	42
Configure TLS	43
Add Forward Proxy Settings	44
Configure Container Security Settings	45
Enable Multi-tenancy	45
Configure SAS Image Staging	46
Deploy SAS Startup Sequencer	46
Configure High Availability	47
Configure PostgreSQL	48
Specify SMP or MPP CAS	49
Configure CAS Settings	49
Configure OpenSearch	65
Configure SAS/CONNECT Settings	69
Configure SAS Programming Run-Time Environment	71
Configure Redis	72
Set Default SAS LOCALE and ENCODING in SAS Launcher Service	72
Change the Location of the NFS Server	73
Enable Access Methods Through LOCKDOWN System Option	73
Configure SSSD	74
Specify PersistentVolumeClaims to Use ReadWriteMany StorageClass	75
Configure Open Source Integration Points	76
Configure SAS/ACCESS	76
<i>Install the Orchestration Tool</i>	77
<i>Create the SASDeployment Custom Resource</i>	77
Add an imagePullSecret for the SAS Viya Platform Namespace	77
Run the create sas-deployment-cr Command	78
Revise the Custom Resource for Proxy Information	83
Revise the Custom Resource for Red Hat OpenShift	84

Deploy the Software	84
Deployment Using the SAS Viya Platform Deployment Operator	84
Deployment Using the sas-orchestration Command	87
Deployment Using Kubernetes Commands	90
Save the \$deploy Directory	92
Readiness Service	92
Checks	92
Usage	93
Review the Logs	94
Customize the Readiness Check Period	94
Sign In as the sasboot User	94
Promotion and Migration	95

Initial kustomization.yaml File

Overview

The base kustomization.yaml file (`$deploy/kustomization.yaml`) is where you customize your Kubernetes deployment and allocate resources. The manifest that is used to deploy the software is based on the declarations in the kustomization.yaml file.

Note: Do not rename the kustomization.yaml file. The name of the file must be kustomization.yaml so that it is recognized by the Kustomize tool.

In this section, you are instructed to create a base kustomization.yaml file that deploys SAS with CAS in SMP mode and with an internal instance of PostgreSQL. It is intended to deploy an initial basic SAS Viya platform to help you understand how the base kustomization.yaml file works before adding more customizations. After you have deployed using the initial kustomization.yaml file, you can update it with more customizations for a new deployment or to change your initial deployment, or you can create your own kustomization.yaml file from scratch.

“[Common Customizations](#)” on [page 36](#) contains the most commonly used configurations for deployments. To see the full list of configurations that are available for your order, go to `/$deploy/sas-bases/README.md` for README files in Markdown language or `$deploy/sas-bases/docs/index.htm` for README files in HTML. Failure to review the READMEs in your order and perform the tasks described in them will lead to failures in deployment and usage.

IMPORTANT Before using an existing base kustomization.yaml file for deployments of subsequent versions of the SAS Viya platform, review the deployment notes for changes that could invalidate the continued use of that file for the current version of the software. Although the deployment notes focus on updates, the same issues apply to a base kustomization.yaml file

(and other .yaml files) that was created for the previous version of the software but is being used for the current version of the software.

Note: The README files that are written in the Markdown language can be read as plain text, but they are more useful if they are rendered as Markdown. SAS recommends that you use a tool that renders Markdown when using the README files.

Create the File

- 1 Change to the directory where you saved the deployment assets in “Retrieve Required Files” on page 18:

```
cd /$deploy
```

- 2 Create a kustomization.yaml file. Copy the following code example as the basis of the kustomization.yaml file. For each variable that is enclosed in braces ({}), replace the entire variable (including braces, text, and spaces) with the value that is appropriate for your deployment. For example, {{ NAME-OF-NAMESPACE }} should be replaced with a namespace name, such as **deploy**.

```
namespace: {{ NAME-OF-NAMESPACE }} 1
resources:
- sas-bases/base
- sas-bases/overlays/network/networking.k8s.io 2
- site-config/security/openssl-generated-ingress-certificate.yaml 3
- sas-bases/overlays/cas-server
- sas-bases/overlays/crunchydata/postgres-operator
- sas-bases/overlays/postgres/platform-postgres
# If your deployment contains SAS Viya Programming, comment out the next line
- sas-bases/overlays/internal-elasticsearch
- sas-bases/overlays/update-checker
- sas-bases/overlays/cas-server/auto-resources 4
configurations:
- sas-bases/overlays/required/kustomizeconfig.yaml
transformers:
# If your deployment does not support privileged containers or if your deployment
# contains SAS Viya Programming, comment out the next line
- sas-bases/overlays/internal-elasticsearch/sysctl-transformer.yaml
- sas-bases/overlays/required/transformers.yaml
- sas-bases/overlays/cas-server/auto-resources/remove-resources.yaml 4
# If your deployment contains SAS Viya Programming, comment out the next line
- sas-bases/overlays/internal-elasticsearch/internal-elasticsearch-transformer.yaml
# Mount information
# - site-config/{{ DIRECTORY-PATH }}/cas-add-host-mount.yaml
components:
- sas-bases/components/crunchydata/internal-platform-postgres 5
- sas-bases/components/security/core/base/full-stack-tls 6
- sas-bases/components/security/network/networking.k8s.io/ingress/
  nginx.ingress.kubernetes.io/full-stack-tls 6
patches:
- path: site-config/storageclass.yaml 7
```

```

target:
  kind: PersistentVolumeClaim
  annotationSelector: sas.com/component-name in (sas-backup-job,sas-data-quality-
services,sas-commonfiles,sas-cas-operator,sas-pyconfig)
# License information
# secretGenerator:
# - name: sas-license
#   type: sas.com/license
#   behavior: merge
#   files:
#     - SAS_LICENSE=license.jwt
configMapGenerator:
- name: ingress-input
  behavior: merge
  literals:
    - INGRESS_HOST={{ NAME-OF-INGRESS-HOST }}
- name: sas-shared-config
  behavior: merge
  literals:
    - SAS_SERVICES_URL=https://{{ NAME-OF-INGRESS-HOST }}:{{ PORT }} 8
# - SAS_URL_EXTERNAL_VIYA={{ EXTERNAL-PROXY-URL }} 9

```

- 1 If you are deploying with the SAS Viya Platform Deployment Operator, the name of the namespace is not required.
- 2 If you are deploying on Red Hat OpenShift, replace this line with `- sas-bases/overlays/network/route.openshift.io`.
- 3 This line is appropriate for deployments using openssl to generate the ingress certificate. For information about using cert-manager instead of openssl, see the README file at `$deploy/sas-bases/examples/security/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configure_network_security_and_encryption_using_sas_security_certificate_framework.htm` (for HTML format).
- 4 The nodes in the CAS node pool require taints in order to prevent Kubernetes from scheduling non-CAS workloads on them. If you enable the CAS auto-resources option, your designated CAS nodes must already be labeled. For more information, see [“Plan the Workload Placement” on page 4](#). If the nodes are not labeled, you must disable auto-resources and use hardcoded default resource settings. To disable this option, remove the following line from the resources section:

```
- sas-bases/overlays/cas-server/auto-resources
```

And remove the following line from the transformers section:

```
- sas-bases/overlays/cas-server/auto-resources/remove-resources.yaml
```

- 5 Any lines added to the components block that are not for TLS should be placed before the TLS lines. For example, this line for PostgreSQL comes before the two TLS lines.
- 6 These lines are appropriate for ingress-nginx on a full-stack TLS deployment using openssl as the certificate generator. For the other options available for these lines, see [“Configure TLS” on page 43](#).
- 7 For information about the storageclass.yaml file, see [“Specify PersistentVolumeClaims to Use ReadWriteMany StorageClass” on page 75](#).

- 8 The `SAS_SERVICES_URL` variable is used to provide the value for the `SERVICESBASEURL` option in compute server sessions. It should point to the host name and port of the cluster's Ingress for this deployment. The Ingress must be reachable from pods running inside the cluster. The `{{ NAME-OF-INGRESS-HOST }}` must match the value of the `INGRESS_HOST` variable.

If you are working with external reverse proxies or application gateways, such as Azure Application Gateway, `SAS_SERVICES_URL` and `INGRESS_HOST` can be set to a load balancer. The load balancer can be for the Ingress or external.
- 9 The `SAS_URL_EXTERNAL_VIYA` variable is optional. Use it to define an external proxy that should be used for access from outside your internal network (for example, in generated email messages or text messages). If you choose to use the option, you must uncomment the line by removing the number sign (#).
- 3 Because this kustomization.yaml deploys Full-stack TLS using openssl as the certificate generator, follow the steps at [“Configure TLS” on page 43](#) for certificates for Ingress provided by openssl.
- 4 Because this kustomization.yaml deploys an internal instance of PostgreSQL, follow the steps for internal instances at [“Configure PostgreSQL” on page 48](#) to modify the appropriate .yaml file.
- 5 If you want to make changes to your configuration, see [“Common Customizations” on page 36](#), which is a description of the most frequently used configuration options. To see a list of all the configuration options for your software, go to `/deploy/sas-bases/README.md` for Markdown versions of the README files or `/deploy/sas-bases/docs/index.htm` for HTML versions of the README files.

Note: The README files written in the Markdown language can be read as plain text, but they are more useful if they are rendered as Markdown. SAS recommends that you use a tool that renders Markdown when using the README files.

- 6 To deploy your software:
 - If you are deploying your software with the SAS Viya Platform Deployment Operator, continue to [“Create the SASDeployment Custom Resource” on page 77](#).
 - If you are deploying your software with the `sas-orchestration` command, go to [“Deployment Using the sas-orchestration Command” on page 87](#).
 - If you are deploying your software with Kubernetes commands, go to [“Deployment Using Kubernetes Commands” on page 90](#).

Common Customizations

Overview

Configuring your SAS Viya platform deployment, including setting environment variables, is performed by modifying example and overlay files in your deployment assets and then adding references to those files to the base `kustomization.yaml` file. This section describes the most common deployment-wide configuration decisions that need to be made for the deployment and how to enact them.

IMPORTANT This section contains customizations generic to the SAS Viya platform. Additional product-specific configuration tasks are described in the README files included in your deployment assets. To see the full list of configurations that are available for your order, including the ones described in this section, go to `/deploy/sas-bases/README.md` for README files in Markdown language or `deploy/sas-bases/docs/index.htm` for README files in HTML. Failure to review the READMEs in your order and perform the tasks described in them will lead to failures in deployment and usage.

IMPORTANT Although the steps to configure the deployment are based on Kubernetes tasks and the Kustomize tool, the decisions behind them are based on how you want the SAS Viya platform to run. Successfully configuring and deploying your software to run with the features you want requires conversation between the Kubernetes administrator and the SAS administrator. The Kubernetes administrator for your site should inform the SAS administrator of the options available for the deployment based on the README files in the deployment assets, perhaps even sharing the README files and their descriptions of the options. The Kubernetes administrator should then use the instructions in the README files to enable the options that the SAS administrator wants to use.

For more information about example and overlay files, see “[\\$deploy/sas-bases Directory](#)” on page 27.

Using a Mirror Registry

Note: For introductory information about mirror registries, including how to create them, see “[Create a Mirror Registry](#)” on page 133.

Add a Mirror Registry to Your SAS Viya Platform Deployment

Note: Complete this task only if you are deploying your SAS Viya platform with Kubernetes commands or the `sas-orchestration` command. If you are deploying with the SAS Viya Platform Deployment Operator, you have already completed this task while deploying the operator.

Any Cloud Provider Except Red Hat OpenShift

- 1 Copy the `mirror.yaml` file from `sas-bases/examples/mirror/` to the `site-config` directory.
- 2 Open the `mirror.yaml` file and replace each instance of `{{ MIRROR-HOST }}` with the fully qualified domain name (FQDN) of the mirror registry and the registry namespace in this format:

`FQDN/registry-namespace/sasdeployment`

Here is an example:

`example.company.com/registry-namespace/sasdeployment`

After replacing all instances of `{{ MIRROR-HOST }}`, save and close the file.

- 3 In the base `kustomization.yaml` file, add the path to the `mirror.yaml` immediately to the `transformers` block, immediately after the entry for `sas-bases/overlays/required/transformers.yaml`. Here is an example:

```
transformers:
...
- sas-bases/overlays/required/transformers.yaml
- site-config/mirror.yaml
```

- 4 Add the following content to the `configMapGenerator` block in the base `kustomization.yaml` file:

```
configMapGenerator:
- name: input
  behavior: merge
  literals:
- IMAGE_REGISTRY=registry-host:registry-port/registry-namespace
...
```

In the value for `IMAGE_REGISTRY`:

- `registry-host` is required.
- `registry-port` is optional. If `registry-port` is omitted, the colon (:) that precedes it should be omitted also.
- `registry-namespace` is optional. `registry-namespace` is not the same as the Kubernetes namespace. If `registry-namespace` is omitted, the slash (/) that precedes it should be omitted also.

Using Kubernetes Commands on Red Hat OpenShift

- 1 Copy the `mirror-flattened.yaml` file from `sas-bases/examples/mirror/` to the `site-config` directory.
- 2 Modify the `mirror-flattened.yaml` file.

- a Find the image registry name and port:

```
oc get svc -n name-of-registry-namespace
```

The output looks like this:

Name	Type	Cluster-IP	External-IP	Ports	Age
image-registry	ClusterIP	172.30.146.74	<none>	5000/TCP	174m

- b Open the `mirror-flattened.yaml` file and replace each instance of `{{ MIRROR-HOST }}` with the internal route to the registry from inside the cluster followed by the name of the SAS Viya platform namespace. Use this format:

```
service-name.name-of-registry-namespace.svc:port/platform-namespace
```

The *service-name* and *port* values are in the output from the command in the previous step. The *name-of-registry-namespace* is the same one that you used in the command itself. The *platform-namespace* is the namespace where you will deploy the SAS Viya platform. Here is an example using the sample output from the previous step:

```
image-registry.openshift-image-registry.svc:5000/myviya
```

- c After replacing all instances of `{{ MIRROR-HOST }}`, save and close the file.

- 3 In the base `kustomization.yaml` file, add the path to the `mirror-flattened.yaml` file to the `transformers` block, immediately after the entry for `sas-bases/overlays/required/transformers.yaml`. Here is an example:

```
transformers:
...
- sas-bases/overlays/required/transformers.yaml
- site-config/mirror-flattened.yaml
```

- 4 Add the following content to the `configMapGenerator` block in the base `kustomization.yaml` file:

```
configMapGenerator:
- name: input
  behavior: merge
  literals:
- IMAGE_REGISTRY=service-name.name-of-registry-namespace.svc:port/
platform-namespace
...
```

The value for `IMAGE_REGISTRY=` is the same value that you used for `{{ MIRROR-HOST }}` in step 2 above.

Using the sas-orchestration Tool on Red Hat OpenShift

Add content to the configMapGenerator block in the base kustomization.yaml file.

- 1 Find the image registry name and port:

```
oc get svc -n name-of-registry-namespace
```

The output looks like this:

Name	Type	Cluster-IP	External-IP	Ports	Age
image-registry	ClusterIP	172.30.146.74	<none>	5000/TCP	174m

- 2 Add the following content to the base kustomization.yaml file

```
configMapGenerator:
- name: input
  behavior: merge
  literals:
- IMAGE_REGISTRY=service-name.name-of-registry-namespace.svc:port/
platform-namespace
...
```

The *service-name* and *port* values are in the output from the command in the previous step. The *name-of-registry-namespace* is the same one that you used in the command itself. The *platform-namespace* is the namespace where you will deploy the SAS Viya platform. Here is an example using the sample output above:

```
configMapGenerator:
- name: input
  behavior: merge
  literals:
- IMAGE_REGISTRY=image-registry.openshift-image-registry.svc:5000/
myviya
...
```

Use ImagePullSecrets to Access the Mirror Registry

If the target registry that you deposit the images in requires imagePullSecrets to access them, you must configure an imagePullSecret and add a reference to it in the base kustomization.yaml file.

All Cloud Providers

- 1 Create a representation of an imagePullSecret with kubectl's `--dry-run` option. For the steps to do so, see [Pull an Image from a Private Registry](#).

- 2 Here is the command to obtain the unencoded value of the `.dockerconfigjson` field of `imagePullSecret` required for `{{ MY-IMAGEPULLSECRET }}`. The credentials used in these examples vary by cloud provider and by local cloud provider conventions adopted by your organization.

```
kubectl create secret docker-registry --dry-run=client regcred \
  --docker-server=cloud-provider-specific-registry-location \
  --docker-username=Docker-username \
  --docker-password=Docker-password \
  --docker-email=youremail@example.com \
  --output=jsonpath={.data\.dockerconfigjson}" | base64 --decode
```

Note: Use the following for the *Docker-username*, based on your cloud provider:

- For Microsoft Azure, use `00000000-0000-0000-0000-000000000000`.
 - For Amazon ECR, use `AWS`.
 - For Google Cloud Platform, use `oauth2accesstoken`.
 - For Red Hat OpenShift, use the user name of the user. If you are unsure of the user name, use the command `oc whoami`.
-

You should be familiar with the *cloud-provider-specific-registry-location* because you used it to create your mirror registry. For example, if you are using Microsoft Azure as your cloud provider, *cloud-provider-specific-registry-location* is `myregistry.azurecr.io`.

- 3 Add the following content to the `secretGenerator` block:

```
secretGenerator:
- name: sas-image-pull-secrets
  behavior: replace
  type: kubernetes.io/dockerconfigjson
  literals:
  - .dockerconfigjson={{ MY-IMAGEPULLSECRET }}
```

`{{ MY-IMAGEPULLSECRET }}` should be replaced with the output of the command in step 2.

You can also optionally save `.dockerconfigjson` to a local file and reference it in your `secretGenerator` with the “files” keyword instead of “literals”.

Red Hat OpenShift Alternative

If you are deploying on Red Hat OpenShift, you can use the contents of an existing secret to generate the `imagePullSecret` instead of creating your own secret.

- 1 To get the contents of the `docker-dockercfg` secret, find the name of the secret:

```
kubectl -n name-of-namespace get secret | grep default-dockercfg
```

The output looks like this:

```
default-dockercfg-#####
```

- 2 Create a file that contains the contents of the secret:

```
kubectl -n name-of-namespace get secret output-from-step-1 --
output="jsonpath={.data.\.dockercfg}" | base64 --decode > name-of-
namespace.default.dockercfg.json
```

3 Edit the secretGenerator block of the base kustomization.yaml file:

```
secretGenerator:
- name: sas-image-pull-secrets
  behavior: replace
  type: kubernetes.io/dockercfg
  files:
  - '.dockercfg=name-of-namespace.default.dockercfg.json'
...
```

Add the Update Checker to Your Deployment

Overview

The Update Checker cron job builds a report comparing the currently deployed release with available releases in the upstream repository. The report is written to the stdout of the launched job pod and indicates when new content related to the deployment is available.

If you use the initial kustomization.yaml file from this guide, the Update Checker is deployed by default. If you do not want to deploy the Update Checker, remove the `- sas-bases/overlays/update-checker` line from the resources block of the initial kustomization.yaml file.

For more information about using the Update Checker, see [“View the Update Checker Report” in SAS Viya Platform Operations: Updating Software](#).

(Optional) Define Proxy Environment Variables

The Update Checker cron job is defined in the `sas-bases/overlays/update-checker/cronjobs.yaml` file. If you are using a proxy server for your SAS Viya platform deployment, these environment variables must be defined on the Update Checker cron job:

- HTTP_PROXY
- HTTPS_PROXY
- NO_PROXY

- 1 Copy `$deploy/sas-bases/examples/update-checker/proxy_transformer.yaml` file to `$deploy/site-config/update-checker/proxy_transformer.yaml`.

- 2 In the copied file, replace the value for `HTTP_PROXY` with the host name, the fully qualified host name, or the IP address of the HTTP proxy. Replace the value for `HTTPS_PROXY` with the host name, the fully qualified host name, or the IP address of the HTTPS proxy. Replace the value for `NO_PROXY` with a comma-separated list of host names, fully qualified host names, and IP addresses that the proxy server should ignore.
- 3 Save and close `proxy_transformer.yaml`.
- 4 In the transformers section of the base `kustomization.yaml` file, add a reference to the `proxy_transformer.yaml` file:

```
transformers:
...
- site-config/update-checker/proxy_transformer.yaml
```

Add a sitedefault File to Your Deployment

Note: The extension for the sitedefault file for SAS Viya 4 is different from the file extension used in SAS Viya 3.x.

- SAS Viya 3.x uses `sitedefault.yml`.
- SAS Viya 4 uses `sitedefault.yaml`.

Instructions in this section include references to both sitedefault files, depending on which version is being referred to. Take note of the file extension when it is used.

The `sitedefault.yaml` file, which is located in the `$deploy/sas-bases/examples/configuration` directory, is used primarily for the bulk loading of configuration values. After the initial deployment, the `sitedefault.yaml` file cannot be used to modify an existing value and to deploy the software again. The `sitedefault.yaml` file can be used only to set new property values that have not already been set.

Therefore, SAS recommends that you do **not** use the `sitedefault.yaml` file for the initial deployment of your SAS Viya platform, except where specifically described in this guide. You can, however, use a sitedefault file from a previous SAS Viya release.

To load a SAS Viya 3.x `sitedefault.yml` file in your deployment:

- 1 Ensure that the `sitedefault.yml` file is in the `site-config` directory.
- 2 Add the following content to the `kustomization.yaml` file in the `secretGenerator` block:

```
- name: sas-consul-config
  behavior: merge
  files:
    - SITEDEFAULT_CONF={{ DIRECTORY-PATH }}/sitedefault.yml
```

Configure TLS

Follow these steps to deploy using openssl as the certificate generator.

Note: The example kustomization.yaml file, located at [“Create the File” on page 33](#), represents a deployment with NGINX TLS for full-stack TLS mode using openssl to generate the ingress certificate.

- 1 Copy the `$deploy/sas-bases/examples/security/openssl-generated-ingress-certificate.yaml` file to `$deploy/site-config/security/openssl-generated-ingress-certificate.yaml`.
- 2 In the resources block of the base kustomization.yaml, add a reference to the `openssl-generated-ingress-certificate.yaml` file:

```
resources:
...
- site-config/security/openssl-generated-ingress-certificate.yaml
...
```

- 3 Add the following lines to the components block of the base kustomization.yaml file. Create the components block if it does not already exist.

- a Add a reference to the appropriate subdirectory in `$deploy/sas-bases/components/security/core/base/`:

```
components:
- sas-bases/components/security/core/base/TLS-mode
```

Replace *TLS-mode* with one of the following based on the TLS mode you are deploying:

- `front-door-tls`
- `full-stack-tls`
- `truststores-only`

Note: If you use `truststores-only` as the *TLS-mode*, skip part b of this step.

- b Add a reference to the appropriate subdirectory in `$deploy/sas-bases/components/security/network/`:

```
components:
- sas-bases/components/security/network/ingress-technology/TLS-mode
```

Replace *ingress-technology* with the technology that handles traffic entering your cluster:

- `networking.k8s.io/ingress/nginx.ingress.kubernetes.io`
- `route.openshift.io/route`

Replace *TLS-mode* with one of the following based on the TLS mode you are deploying:

- front-door-tls
- full-stack-tls

Note: For information about the required additions to the base kustomization.yaml for using cert-manager as the certificate generator, see the README file at `$deploy/sas-bases/examples/security/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configure_network_security_and_encryption_using_sas_security_certificate_framework.htm` (for HTML format).

SAS strongly recommends that you safeguard your passwords and data by securing network communication. You can choose not to use TLS, but communication between the pods in your deployment will be unsecured. If you accept that risk or want to conduct experiments using fake data and credentials, you can eliminate network security by deleting the following lines from the example kustomization.yaml:

- From the resources block:
 - - site-config/security/openssl-generated-ingress-certificate.yaml
- From the components block:
 - - sas-bases/components/security/core/base/full-stack-tls
 - - sas-bases/components/security/network/networking.k8s.io/ingress/nginx.ingress.kubernetes.io/full-stack-tls

For more information about TLS and your SAS Viya platform deployment, see the README file at `$deploy/sas-bases/examples/security/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configure_network_security_and_encryption_using_sas_security_certificate_framework.htm` (for HTML format).

Add Forward Proxy Settings

Note: If you are not using a forward proxy while deploying your software, skip this section.

SAS Viya Platform Deployment Operator

Proxy settings, when using the SAS Viya Platform Deployment Operator, can be set for all SAS Viya platform deployments in a cluster or for individual SAS Viya platform deployments within the cluster.

To configure proxy settings for all operator-based Viya deployments in the cluster, environment variables should be added to the deployment operator manifest file. For details, see [“Configure Proxy Information” on page 17](#).

To configure proxy settings for a deployment, add proxy information to the SASDeployment Custom Resource. For more information, see [“Revise the Custom Resource for Proxy Information” on page 83](#). If the settings for the cluster-wide proxy are different than the settings for the deployment proxy, the values for the deployment proxy are used.

sas-orchestration Command and Kubernetes Commands

If you are using a proxy and deploying with the sas-orchestration tool or with Kubernetes commands, ensure that the appropriate environment variables are set appropriately for the Update Checker. For details, see [“\(Optional\) Define Proxy Environment Variables” on page 41](#).

Configure Container Security Settings

Note: If you are deploying on Red Hat OpenShift and have completed the steps in [“Additional Security” on page 27](#), you have performed the necessary steps and should skip this section.

You can change the default container security settings in a SAS Viya platform deployment, such as removing, adding, or updating settings in the podSpecs. SAS has provided example and overlay files to manage the fsGroup field and the secure computing mode (seccomp) profile for your deployment.

- The fsGroup field defines a special supplemental group that assigns a group ID (GID) for all containers in the pod.
- Seccomp is a Linux kernel feature used to restrict actions available within a container.

There are many reasons why an administrator might want to modify these settings. However, if you are deploying on Red Hat OpenShift, they must be modified in order to take advantage of OpenShift's built-in security context constraints. For more information about these settings, see the README file at `$deploy/sas-bases/examples/security/container-security/README.md` (for Markdown format) or `$deploy/sas-bases/docs/modify_container_security_settings.htm` (for HTML format).

Enable Multi-tenancy

Note: Multi-tenancy is not supported on Red Hat OpenShift at this time.

By default, your SAS Viya platform deployment is not multi-tenant. To make your SAS Viya platform deployment multi-tenant, follow the instructions in the README file at `$deploy/sas-bases/examples/multi-tenant/README.md` (for Markdown

format) or at `$deploy/sas-bases/docs/multi-tenant_deployment.htm` (for HTML format).

Note: The decision to enable multi-tenancy must be made before deployment. You cannot change the multi-tenancy status of your deployment after the software has been deployed. The only way to change the status of multi-tenancy in a deployment is to re-deploy the software.

Configure SAS Image Staging

By default, SAS Image Staging starts pods on nodes via a daemonset at approximate two-minute intervals to ensure that relevant images have been pulled to hosts. While this behavior accomplishes the goal of pulling images to nodes and decreasing start-up times, some users may want more intelligent and specific control with less churn in Kubernetes. To accomplish these goals, configure SAS Image Staging to take advantage of a node list to further decrease start-up times and target specific nodes for pulling.

For information about both methods of using SAS Image Staging, including a comparison of their relative advantages and disadvantages, see the README file at `$deploy/sas-bases/examples/sas-prepull/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/sas_image_staging_configuration_option.htm` (for HTML format).

Deploy SAS Startup Sequencer

Although the SAS Viya platform comprises components that are designed to start in any order, in some scenarios it is more efficient for the components to start in an ordered sequence. SAS Startup Sequencer inserts an Init Container into the pods within the Deployments and StatefulSets within the SAS Viya platform. The Init Container ensures that a predetermined, ordered start-up sequence is respected by forcing a pod's start-up to wait until that particular pod can be efficiently started relative to the other pods. The Init Container gracefully exits when it detects the appropriate time to start its accompanying pod, allowing the pod to start. This design ensures that certain components start before others and allows Kubernetes to pull container Images in a priority-based sequence. This design also provides a degree of resource optimization, in that resources are more efficiently spent during the SAS Viya platform start-up with a priority given to starting essential components first.

If you prefer to not use SAS Startup Sequencer, see the README file at `$deploy/sas-bases/overlays/startup/README.md` (for Markdown format) or `$deploy/sas-bases/docs/disabling_the_sas_viya_start-up_sequencer.htm` (for HTML format).

Configure High Availability

The SAS Viya platform can be deployed as a High Availability (HA) system. In this mode, the SAS Viya platform has redundant stateless and stateful services to handle service outages, such as an errant Kubernetes node. A Kustomize transformer enables HA in the SAS Viya platform among the stateless microservices. Stateful services, with the exception of SMP and OpenSearch, are enabled as HA at initial deployment.

To enable HA, add a reference to the `enable-ha-transformer.yaml` file to the base kustomization.yaml file:

```
...
transformers:
...
- sas-bases/overlays/scaling/ha/enable-ha-transformer.yaml
...
```

Note: To enable HA for OpenSearch, see the README file located at `$deploy/sas-bases/examples/configure-elasticsearch/internal/topology/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configure_a_default_topology_for_opensearch.htm` (for HTML format).

For more information about this transformer file and disabling HA, see the README file at `$deploy/sas-bases/examples/scaling/ha/README.md` (for Markdown) or at `$deploy/sas-bases/examples/docs/high_availability_ha_in_the_sas_viya_platform.htm` (for HTML).

Note: The instructions in this section increase the number of replicas of your SAS Viya platform deployment, making it more resilient to pod or node failure and increasing its availability. However, your SAS Viya platform environment probably has dependencies on software that is running in other namespaces in the same cluster. For example, software like ingress-nginx, and the SAS Viya Platform Monitoring for Kubernetes solution might be critical to the availability of the SAS Viya platform and may have been deployed by default with unique replicas, making them less highly available than the SAS Viya platform itself. If you want to increase the availability of other software, consult the documentation for that software for more information.

Furthermore, an over-tainting of the nodes for the SAS Viya platform can result in third-party software being locked out of the cluster in spite of available spare capacity. In order to achieve maximum overall availability, either dedicate some nodes to this software or add tolerations to it so it can more easily run on the same nodes as the SAS Viya platform.

Configure PostgreSQL

Based on your decision about your PostgreSQL instance (see [“Internal versus External PostgreSQL Instances”](#) in *System Requirements for the SAS Viya Platform*), you must perform steps to deploy PostgreSQL (internal) or connect to an existing PostgreSQL instance (external).

Internal Instance of PostgreSQL

If you are using an internal instance of PostgreSQL:

- 1 Go to the base kustomization.yaml file. In the resources block of that file, add the following lines:

```
resources:
- sas-bases/overlays/crunchydata/postgres-operator
- sas-bases/overlays/postgres/platform-postgres
```

- 2 In the components block of the base kustomization.yaml file, add the following content. The new line should be listed before any entries that do not relate to Crunchy Data, such as TLS..

```
components:
- sas-bases/components/crunchydata/internal-platform-postgres
```

The kustomization.yaml file from [“Initial kustomization.yaml File”](#) on page 32 includes these references. For additional information, see the README file located at `$deploy/sas-bases/examples/postgres/README.md` (for Markdown format) or `$deploy/sas-bases/docs/configure_postgresql.htm` (for HTML format).

External Instance of PostgreSQL

For the steps to configure include an external instance of PostgreSQL in your deployment, see the README file located at `$deploy/sas-bases/examples/postgres/README.md` (for Markdown format) or `$deploy/sas-bases/docs/configure_postgresql.htm` (for HTML format).

Configure CDS PostgreSQL

Several offerings on the SAS Viya platform include a second instance of PostgreSQL referred to as CDS PostgreSQL. The CDS PostgreSQL instance is used because the character of the data used by those offerings is hierarchically different than the data generally stored in the platform PostgreSQL database. The separation into two different databases allows them to be tuned individually, in turn enhancing the performance of both.

The list of software offerings that include CDS PostgreSQL is located at “[SAS Common Data Store Requirements](#)” in *System Requirements for the SAS Viya Platform*. If your software order includes at least one of these offerings, CDS PostgreSQL must be configured as well. To configure an instance of CDS PostgreSQL, see the README file located at `$deploy/sas-bases/examples/postgres/README.md` (for Markdown format) or `$deploy/sas-bases/docs/configure_postgresql.htm` (for HTML format).

Specify SMP or MPP CAS

Your deployment of SAS Cloud Analytic Services (CAS) consists of either a single node (SMP) or a set of nodes that include one controller, optionally one backup controller, and multiple workers (MPP). (Although a one-worker MPP configuration is supported, it is not an efficient allocation of resources.) The base `kustomization.yaml` file from “[Initial kustomization.yaml File](#)” on [page 32](#) includes the reference that is required for deploying SMP or MPP CAS. If you do not make any changes to the files in `$deploy/sas-bases/overlays/cas-server`, CAS is deployed as SMP.

To deploy MPP CAS, follow the instructions in the README file at `$deploy/sas-bases/overlays/cas-server/README.md` (for Markdown) or `$deploy/sas-bases/docs/cas_server_for_the_sas_viya_platform.htm` (for HTML) to modify the appropriate files.

Note: Deployments that enable multi-tenancy should use SMP CAS. When additional tenants are onboarded, the decision whether to use SMP or MPP CAS should be made for each tenant.

Configure CAS Settings

Mount hostPaths and Data Connectors for the CAS Server

Data storage in containers is ephemeral. When a container is deleted, the stored data is lost. For durable storage, data should be maintained in persistent volumes outside the Kubernetes cluster. Data remains intact, regardless of whether the containers that the storage is connected to are terminated.

To connect data storage outside the cluster to your SAS Viya deployment:

- 1 Copy `$deploy/sas-bases/examples/cas/configure/cas-add-host-mount.yaml` to your `/site-config` directory.

Note: For more information about the /site-config directory and its structure, see [“\\$deploy/site-config Directory” on page 28](#).

- 2 In the new cas-add-host-mount.yaml file, replace the variables with actual values. Variables are enclosed in braces ({ }) and spaces. When replacing a variable with an actual value, ensure that the braces, spaces, and the hyphenated variable name are removed.
- 3 Save and close the new cas-add-host-mount.yaml file.
- 4 In the base kustomization.yaml file, add the path to your new cas-add-host-mount.yaml file to the transformers block:

```
...
transformers:
- site-config/{{ DIRECTORY-PATH }}/cas-add-host-mount.yaml
...
```

- 5 Save and close the kustomization.yaml file.

Note: You need a version of the cas-add-host-mount.yaml file for each persistent volume that you want to use. Repeat the steps in this section for each persistent volume. Use a different name for the .yaml file each time. Try to use a name that indicates the purpose of the file.

Change accessMode

The default accessMode for the cas-default-data (referred to as “CASDATADIR” in SAS Viya 3.X) and cas-default-permstore persistentVolumeClaims is ReadWriteMany, because it is required for any backup controllers for CAS. It is not required for deployments with SMP CAS, but changing the access mode complicates a possible transition from SMP to MPP in the future.

To change the access mode for either cas-default-data or cas-default-permstore:

- 1 Copy \$deploy/sas-bases/examples/cas/configure/cas-storage-access-modes.yaml to your /site-config directory.

Note: For more information about the /site-config directory and its structure, see [“\\$deploy/site-config Directory” on page 28](#).

- 2 In the new cas-storage-access-modes.yaml file, replace the variables with actual values. Variables are enclosed in braces ({ }) and spaces. To replace a variable with an actual value, ensure that the braces, spaces, and hyphenated variable name are removed.
- 3 Save and close the new cas-storage-access-modes.yaml file.
- 4 In the base kustomization.yaml file, add the path to your new cas-storage-access-modes.yaml file to the transformers block:

```
...
transformers:
```

```
- site-config/{{ DIRECTORY-PATH }}/cas-storage-access-modes.yaml
...
```

- 5 If you are using the initial kustomization.yaml file, go to the patches block. Remove `sas-cas-operator` from the parenthetical list in the `annotationSelector` value.

Note: For more information about the initial kustomization.yaml file, see [“Initial kustomization.yaml File” on page 32](#).

- 6 Save and close the kustomization.yaml file.

Adjust RAM and CPU Resources for CAS Servers

If you use the initial kustomization.yaml file, the CAS operator applies auto-resourcing by default in order to manage the RAM and CPU resources of the nodes where CAS is running. When you instead want to allocate node resources manually, you can disable auto-resourcing and manually modify resourcing requests. For example, you might want to configure guaranteed QoS for CAS server pods. If you manually allocate resources, you must set both the RAM and CPU resources manually.

Note: For auto-resourcing to work appropriately, you must have set labels on your node. See [“Plan the Workload Placement” on page 4](#) for more information.

If you prefer to set your own RAM and CPU resources, perform the following steps.

- 1 Copy `$deploy/sas-bases/examples/cas/configure/cas-manage-cpu-and-memory.yaml` to your `/site-config` directory.

Note: For more information about the `/site-config` directory and its structure, see [“\\$deploy/site-config Directory” on page 28](#).

- 2 In the new `cas-manage-cpu-and-memory.yaml` file, replace the variables with actual values. Variables are enclosed in braces (`{ }`) and spaces. To replace a variable, ensure that the braces, spaces, and hyphenated variable name are removed.
- 3 Save and close the new `cas-manage-cpu-and-memory.yaml` file.
- 4 In the base kustomization.yaml file, remove `- sas-bases/overlays/cas-server/auto-resources` from the resources block. Also remove `- sas-bases/overlays/cas-server/auto-resources/remove-resources.yaml` from the transformers block.
- 5 In the base kustomization.yaml file, add the path to your new `cas-manage-cpu-and-memory.yaml` file to the transformers block:

```
...
transformers:
```

```
- site-config/{{ DIRECTORY-PATH }}/cas-manage-cpu-and-memory.yaml
...
```

- 6 Save and close the kustomization.yaml file.

Change the Number of Workers for MPP CAS

Note: This customization can be performed only for deployments enabling MPP CAS.

By default, MPP CAS has two workers. Perform the following steps to change the number of workers before or after the initial deployment of your SAS Viya platform.

Note: If you want to change the number of workers after the initial deployment of the SAS Viya platform, adding workers and having them join the grid does not require a restart. However, existing SAS sessions will not reallocate or load balance to use the new workers. New sessions should take advantage of the new workers.

Removing workers after the initial deployment requires deleting the CAS deployment, modifying the YAML file, restarting the CAS server, reloading your data, and starting new SAS sessions.

- 1 Copy `$deploy/sas-bases/examples/cas/configure/cas-manage-workers.yaml` to your `/site-config` directory.

Note: For more information about the `/site-config` directory and its structure, see [“\\$deploy/site-config Directory” on page 28](#).

- 2 In the new `cas-manage-workers.yaml` file, replace the variables with actual values. Variables are enclosed in braces (`{ }`) and spaces. To replace a variable with a value, ensure that the braces, spaces, and hyphenated variable name are removed.
- 3 Save and close the new `cas-manage-workers.yaml` file.
- 4 In the base `kustomization.yaml` file, add the path to your new `cas-manage-workers.yaml` file to the transformers block:

```
...
transformers:
- site-config/{{ DIRECTORY-PATH }}/cas-manage-workers.yaml
...
```

- 5 Save and close the kustomization.yaml file.

Add a Backup Controller for MPP CAS

Note: This customization can be performed only for deployments with CAS in MPP mode.

- 1 Copy `$deploy/sas-bases/examples/cas/configure/cas-manage-backup.yaml` to your `/site-config` directory.

Note: For more information about the `/site-config` directory and its structure, see [“\\$deploy/site-config Directory” on page 28](#).

- 2 In the new `cas-manage-backup.yaml` file, replace the variable with the value 0 or 1. The value 0 indicates that you do not want a backup controller, and the value 1 indicates that you want a backup controller.
- 3 Save and close the new `cas-manage-backup.yaml` file.
- 4 In the base `kustomization.yaml` file, add the path to your new `cas-manage-backup.yaml` file to the transformers block:

```
...
transformers:
- site-config/{{ DIRECTORY-PATH }}/cas-manage-backup.yaml
...
```

- 5 Save and close the `kustomization.yaml` file.

Tune CAS_DISK_CACHE

About CAS_DISK_CACHE

The CAS server uses the directory or directories referred to as the CAS Disk Cache as a scratch area. It is associated with the environment variable `CASENV_CAS_DISK_CACHE` and has two primary purposes:

- 1 As data is loaded into memory, it is organized in blocks. Each time a block reaches the default block size of 16Mb, the block is copied to the CAS Disk Cache. The copied block can be re-read back into memory quickly if memory use becomes high and the original data must be freed from memory.
- 2 For a distributed CAS server (MPP), copies of the blocks are transferred to CAS worker pod for fault tolerance. Those copies are also stored in the CAS Disk Cache of the receiving CAS Worker.

A secondary use of the cache is for files that are uploaded to the server. By default, a copy of the file is temporarily stored on the CAS controller in its CAS Disk Cache.

To specify a different location, see [“Storage Location for Uploaded Files”](#) on page 57.

About the Default Configuration

By default, the server is configured to use a directory that is named `/cas/cache` on each controller and worker pod. This directory is provisioned as a Kubernetes `emptyDir` and uses disk space from the root volume of the Kubernetes node.

The default configuration is acceptable for testing and evaluation, but not for production workloads. If disk space in the root volume of the node becomes low, then Kubernetes begins evicting pods. The pod is unlikely to be rescheduled.

When the server stores a block in the cache, the server uses a configure technique that involves opening a file, deleting the file, and then holding the handle to the deleted file. The negative consequence to this technique is that Kubernetes cannot monitor the disk use in the cache.

Choose the Best Storage

The server uses memory mapped I/O for the blocks in the cache. The best performance is provided by using disks that are local to the node for each controller and worker pod. If possible, use disks that provide high data transfer rates such as NVMe or SSD.

If you follow the best practices for workload placement, then no other pods are scheduled on a node that is used by CAS. Even if the root volume is sufficiently large, it is likely that the performance yielded by the root volume will be lower than that of an Ephemeral drive, assuming one is available to the node.

A better strategy is to use a disk that is attached to the node. If the server fills the disk with blocks, the server logs an error rather than Kubernetes evicting the pod. An end user receives the following message when the server runs out of disk space used for the cache on any node.

```
Cloud Analytic Services failed writing to system disk space. Please contact your
administrator.
```

Note: The disk that is used does not need to persist beyond the duration of the pod and does not need to be backed up. Ephemeral storage is ideal.

Use a hostPath for CAS Disk Cache

Most cloud providers offer virtual machines that include a temporary disk for ephemeral storage. Typically, the disk is available at `/dev/sdb1` or a similarly named device. Some cloud providers automatically mount the device on the `/mnt` directory for the VM.

In order to leverage those alternate disks, you can use a Kubernetes `hostPath` instead of an `emptyDir`. The SAS Viya platform deployment requires that those temporary disks are already mounted and available on the CAS nodes and that the path is identical on all nodes.

Single Disk for CAS Disk Cache

- 1 In your `$deploy/site-config/` directory, create a file named `cas_disk_cache-config.yaml`.
- 2 Use the following content in the `cas_disk_cache-config.yaml` file. Replace the variables in the brackets, and the brackets themselves, with values that match your environment.

```
# # this defines the volume and volumemount for CAS DISK CACHE location
---
apiVersion: builtin
kind: PatchTransformer
metadata:
  name: cas-cache-hostpath
patch: |-
  - op: add
    path: /spec/controllerTemplate/spec/volumes/-
    value:
      name: cas-cache-nvme0
      hostPath: # # hostPath, is the path on the host, outside the pod
        path: {{/mnt-nvme0}}
  - op: add
    path: /spec/controllerTemplate/spec/containers/0/volumeMounts/-
    value:
      name: cas-cache-nvme0
      mountPath: /cas/cache-nvme0 # # mountPath is the path inside the
pod that CAS will reference
  - op: add
    path: /spec/controllerTemplate/spec/containers/0/env/-
    value:
      name: CASENV_CAS_DISK_CACHE
      value: "/cas/cache-nvme0" # # This has to match the value that
is inside the pod
target:
  version: v1alpha1
  group: viya.sas.com
  kind: CASDeployment
  # # Target filtering: chose/uncomment one of these option:
  # #   To target only the default CAS server (cas-shared-default) :
  labelSelector: "sas.com/cas-server-default"
  # #   To target only a single CAS server (e.g. MyCAS) other than
default:
  # name: {{MyCAS}}
  # #   To target all CAS Servers
  # name: .*
```

- 3 In the base `kustomization.yaml` file, add the path to your new `cas_disk_cache-config.yaml` file to the transformers block:

```
...
transformers:
...
- site-config/cas_disk_cache-config.yaml
...
```

Microsoft Azure and other cloud providers offer VMs with NVMe storage. Make sure the volume is formatted with an xfs or ext4 file system and is mounted by the VM.

Multiple Disks for CAS Disk Cache

If you use nodes with more than one high-performance disk, you can use more than one disk for the CAS Disk Cache. The server uses a round-robin algorithm for storing blocks on multiple disks.

- 1 In your `$deploy/site-config/` directory, create a file named `cas_disk_cache-config.yaml`.
- 2 Use the following content in the `cas_disk_cache-config.yaml` file. Replace the variables in the brackets, and the brackets themselves, with values that match your environment.

```
# # this defines the volume and volumemount for CAS DISK CACHE location
---
apiVersion: builtin
kind: PatchTransformer
metadata:
  name: cas-cache-hostpath
patch: |-
  - op: add
    path: /spec/controllerTemplate/spec/volumes/-
    value:
      name: cas-cache-nvme0
      hostPath: # # hostPath, is the path on the host, outside the pod
        path: {{/mnt-nvme0}}
  - op: add
    path: /spec/controllerTemplate/spec/volumes/-
    value:
      name: cas-cache-nvme1
      hostPath: # # hostPath, is the path on the host, outside the pod
        path: {{/mnt-nvme1}}
  - op: add
    path: /spec/controllerTemplate/spec/containers/0/volumeMounts/-
    value:
      name: cas-cache-nvme0
      mountPath: /cas/cache-nvme0 # # mountPath is the path inside the
pod that CAS will reference
  - op: add
    path: /spec/controllerTemplate/spec/containers/0/volumeMounts/-
    value:
      name: cas-cache-nvme1
      mountPath: /cas/cache-nvme1 # # mountPath is the path inside the
pod that CAS will reference
  - op: add
    path: /spec/controllerTemplate/spec/containers/0/env/-
    value:
      name: CASENV_CAS_DISK_CACHE
      value: "/cas/cache-nvme0:/cas/cache-nvme1" # # This has to match
the value that is inside the pod
target:
  version: v1alpha1
  group: viya.sas.com
```

```

kind: CASDeployment
# # Target filtering: chose/uncomment one of these option:
# #   To target only the default CAS server (cas-shared-default) :
labelSelector: "sas.com/cas-server-default"
# #   To target only a single CAS server (e.g. MyCAS) other than
default:
# name: {{MyCAS}}
# #   To target all CAS Servers
# name: .*

```

- 3 In the base kustomization.yaml file, add the path to your new cas_disk_cache-config.yaml file to the transformers block:

```

...
transformers:
...
- site-config/cas_disk_cache-config.yaml
...

```

The preceding sample suggests that two NVMe disks are mounted on the node at `/mnt-nvme0` and `/mnt-nvme1`. Steps to perform that action are not shown in this documentation.

Configure Block Size

By default, the server uses a 16 MB block size. If the site accesses very large tables exclusively, you can configure a larger block size to reduce the chance of running out of file handles. Set the `CASCFG_MAXTABLEMEM` environment variable to the preferred value by adding the following block of code to the end of the patch block of your `cas_disk_cache-config.yaml` file.

```

- op: add
  path: /spec/controllerTemplate/spec/containers/0/env/-
  value:
    name: CASCFG_MAXTABLEMEM
    value: {{ BLOCKSIZE }}

```

The value for `{{ BLOCKSIZE }}` should be a numerical value followed by units (K=kilobytes, M=megabytes, or G=gigabytes). The default is 16M.

If a variety of table sizes is used, then individual users can set the `MAXTABLEMEM` session option on a case-by-case basis.

Storage Location for Uploaded Files

An upload is a data transfer of an entire file to the server, such as a SAS data set in SAS7BDAT format or a CSV file. The client, such as SAS, Python, or a web browser, performs no processing on the file. The server performs any processing that is needed, such as parsing records from a CSV file.

```

- op: add
  path: /spec/controllerTemplate/spec/containers/0/env/-
  value:
    name: CASENV_CAS_CONTROLLER_TEMP
    value: {{ MOUNT-PATH-TO-VOLUME }}

```

Ensure that the path you use for `{{ MOUNT-PATH-TO-VOLUME }}` is enclosed by double quotation marks, such as `"/cas/cache-nvme0"`.

Configure External Access to CAS

Overview of CAS Connectivity

By default, a single CAS server is configured during the deployment process and is accessible to SAS services and web applications that are deployed in the Kubernetes cluster. For example, SAS Visual Analytics, SAS Studio, and other SAS software can work with CAS and do not require any additional configuration.

In addition, an HTTP Ingress is enabled that provides access to CAS from outside the cluster to clients that use REST. This Ingress can be used with clients such as Python SWAT.

The Ingress Controller that is configured for your cluster enables connectivity to CAS at an HTTP path like `https://www.example.com/cas-shared-default-http/`.

Note: Use the path as shown for clients such as Python SWAT. For curl, use a path such as `/cas-shared-default-http/cas`. This document shows the path that is appropriate for Python SWAT.

Note: The default instance of the CAS server is referenced in this example and the rest of this topic. If you add more than one server, then the Ingress or Service name uses the server instance name instead of the word “default”.

Optional Connectivity

There are two uses of CAS that require additional configuration:

- **Connections from SAS 9.4, SAS Viya 3.5, or other binary clients.** If you want to connect to CAS from SAS Viya 3.5, SAS 9.4, or use a binary connection with open programming clients such as Python, R, and Java, you can enable a binary connection.
- **Connections to CAS from SAS Data Connectors.** For information about enabling connectivity for SAS/ACCESS and data connectors, see the README file `$deploy/sas-bases/examples/data-access/README.md` (for Markdown) or `$deploy/sas-bases/docs/configuring_sasaccess_and_data_connectors_for_sas_viya_4.htm` (for HTML).

About Binary Connectivity

Most clients can use a binary connection to the CAS server. Typically, performance is better than HTTP because the data stream is more compact than REST.

If you want to connect from SAS Viya 3.5 or SAS 9.4, then you must enable binary communication. You can use the node port or load balancer as described here or

you can configure a custom Ingress to proxy TCP port 5570. Configuring a custom Ingress is not described in this documentation.

Optional Binary and HTTP Services

You can enable two services that provide external access to CAS for programmers. One service provides binary communication and the other service provides HTTP communication for REST. The HTTP service is an alternative to using the HTTP Ingress that is enabled by default.

The binary communication provides better performance and can be used by SAS Viya 3.5 or SAS 9.4. Open source clients such as Python SWAT require C language libraries to use the binary connection. Refer to the documentation for the open source client for information about the libraries.

If you enable either of these services, they are enabled as NodePorts by default. To use the services as LoadBalancers, you must specify LoadBalancer as the type. You can also restrict traffic by setting ranges of IP addresses for the load balancers to accept traffic on.

Note: The CAS operator supports setting the binary and HTTP services to either NodePort or LoadBalancer. Setting a combination of service types is not supported by the operator. In addition, the DC and EPCS services that are part of SAS/ACCESS and Data Connectors are also affected.

Configuration

- 1 Copy the `$deploy/sas-bases/examples/cas/configure/cas-enable-external-services.yaml` to your `$deploy/site-config` directory.
- 2 In the copied file, set the `publishBinaryService` key to true to enable binary communication for clients from outside the Kubernetes cluster:


```
- op: replace
  path: /spec/publishBinaryService
  value: true
```
- 3 If you want to enable the HTTP service, set the `publishHTTPService` key to true. This enables a service for REST access from outside the Kubernetes cluster. Be aware that REST access is enabled by default through a Kubernetes Ingress. If you have access through the Ingress, then enabling this HTTP service is redundant.


```
- op: replace
  path: /spec/publishHTTPService
  value: true
```
- 4 The services are configured as NodePort by default. For deployments in Microsoft Azure or Amazon Web Services (AWS), NodePort is not supported and you must configure the services as LoadBalancer services.

To configure them as LoadBalancer services, uncomment the `serviceTemplate`. Setting source ranges is optional. Delete the lines if you do not want them. Here is an example:

```
- op: add
  path: /spec/serviceTemplate
```

```

value:
  spec:
    type: LoadBalancer
    loadBalancerSourceRanges:
      - 192.168.0.0/16
      - 10.0.0.0/8

```

Note: SAS supports setting the `type` and `loadBalancerSourceRanges` keys in the service specification. Adding any other key such as `port` or `selector` can result in poor performance or prevent connectivity.

- 5 Set the `publishExtHostnameSuffix` key if you set the service to use `LoadBalancer`, your deployment is in Microsoft Azure or in AWS, and you meet either of these conditions:

- if you are using the DC or EPCS service.
- if the deployment is configured to use TLS.

When you set the key, the CAS Operator adds a subject alternative name (SAN) for each service to the certificate that is created by `sas-certframe`. The operator also adds a DNS label annotation to the service.

```

- op: add
  path: /spec/publishExtHostnameSuffix
  value: "-unique-name.subdomain-name"

```

For Microsoft Azure, replace *subdomain-name* with your Azure region name, such as `"eastus2.cloudapp.azure.com"`. The text in the value, up to the first period, is appended to the service name to create a unique DNS name. For example, the default value for the binary service is `sas-cas-server-default-bin`. If `-orion.eastus2.cloudapp.azure.com` is specified, then the operator creates the following annotation and publishes a DNS record for it.

```

apiVersion: v1
kind: Service
metadata:
  annotations:
    service.beta.kubernetes.io/azure-dns-label-name: sas-cas-server-default-bin-orion
...

```

For the example, the DNS record is `sas-cas-server-default-bin-orion.eastus2.cloudapp.azure.com`.

For AWS, replace *subdomain-name* with your subdomain of choice, such as `"viya.acme.com"`. For example the value you supply for `publishExtHostnameSuffix` could be `"-pisces.viya.acme.com"`. The `service.beta.kubernetes.io/azure-dns-label-name` annotation will be added to the deployment, but will be ignored by AWS. No DNS record will be generated. The administrator must create a DNS alias/CNAME record for each external service, including each node of the SAS Data Connect Accelerators, after deployment. See [“Configure External Access to Amazon Web Services CAS Services” on page 106](#) for details.

- 6 In the base `kustomization.yaml` file, add the path to your new `cas-enable-external-services.yaml` file to the `transformers` block:

```

transformers:
...
- site-config/{{ DIRECTORY-PATH }}/cas-enable-external-services.yaml

```


7 Save and close the kustomization.yaml file.

Note: If you configure direct access to the CAS server via HTTP or binary and are using full-stack TLS, the subject alternative names (SAN) in the certificate generated by cert-manager must include the host name or IP address being used to access that service. For the steps to include the host name or IP address, see [“Add the External Host Name or IP Address to the SAN in the Signed Certificate” in SAS Viya Platform Encryption: Data in Motion](#).

If you are making these changes after the initial deployment, the binary and HTTP services do not require that you restart CAS. For other services related to CAS, refer to the documentation to determine if a restart is required.

Use one of the next two sections to identify the connection information that programmers need to connect to CAS from outside the Kubernetes cluster.

Connection Information for Programmers: NodePort

You can use the following commands to identify the network port that maps to the service.

```
kubectl -n name-of-namespace get svc sas-cas-server-default-http
kubectl -n name-of-namespace get svc sas-cas-server-default-bin
```

For a NodePort, find the network port that programmers connect to.

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
sas-cas-server-default-bin	NodePort	10.0.5.236	<none>	5570:31066/TCP

Programmers need to know the host name of one of the Kubernetes nodes. You can use the following command to list the node names.

```
kubectl -n name-of-namespace get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
host02398.example.com	Ready	<none>	24d	v1.18.4
host02483.example.com	Ready	<none>	24d	v1.18.4
host02656.example.com	Ready	master	24d	v1.18.4
host02795.example.com	Ready	<none>	24d	v1.18.4
host02854.example.com	Ready	<none>	24d	v1.18.4

To connect from SAS Viya 3.5 or SAS 9.4 to the NodePort, run a CAS statement like the following example:

```
options CASHOST="host02398.example.com" CASPORT=31066;
cas casauto;
```

For the sas-cas-server-default-http service, a REST client connects to one of the Kubernetes nodes, such as host02398.example.com, and the port that is mapped to 8777.

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
sas-cas-server-default-http	NodePort	10.107.219.118	<none>	8777:31535/TCP

A REST client can connect to a resource such as `https://host02398.example.com:31535/cas-shared-default-http/`.

Connection Information for Programmers: LoadBalancer

You can use the following commands to identify the external IP address for the load balancer.

```
kubectl -n name-of-namespace get svc sas-cas-server-default-http
kubectl -n name-of-namespace get svc sas-cas-server-default-bin
```

The output includes the IP address of the load balancer and programmers connect to native port, 5570.

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
sas-cas-server-default-bin	LoadBalancer	10.0.44.57	52.247.0.1	5570:32215/TCP

To connect from SAS Viya 3.5 or SAS 9.4 to the LoadBalancer, run a CAS statement like the following example:

```
options CASHOST="sas-cas-server-default-bin-
orion.eastus2.cloudapp.azure.com";
options CASPORT=5570;

cas casauto;
```

Substitute your deployment-specific information for the sample unique value, orion, and the sample region, eastus2.

For the `sas-cas-server-default-http` service, a REST client connects to the load balancer on port 80 for HTTP or port 443 for HTTPS if TLS is configured. Only one of the two ports is operational, depending on whether TLS is configured.

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
sas-cas-server-default-http	LoadBalancer	10.0.61.68	52.247.0.168	8777:31979/ TCP,80:30707/TCP,443:30032/TCP

A REST client can connect to a resource such as `https://sas-cas-server-default-http-orion.eastus2.cloudapp.azure.com:443/cas-shared-default-http/`.

SAS Data Connect Accelerators

The SAS Data Connect Accelerators enable parallel data transfer between a distributed CAS server (MPP) and some data sources such as Teradata and Hadoop. For information about enabling connectivity for SAS/ACCESS and Data Connectors, see the README file at `$deploy/sas-bases/examples/data-access/README.md` (for Markdown) or `$deploy/sas-bases/docs/configuring_sasaccess_and_data_connectors_for_sas_viya_4.htm` (for HTML).

More Documentation

For more information about the services, see [“Kubernetes Services for CAS” in SAS Viya Platform Operations: Servers and Services](#).

Enable Host Launch

By default, CAS cannot launch sessions under a user's host identity. All sessions run under the cas service account instead. CAS can be configured to allow for host identity launches by including a patch transformer in the kustomization.yaml file. To enable host launch for CAS, see the “Enable Host Launch in the CAS Server” section of the README file located at `$deploy/sas-bases/examples/cas/configure/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configuration_settings_for_cas.htm` (for HTML format).

Enable State Transfer for CAS Servers

Note: If you are not using the SAS Viya Platform Deployment Operator or sas-orchestration deploy to manage your deployment, skip this section.

Enabling state transfers preserves the sessions, tables, and state of a running CAS server for a new CAS server instance that is being started as part of a CAS server update. To enable the state transfer of CAS servers, see the README at `$deploy/sas-bases/overlays/cas-server/state-transfer/README.md` (for Markdown format) or `$deploy/sas-bases/docs/state_transfer_for_cas_server_for_the_sas_viya_platform.htm` (for HTML format).

Note: You cannot enable state transfer and CAS auto-restart in the same SAS Viya platform deployment. If you want to enable state transfer for a deployment that already has CAS auto-restart enabled, you must first disable CAS auto-restart before enabling state transfer.

Enable CAS Auto-Restart After Updates

Note: If you are not using the SAS Viya Platform Deployment Operator or sas-orchestration deploy to manage your deployment, skip this section.

By default, CAS does not automatically restart during version updates performed by the SAS Viya Platform Deployment Operator or sas-orchestration deploy. To change the default to enable auto-restart, see the “CAS Auto-Restart During Version Updates” section of the README file located at `$deploy/sas-bases/overlays/`

`cas-server/README.md` (for Markdown format) or `$deploy/sas-bases/docs/cas_server_for_the_sas_viya_platform.htm` (for HTML format).

Note: You cannot enable CAS auto-restart and state transfer in the same SAS Viya platform deployment.

Configure GPUs for CAS

The SAS GPU Reservation Service aids SAS processes in resource sharing and utilization of the Graphic Processing Units (GPUs) that are available in a Kubernetes Pod. It is required in every SAS Cloud Analytic Services (CAS) Pod that is GPU-enabled. For information about implementing the SAS GPU Reservation Service, see the README located at `$deploy/sas-bases/examples/gpu/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/sas_gpu_reservation_service.htm` for HTML format.

Note: This README only describes how to use the SAS GPU Reservation Service for CAS. To learn about how other products use GPUs, consult the READMEs for each product.

Create a Personal CAS Server

For development purposes in applications such as SAS Studio, you might need to allow data scientists the ability to work with a CAS server that is local to their SAS Compute session. This personal CAS server is just like a regular (shared) CAS server except it is simpler, relatively short-lived, and is only for one person.

To set up a personal CAS server, see the README file at `$deploy/sas-bases/overlays/sas-programming-environment/personal-cas-server/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configuring_sas_compute_server_to_use_a_personal_cas_server.htm` (for HTML format).

To set up a personal CAS server that uses a GPU, see the README file at `$deploy/sas-bases/overlays/sas-programming-environment/personal-cas-server-with-gpu/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configuring_sas_compute_server_to_use_a_personal_cas_server_with_gpu.htm` (for HTML format).

To configure a personal CAS server with or without a GPU, see the README file at `$deploy/sas-bases/examples/sas-programming-environment/personal-cas-server/README.md` (for Markdown format) or at `configuration_settings_for_the_personal_cas_server.htm` (for HTML format).

Configure OpenSearch

Note: The SAS Viya Programming offering does not include OpenSearch. If your SAS Viya platform order contains SAS Viya Programming, skip this section.

Based on your decision about your OpenSearch instance (see [“Internal versus External OpenSearch Instances”](#) in *System Requirements for the SAS Viya Platform*), you must perform steps to deploy OpenSearch (internal) or connect to an existing OpenSearch instance (external).

Internal Instances of OpenSearch

Initial Customizations

OpenSearch is an Apache 2.0-licensed distribution with enterprise security. The SAS Viya platform includes OpenSearch and uses its distributed search cluster in infrastructure and solution services. Some additions to the base `kustomization.yaml` file must be made to configure OpenSearch.

Note: The example `kustomization.yaml` file, located at [“Initial kustomization.yaml File”](#) on page 32, includes these customizations.

- 1 Add the following line to the resources block of the base `kustomization.yaml` file:

```
resources:
...
- sas-bases/overlays/internal-elasticsearch
...
```

- 2 Add the following line to the transformers block of the base `kustomization.yaml` file:

```
transformers:
...
- sas-bases/overlays/internal-elasticsearch/internal-elasticsearch-transformer.yaml
...
```

Configure Default Virtual Memory Resources

Note: If you are deploying on Red Hat OpenShift and have completed the steps in [“Security Context Constraints and Service Accounts”](#) on page 21, you have performed the necessary steps and should skip this section.

The OpenSearch pods require additional virtual memory resources. In order to provide these memory resources, a transformer uses a privileged container to set the virtual memory for the `mmapfs` directory to the required level. Therefore, privileged containers must be permitted by your Pod security policies. For more information about Pod security policies, see <https://kubernetes.io/docs/concepts/policy/pod-security-policy/>.

You have three options:

- If privileged containers are enabled, add a reference to the `sysctl-transformer.yaml` file to the transformers block of the base `kustomization.yaml` file. This transformer must be included after any TLS transformers and before the `sas-bases/overlays/required/transformers.yaml` transformer.

Note: The `sysctl-transformer.yaml` transformer uses a privileged container to set `vm.max_map_count`. If privileged containers are not allowed in your deployment, do not add this line.

Here is an example:

```
transformers:
...
- sas-bases/overlays/network/ingress/security/transformers/...
- sas-bases/overlays/internal-elasticsearch/sysctl-transformer.yaml
- sas-bases/overlays/required/transformers.yaml
...
```

Note: Using this option requires modifying the OpenShift SCC for `sas-opensdistro` to allow it. For more information, see the README file at `$deploy/sas-bases/examples/configure-elasticsearch/internal/openshift/README.md` (for Markdown format) or `$deploy/sas-bases/docs/openshift_on_red_hat_openshift.htm` (for HTML format).

- If privileged containers are not allowed in your environment, a Kubernetes administrator with elevated permissions can set the virtual memory manually before performing the SAS Viya platform deployment. All nodes that run workloads in a class that is tolerated by the [stateful workload class](#) are affected by this requirement.

To configure the virtual memory settings for `mmapfs` manually:

- 1 Log on to the first stateful node as root or with a `sudoers` account.
- 2 Set the virtual memory using the appropriate method:
 - To set the value permanently, use your preferred text editor to modify `/etc/sysctl.conf` or the equivalent in your environment. Update the `vm.max_map_count` setting to `262144` and save the file.
 - To set the value temporarily, run the following command:

```
sysctl -w vm.max_map_count=262144
```

- 3 (Optional) Verify the modified setting:

```
sysctl vm.max_map_count
```

- 4 Repeat the previous steps on each node that is labeled for stateful workloads.

If you are using a managed Kubernetes cluster, your cloud provider probably provisions the nodes dynamically. In this instance, be aware that manual modifications do not persist after a restart of a Kubernetes node. The cluster administrator must use an alternative method to save the `vm.max_map_count` setting.

- You can disable the use of `mmap` at a cost of performance and memory usage. To disable `mmap`, include a reference to the `disable-mmap-transformer.yaml` overlay in the transformers block of the base `kustomization.yaml` file.

```
transformers:
...
- sas-bases/overlays/internal-elasticsearch/disable-mmap-transformer.yaml
```

Configure a StorageClass

Deploying OpenSearch requires a StorageClass that provides block storage (such as virtual disks) or a local file system mount to store the search indices. For the instructions to configure such a StorageClass for all cloud providers, see the “Configure a Default StorageClass for OpenSearch” README, located at `$deploy/sas-bases/examples/configure-elasticsearch/internal/storage/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configure_a_default_storageclass_for_opensearch.htm` (for HTML format).

Configure High Availability

To enable HA for OpenSearch, see the README file located at `$deploy/sas-bases/examples/configure-elasticsearch/internal/topology/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configure_a_default_topology_for_opensearch.htm` (for HTML format).

Configure a Run User

A fixed user ID (UID) is required so that files that are written to storage for the search indices can be read after subsequent restarts. In a default deployment of the SAS Viya platform, the OpenSearch JVM process runs under the fixed UID of 1000. However, on some environments, using a UID of 1000 can lead to a conflict between users within the container and those on the host. At the initial deployment, you can select a new run user for the OpenSearch pods.

Note: This task can only be performed at the initial deployment.

To configure a UID other than 1000, see the README file located at `$deploy/sas-bases/examples/configure-elasticsearch/internal/run-user/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configure_a_run_user_for_opensearch.htm` (for HTML format).

External Instances of OpenSearch

For the steps to configure an external instance of OpenSearch in your deployment, see the README file located at `$deploy/sas-bases/examples/configure-elasticsearch/external/README.md` (for Markdown format) or `$deploy/sas-bases/docs/configure_an_external_opensearch_instance.htm` (for HTML format).

Additional Configuration for FIPS Compliance

Starting with 2023.06, the SAS Viya platform supports deployments in a FIPS-compliant environment. However, neither the internal nor the external instance of OpenSearch supports FIPS at this time. In order to enable OpenSearch to start and run in a FIPS-enabled environment, you must apply a transformer to your deployment manifest.

- 1 Create a transformer named `opendistro-disable-fips-transformer.yaml` in the directory `$deploy/site-config` with the following contents:

```
---
apiVersion: builtin
kind: PatchTransformer
metadata:
  name: sas-opendistro-disable-fips-transformer
patch: |-
  - op: add
    path: /spec/config/jvm/-
    value: -Dcom.redhat.fips=false
target:
  kind: OpenDistroCluster
  name: sas-opendistro
```

Note: The formatting for this file is important. Be sure to copy and paste the content exactly as it appears here.

- 2 Add the `opendistro-disable-fips-transformer.yaml` file to the transformers block of the base `kustomization.yaml` file. Here is an example:

```
transformers:
...
- site-config/opendistro-disable-fips-transformer.yaml
```


Configure SAS/CONNECT Settings

Support External Sign-on

To enable NodePort or LoadBalancer, see the README file at `$deploy/sas-bases/examples/sas-connect-spawner/README.md` (for Markdown language) or `$deploy/sas-bases/docs/configure_sasconnect_spawner_in_the_sas_viya_platform.htm` (for HTML).

Note: In managed environments like Microsoft Azure, you cannot access the NodePort service from a client outside of the cluster.

Spawn SAS/CONNECT Servers Within the Spawner Pod

By default, SAS/CONNECT servers cannot be spawned within the spawner pod. Instead they are spawned in the SAS/CONNECT server pods. However, SAS clients at 9.4 M6 or older and 9.4 M7 clients that do not have the hot fix linked to [SAS Note 68611](#) cannot reach launched SAS/CONNECT server pods. Those clients must enable spawning SAS/CONNECT servers within the spawner pods by applying the security settings in the `enable-spawned-servers.yaml` example file. For details, see the "Allow the Ability to Spawn Servers within the Spawner Pod" section of the "Configure SAS/CONNECT Spawner in SAS Viya" README file located at `$deploy/sas-bases/examples/sas-connect-spawner/README.md` (for Markdown format) and at `$deploy/sas-bases/docs/configure_sasconnect_spawner_in_the_sas_viya_platform.htm` (for HTML format).

Connection Information for Programmers: NodePort

To sign on when a NodePort is specified:

- 1 Get the NodePort value that is mapped to the service port.

```
kubectl describe service/sas-connect-spawner-nodeport
```

or

```
kubectl get service/sas-connect-spawner-nodeport -o yaml
```

```
- name: service
  nodePort: 24133 // port that is exposed externally
  port: 17551
  protocol: TCP
```

- 2 Determine the host name of one of the Kubernetes nodes. If you are using no TLS or using TLS with self-signed certificates with the nodes in the DNS list or using a wildcard to match the nodes, you can use the following command to list the node names.

```
kubectl -n name-of-namespace get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
host02398.example.com	Ready	<none>	24d	v1.18.4
host02483.example.com	Ready	<none>	24d	v1.18.4
host02656.example.com	Ready	master	24d	v1.18.4
host02795.example.com	Ready	<none>	24d	v1.18.4
host02854.example.com	Ready	<none>	24d	v1.18.4

If you are using TLS with a cert-manager (such as sas-viya-issuer), sas-certframe adds the node name that the pod is running on to the certificate for the service. Use this command to find the log entry describing the addition of the node name to the certificate:

Note: Because of the length of the command and the margin of the page, this command appears as more than one line. The command should be entered as a single line.

```
kubectl -n name-of-namespace logs deployment/sas-connect-spawner sas-
certframe | grep KUBE_NODE_NAME
```

In the output, the node name is listed and can be provided to programmers. Here is an example:

```
2020-09-03 23:11:45 - [INFO] - Adding KUBE_NODE_NAME host02398.example.com to
SAS_CERTIFICATE_SAN_DNS
```

- 3 Sign on from an external client machine.

```
%let rem=node-name-from-step-2 nodeport-from-step-1;
signon rem user='user-ID' password='password';
```

Using the examples from step 1 and 2, the command would look like this:

```
%let rem=host02398.example.com 24133;
signon rem user='myuserid' password='mypassword';
```

Connection Information for Programmers: LoadBalancer

- 1 Determine the DNS name for the IP address that was provided by the load balancer. If you have not already registered a DNS name, you should do so now.

The requirement to register a DNS name is described at “[Kubernetes Cluster Requirements](#)” in *System Requirements for the SAS Viya Platform*.

Note: For information about DNS names while using Azure, see [Apply a DNS label to the service](#).

- 2 Sign on from an external client machine.

```
%let rem=DNS-name-from-step-1 17551;
signon rem user='user-ID' password='password';
```

Configure SAS Programming Run-Time Environment

External Storage Class for SAS Programming Run-Time Environment

All SAS Viya Platform Servers

The Batch Server, Compute Server, and SAS/CONNECT Server are SAS Viya platform servers that use the SAS Programming Run-time Environment. They create a number of temporary files for run-time information in a location that is local to the `sas-programming-environment` pod. By default, these pods are backed by an `emptyDir` volume named `viya`, which is mounted automatically. However, using the default `emptyDir` volume is not recommended because SAS programming components can consume large amounts of storage quickly and cause nodes to shut down.

To configure different storage classes for the `viya` volume, see the README file at `$deploy/sas-bases/examples/sas-programming-environment/storage/README.md` (for Markdown format) or `$deploy/sas-bases/docs/sas_programming_environment_storage_tasks.htm` (for HTML format).

Batch Server Only

If you want the Batch Server to have storage that is different than the Compute Server and the SAS/CONNECT Server, such as using persistent storage rather than ephemeral storage, see the README file at `$deploy/sas-bases/examples/sas-batch-server/storage/README.md` (for Markdown format) or `$deploy/sas-bases/docs/sas_batch_server_storage_task_for_checkpoint_restart.htm` (for HTML format).

GPUs for SAS Programming Run-Time Environment

For large amounts of data, some procedures for SAS Programming Run-Time Environment run faster on a Graphic Processing Unit (GPU) than on a CPU with multiple threads. The SAS GPU Reservation Service aids SAS processes in resource sharing and utilization of the GPUs that are available in a Kubernetes pod. The SAS Programming Environment container image makes this service available, but it must be enabled in order to take advantage of the GPUs in your cluster. To enable the SAS GPU Reservation Service for the SAS Programming Run-time Environment, see the README file at `$deploy/sas-bases/overlays/sas-programming-environment/gpu/README.md` (for Markdown format) or `$deploy/sas-bases/docs/sas_gpu_reservation_service_for_sas_programming_environment.htm` (for HTML format).

Note: This README only describes how to use the SAS GPU Reservation Service for the SAS Programming Run-Time Environment. To learn about how other products use GPUs, consult the READMEs for each product.

Configure Redis

The SAS Viya platform uses Redis to provide a distributed cache technology in its deployments. For information about configuring Redis, see the README file at `$deploy/sas-bases/examples/redis/server/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configuration_settings_for_redis.htm` (for HTML format).

Set Default SAS LOCALE and ENCODING in SAS Launcher Service

Setting the default locale and encoding for the SAS Launcher Service controls the default SAS LOCALE and ENCODING for SAS Compute Server, SAS/CONNECT, and SAS Batch Server, unless overridden by another specification. In order to set or modify these settings, see the “Locale and Encoding Defaults” section of the README file at `$deploy/sas-bases/examples/sas-launcher/configure/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configuration_settings_for_sas_launcher_service.htm` (for HTML format).

Change the Location of the NFS Server

SAS provides a transformer that allows you to change the location of the NFS server hosting the user's home directories. For information about using the transformer, see the “NFS Server Location” section of the README file located at `$deploy/sas-bases/examples/sas-launcher/configure/README.md` (for Markdown format) or `$deploy/sas-bases/docs/configuration_settings_for_sas_launcher_service.htm` (for HTML format).

Enable Access Methods Through LOCKDOWN System Option

The SAS Viya platform uses the LOCKDOWN system option to limit access to files and features. By default, the following methods cannot be used to access files and specific SAS features for a SAS session that is executing in batch mode or server processing mode:

- EMAIL
- FTP
- HADOOP
- HTTP
- PYTHON
- PYTHON_EMBED
- SOCKET
- TCPIP
- URL

To enable any of these access methods, see the README file at `$deploy/sas-bases/examples/sas-programming-environment/lockdown/README.md` (for Markdown format) or `$deploy/sas-bases/docs/lockdown_settings_for_the_sas_programming_environment.htm` (for HTML format). For more information about the LOCKDOWN system option, see “[LOCKDOWN System Option](#)” in *SAS Viya Platform: Programming Run-Time Servers*.

Configure SSSD

Enable SSSD

The configuration of SSSD is, by default, performed automatically for you. The automatic process uses the configuration service LDAP settings (if they exist) to construct an `sssd.conf` file. However, until that configuration is enabled, SSSD will not be available to your SAS Viya platform deployment. To enable that SSSD configuration:

- 1 Add a reference to the `sas-bases/overlays/cas-server/cas-sssd-sidecar.yaml` file to the transformers block in the base `kustomization.yaml` file. The new line must precede any lines for TLS transformers and the line for required transformers. Here is an example:

```
transformers:
- sas-bases/overlays/cas-server/cas-sssd-sidecar.yaml
...
- sas-bases/overlays/required/transformers.yaml
...
```

- 2 Follow the steps described in the "Disable Cloud Native Mode" in the "Configuration Settings for CAS" README file. The README is located at `$deploy/sas-bases/examples/cas/configure/README.md` (for Markdown format) and `$deploy/sas-bases/docs/configuration_settings_for_cas.htm` (for HTML format).
- 3 Because SSSD requires host authentication, follow the steps described at ["Enable Host Launch" on page 63](#).
- 4 Save and close the `kustomization.yaml` file.

Add a Custom Configuration for SSSD

If you would prefer to use a custom configuration for SSSD instead of the default, after completing the steps in ["Enable SSSD" on page 74](#), perform the following steps:

- 1 Copy the `$deploy/sas-bases/examples/cas/configure/cas-sssd-example.yaml` file to the location of your CAS server overlay, such as `$deploy/site-config/cas-server/cas-sssd-example.yaml`.
- 2 Add the location of the copied `cas-sssd-example.yaml` to the transformers block of the base `kustomization.yaml` file. The new line should go after the required transformers line. Here is an example based on the example used in step 1:

```
transformers:
...
- sas-bases/overlays/required/transformers.yaml
```

```
- site-config/cas-server/cas-sssd-example.yaml
...
```

- 3 Create your sssd.conf file and add your custom SSSD configuration to it. SAS recommends putting the sssd.conf file in the `$deploy/site-config` directory.

- 4 Add the following code to the secretGenerator block of the base kustomization.yaml file using the path to the sssd.conf file you created in step 3. Here is an example using `$deploy/site-config/cas-server/sss.conf` as that path:

```
secretGenerator:
...
- name: sas-sssd-config
  files:
    - SSSD_CONF=site-config/cas-server/sss.conf
  type: Opaque
...
```

- 5 Save and close the kustomization.yaml file.

Specify PersistentVolumeClaims to Use ReadWriteMany StorageClass

The manifest file that the base kustomization.yaml creates must have information about which PVCs in your deployment should take advantage of the StorageClass you created for your cloud provider.

- 1 In the `$deploy/site-config` directory, create a file named `storageclass.yaml`. Use the following content in that file.

```
kind: RWXStorageClass
metadata:
  name: wildcard
spec:
  storageClassName: {{ RWX-STORAGE-CLASS }}
```

Replace `{{ RWX-STORAGE-CLASS }}` with the name of your cluster's StorageClass that provides ReadWriteMany (RWX) access.

- 2 In the base kustomization.yaml file, add a patches block with the following content.

Note: The `annotationSelector` line in the following code is too long for the width of the page. A line break has been added to address the issue. If you copy this code for use, be sure to remove the line break.

```
patches:
- path: site-config/storageclass.yaml
  target:
    kind: PersistentVolumeClaim
    annotationSelector: sas.com/component-name in (sas-backup-job,sas-
data-quality-services,sas-commonfiles,sas-cas-operator,sas-pyconfig)
```

Note: If you are using the example `kustomization.yaml` file included at “[Initial kustomization.yaml File](#)” on page 32, the `patches` block is already present.

Depending on the software that you are deploying, you might have to add more content to the `annotationSelector` line.

- If your order contains SAS Model Risk Management, add `sas-risk-cirrus-search` to the parenthetical list in the `annotationSelector` value.
 - If your order includes SAS Risk Modeling, add `sas-risk-modeling-core` to the parenthetical list in the `annotationSelector` value.
-

Note: If you have changed the `accessMode` for CAS per the instructions at “[Change accessMode](#)” on page 50, remove `sas-cas-operator` from the parenthetical list in the `annotationSelector` value.

Configure Open Source Integration Points

The SAS Viya platform integrates with open source programming languages such as Python and R in both directions, from the SAS Viya platform to open source and back. With this integration, you can call out to open-source engines from within the SAS Viya platform interfaces to leverage code that was previously written in other environments. You can also write open-source code to access powerful SAS Analytics from your coding interfaces of choice, including Jupyter Notebooks and R-Studio. You can use Python or R directly with SAS or integrate SAS into applications using REST APIs to process operations more efficiently on a multithreaded, in-memory, massively parallel processing engine.

For a high-level list of the steps to install, configure, and deploy Python and R to enable integration in the SAS Viya platform, see the README located at `$deploy/sas-bases/examples/sas-open-source-config/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configure_python_and_r_integration_with_sas_viya.htm` (for HTML format).

Configure SAS/ACCESS

To configure and deploy your SAS/ACCESS products, see the “Configuring SAS/ACCESS and Data Connectors for Viya 4” README file at `$deploy/sas-bases/examples/data-access/README.md` (for Markdown) or `$deploy/sas-bases/docs/configuring_sasaccess_and_data_connectors_for_sas_viya_4.htm` (for HTML).

Install the Orchestration Tool

The `sas-orchestration` tool is required when using `sas-orchestration` or SAS Viya Platform Deployment Operator to deploy. If you have not already deployed the orchestration tool, do it now. Follow the instructions in the “Prerequisites” section of the README file at `$deploy/sas-bases/examples/kubernetes-tools/README.md` (for Markdown format) or `$deploy/sas-bases/docs/using_kubernetes_tools_from_the_sas-orchestration_image.htm` (for HTML format).

Create the SASDeployment Custom Resource

Note: This section is required only if you are using the SAS Viya Platform Deployment Operator.

Add an imagePullSecret for the SAS Viya Platform Namespace

If your SAS Viya platform content has been mirrored and the mirror requires authentication, you must create an `imagePullSecret` for the namespace in which you are deploying the SAS Viya platform. The `imagePullSecret` must be named `sas-orchestration-secret`. For more information about the command to create `imagePullSecrets`, see [Pull an Image from a Private Registry](#).

All Cloud Providers

Use the following command to add the `sas-orchestration-secret`.

```
kubectl -n name-of-namespace \
  create secret generic sas-orchestration-secret \
  --type=kubernetes.io/dockerconfigjson \
  --from-file=.dockerconfigjson=file-with-secret-content
```

For example, if you are deploying the SAS Viya platform into a namespace called `viya`, and the secret content is in a file named `site-config/image-pull-secret.json`, the command would look like this:

```
kubectl -n viya \
  create secret generic sas-orchestration-secret \
  --type=kubernetes.io/dockerconfigjson \
  --from-file=.dockerconfigjson=site-config/image-pull-secret.json
```

Red Hat OpenShift Alternative

If you are deploying on Red Hat Openshift, you can create the secret from the existing secret.

1 Find the name of the secret:

```
kubectl -n name-of-namespace get secret | grep default-dockercfg
```

The output looks like this:

```
default-dockercfg-#####
```

2 Create a file that contains the contents of the secret:

```
kubectl -n name-of-namespace get secret output-from-step-1 --
output="jsonpath={.data.\.dockercfg}" | base64 --decode > name-of-
namespace.default.dockercfg.json
```

3 Create the secret named from this file:

```
kubectl -n name-of-namespace \
  create secret generic sas-orchestration-secret \
  --type=kubernetes.io/dockercfg \
  --from-file=.dockercfg=name-of-namespace.name-of-secret.json
```

The *name-of-secret* is a name you choose. It should have meaning to you or your organization and help to identify the secret.

Run the create sas-deployment-cr Command

Note: The container to create the custom resource runs under the sas ID and group. Ensure that the \$(pwd) directory, specified in the create command, has permissions that can accommodate the sas ID and group.

As an administrator with local cluster permissions, run the following command to create the SASDeployment custom resource and to name it \$deploy-sasdeployment.yaml. Make sure that the command is run from the parent directory of the \$license and \$deploy directories. Here is the command format:

```
docker run --rm \
  -v $(pwd):mount-for-working-directory-inside-container \
  sas-orchestration \
  create sas-deployment-cr \
  --deployment-data certificates-information \
  --license license-information \
  --user-content location-of-deployment-files \
  --cadence-name stable-or-lts \
```

```
--cadence-version cadence-version-number \
[--cadence-release cadence-release-number \]
[--image-registry mirror-registry-location \]
[--repository-warehouse repository-warehouse-location \]
> $deploy-sasdeployment.yaml
```

Here is a description of the values to be substituted for the variables in the command:

Note: For information about all the flags available for the `create sas-deployment-cr` command, use the help flag:

```
docker run --rm \
sas-orchestration \
create sas-deployment-cr \
--help
```

mount-for-working-directory-inside-container

The path at which the current working directory should be mounted inside the container.

certificates-information

The location of the `*-certs.zip` file. It can be a directory path, which includes the mount for the working directory, or a go-getter URL.

license-information

The location of the license, which can be a directory path, including the mount for the working directory, or a go-getter URL.

location-of-deployment-files

The location of the `$deploy` directory. This can be a directory path, including the mount for the working directory and the `$deploy` directory name, or a go-getter URL.

stable-or-lts

Use `stable` for software in the Stable cadence, or use `lts` for the Long-Term Support cadence.

cadence-version-number

The cadence version number of the software to be deployed (for example, 2020.1.4).

[cadence-release-number] (optional)

The latest cadence release or a specific cadence release of the cadence version number of the software to be deployed. See the important note that follows.

Note: Because the orchestration tool generates an internal `sas-bases` folder based on the information in this command, you can ensure that the data is consistent by reviewing the `$deploy/sas-bases/.orchestration/cadence.yaml` file. Ensure that the `cadence-type`, `cadence-version-number`, and `cadence-release-number` flags in the command match the `name`, `version`, and `release` fields, respectively, of the `cadence.yaml` file.

[mirror-registry-location] (optional)

The URL for the docker image registry (for example, `registry.example.com`). This flag is required if you are deploying with a mirror registry.

Note: If you are deploying on Red Hat OpenShift, the URL must be in the following format: *service-name.name-of-registry-namespace.svc:port/platform-namespace*. For example, *image-registry.openshift-image-registry.svc:5000/myviya*. Use the same value you used to replace `{{ MIRROR-HOST }}` in the *mirror-flattened.yaml* file (see step 2 of “Configure the Mirror Registry” on page 14).

[*repository-warehouse-location*] (optional)

The URL for the warehouse describing what should be deployed. This flag is needed if you are managing a dark environment.

\$deploy

Precede the name of the *sasdeployment.yaml* file with the name of the directory that the software is being deployed from. For example, if you use *viya1* as *\$deploy*, the file should be named *viya1-sasdeployment.yaml*.

Note: The files being pulled into the custom resource must be text files. If you must use a binary file, it should be added to the custom resource by using a go-getter URL. For information about go-getter URLs, see <https://github.com/hashicorp/go-getter>.

IMPORTANT If you specify the cadence release, its specification affects how the operator reconciles the custom resource. Consider the following when deciding how to use the `--cadence-release` option.

- If you are using a mirror registry, the *cadence-type*, *cadence-version-number*, and *cadence-release-number* must match the `--deployment-assets`, `--cadence`, and `--release` flags used to populate that mirror.
- If you are using a mirror registry and did not use the `--deployment-assets`, `--cadence`, or `--release` flags to populate that mirror, then the *cadence-type* and *cadence-version-number* must be the latest available at the time the mirror was created. The *cadence-release-number* must be either an empty string (" ") or the latest release available at the time the mirror was created.
- Set the value to " " in order to force the operator to use the latest release for the requested *cadence-version-number*. A consequence of using this value is that processes that seem unrelated to software updates, such as renewing your license or making configuration changes, might also spawn updates to the running software.
- Set the value to a specific cadence release value to use that cadence release in the custom resource. Specifying the cadence release of the currently running deployment can be used to make configuration changes without introducing updates. If the specified cadence release does not exist, the user is presented with an error.
- If no cadence release is specified, as in the examples that follow, the operator uses the latest cadence release when the custom resource is initially introduced to a namespace. The operator automatically assigns the chosen release value in the cluster representation of this custom resource. If the user later applies an updated custom resource, without a cadence release specified, Kubernetes preserves the previously assigned value of this field and therefore the versions of any software already

deployed into that namespace. To change the version of software in an existing namespace, assign a cadence release as described above.

Here is an example of the command with the following values.

- The directory should be mounted in `/cwd/` in the container.
- The `*-certs.zip` file is located at `/cwd/license/SASViyaV4_69SWC4_certs.zip`.
- The license file from SAS is located at `/cwd/license/SASViyaV4_69SWC4_lts_2021_license_2020-09-08T105930.jwt`.
- The `$deploy` directory is `/cwd/viya1`.
- The software being deployed is Long-Term Support 2021.1.

```
docker run --rm \
  -v $(pwd):/cwd/ \
  sas-orchestration \
  create sas-deployment-cr \
  --deployment-data /cwd/license/SASViyaV4_69SWC4_certs.zip \
  --license /cwd/license/SASViyaV4_69SWC4_lts_2021_license_2020-09-08T105930.jwt \
  --user-content /cwd/viya1 \
  --cadence-name lts \
  --cadence-version 2021.1 \
> viya1-sasdeployment.yaml
```

Here is an excerpt of the generated custom resource:

```
...
---
apiVersion: orchestration.sas.com/v1alpha1
kind: SASDeployment
metadata:
  name: sas-viya
spec:
  cadenceName: lts
  cadenceVersion: "2021.1"
  license:
    secretKeyRef:
      name: sas-viya
      key: license
  clientCertificate:
    secretKeyRef:
      name: sas-viya
      key: cert
  caCertificate:
    secretKeyRef:
      name: sas-viya
      key: cacert
  userContent:
    files:
      kustomization.yaml: |
        resources:
          - sas-bases/base
          - sas-bases/overlays/cert-manager-issuer
...
```

Notice that the values that were entered in the command are included in the custom resource. The custom resource also includes a transcription of the contents of the base customization.yaml file.

Here is an example of the command that uses references (in the form of go-getter URLs) to the locations for the values:

- The directory should be mounted in `/cwd/dep1oy` in the container.
- The `*-certs.zip` file is located at `https://example.com/SASViyaV4_69SWC4_certs.zip`.
- The license file from SAS is located at `https://example.com/SASViyaV4_69SWC4_lts_2021_license_2020-09-08T105930.jwt`.
- The `$deploy` directory is `git::https://user:token@git.example.com/repository.git//viya1`.
- The software that is being deployed is Long-Term Support 2021.1.

Note: When fetching from a Git repository, in order for the content to be cloned locally by the operator before being used, you must use the annotation `environment.orchestration.sas.com/readOnlyRootFilesystem: "false"`.

```
docker run --rm \
  -v $(pwd):/cwd/dep1oy \
  sas-orchestration \
  create sas-deployment-cr \
  --deployment-data https://example.com/SASViyaV4_69SWC4_certs.zip \
  --license https://example.com/SASViyaV4_69SWC4_lts_2021_license_2020-09-08T105930.jwt \
  --user-content git::https://user:token@git.example.com/repository.git//viya1 \
  --cadence-name lts \
  --cadence-version 2020.1 \
> viya1-sasdeployment.yaml
```

The generated custom resource would include this content:

```
...
---
apiVersion: orchestration.sas.com/v1alpha1
kind: SASDeployment
metadata:
  annotations:
    environment.orchestration.sas.com/readOnlyRootFilesystem: "false"
  creationTimestamp: null
  name: sas-viya
spec:
  caCertificate:
    url: https://example.com/SAS_CA_Certificate.pem
  cadenceName: lts
  cadenceVersion: "2020.1"
  clientCertificate:
    url: https://example.com/entitlement_certificate.pem
  license:
    url: https://example.com/SASViyaV4_69SWC4_lts_2021_license_2020-09-08T105930.jwt
  repositoryWarehouse: {}
  userContent:
    url: git::https://user:token@git.example.com/repository.git//viya1
status:
```

...

Notice that the custom resource contains the information that is included in the command.

Note: For more information about the fields in the SASDeployment custom resource, see [“Fields in the SASDeployment Custom Resource” on page 150](#).

Revise the Custom Resource for Proxy Information

Note: If you are not using a forward proxy for your cluster, skip this section. If you would like to use a proxy for all SAS Viya platform deployments in a cluster, you must modify the deployment operator manifest file. For details, see [“Configure Proxy Information” on page 17](#).

To define the proxy for a single SAS Viya platform deployment, add the following lines to the `metadata/annotations` block of the custom resource:

```
environment.orchestration.sas.com/HTTP_PROXY: proxy-URL-for-HTTP-requests
environment.orchestration.sas.com/HTTPS_PROXY: proxy-URL-for-HTTPS-requests
```

Additionally, you can add a line that defines which requests should not go through the proxy:

```
environment.orchestration.sas.com/NO_PROXY: do-not-proxy-list
```

The *do-not-proxy-list* is a comma-separated list of host names, fully qualified host names, and IP addresses that the proxy should ignore.

Here is an example of a custom resource that includes proxy information:

...

```
apiVersion: orchestration.sas.com/v1alpha1
kind: SASDeployment
metadata:
  annotations:
    environment.orchestration.sas.com/readOnlyRootFilesystem: "false"
    environment.orchestration.sas.com/HTTP_PROXY: http://webproxy.example.com:5000
    environment.orchestration.sas.com/HTTPS_PROXY: http://webproxy.example.com:5000
    environment.orchestration.sas.com/NO_PROXY: localhost,noproxy.example.com,
kubernetes.default.svc,10.96.0.1
  creationTimestamp: null
  name: sas-viya
...
```

Note: The following values must be included in the list of values for the `NO_PROXY` variable:

- `kubernetes.default.svc` (the Kubernetes API server)
- the value of the `KUBERNETES_SERVICE_HOST` environment variable for the cluster

Revise the Custom Resource for Red Hat OpenShift

Note: If your deployment is not running on Red Hat OpenShift, you should skip this section.

If your deployment is running on Red Hat OpenShift, you must add an annotation to the SAS Deployment custom resource. In the `metadata/annotations` block of the custom resource, add the following line:

```
environment.orchestration.sas.com/FLATTENED_IMAGE_REGISTRY: "true"
```

Here is an example:

```
...
---
apiVersion: orchestration.sas.com/v1alpha1
kind: SASDeployment
metadata:
  annotations:
    environment.orchestration.sas.com/readOnlyRootFilesystem: "false"
    environment.orchestration.sas.com/FLATTENED_IMAGE_REGISTRY: "true"
  creationTimestamp: null
  name: sas-viya
...
```

Deploy the Software

Deployment Using the SAS Viya Platform Deployment Operator

Command and Output

Because the operator is actually running as a result of the last command that you performed in [“Apply the SAS Viya Platform Deployment Operator Resources to the Cluster” on page 18](#), the operator responds to any changes to the SASDeployment custom resource by applying those changes. Therefore, to perform the initial deployment, run the following command as an administrator with local cluster permissions to apply the SASDeployment custom resource:

```
kubectl -n name-of-namespace apply -f $deploy-sasdeployment.yaml
```


Note: Because `$deploy/sas-bases` is restricted from modification, the operator generates a `sas-bases` folder based on the cadence information you supplied. This folder plus user-supplied content (such as the base `kustomization.yaml` file) is used for deploying or updating your software.

To determine the status of the deployment, run the following command:

```
kubectl -n name-of-namespace get sasdeployment
```

Here is an example of the output:

NAME	STATE	CADENCENAME	CADENCEVERSION	CADENCERELEASE	AGE
viya1 130m	SUCCEEDED	stable	2020.1.3	20210304.1614817334881	

The `STATE` field cycles through several values. The field value starts with `PENDING`, then `RECONCILING`, and finishes in either `SUCCEEDED` or `FAILED`. For more information about communications from the SAS Viya Platform Deployment Operator, see [“Communications from the Operator” on page 156](#).

Note: SAS recommends that you save a copy of the `SASDeployment` custom resource locally or to Git as a backup.

Initial Troubleshooting

When the SAS Viya Platform Deployment Operator is not working as expected, three different sources can be used to diagnose problems. If you need to contact SAS Technical Support for help, be sure to share the output from all three of these sources.

Note: After the deployment, the log from the SAS Viya Platform Deployment Operator Reconcile Job might contain the following message:

```
Warning: 'vars' is deprecated. Please use 'replacements' instead. [EXPERIMENTAL]
Run 'kustomize edit fix' to update your Kustomization automatically.
```

If this message is displayed, it can safely be ignored.

Log from the SAS Viya Platform Deployment Operator Pod

The log from the SAS Viya Platform Deployment Operator pod can be useful in diagnosing problems that might be preventing the SAS Viya Platform Deployment Operator from deploying the SAS Viya platform. By default, that pod is named `sas-deployment-operator-hash`. The Kustomize tool appends the hash value during the deployment of the SAS Viya Platform Deployment Operator. An example pod name is `sas-deployment-operator-57f567f7bc-drg5z`.

Use the following command to generate log output:

```
kubectl \
```

```
logs \
-n name-of-deployment-operator-namespace \
deployment-operator-pod-name
```

SASDeployment Custom Resource

The `.status` field of a SASDeployment custom resource contains information about the last attempt to deploy the SAS Viya platform. For complete details about this field, see [“Communications from the Operator” on page 156](#). Specifically, the `.status.messages` field contains all the messages from the last Reconcile Job that was started by the SAS Viya Platform Deployment Operator. These messages relate to fetching URLs, running Kustomize, and running kubectl.

Use the following command to generate output for the entire SASDeployment custom resource:

```
kubectl \
get sasdeployments \
-n name-of-SAS-Viya-namespace \
-o yaml
```

Log from the Reconcile Job

The log from the SAS Viya Platform Deployment Operator Reconcile Job can be useful in diagnosing problems with deploying a particular SASDeployment custom resource. By default, that Job is named `sas-deployment-operator-reconcile-hash`. A unique Job is associated with each deployment attempt. All these Jobs are located in the same namespace as the SASDeployment custom resource that they are deploying, providing a historical record of those attempts. The Jobs are removed automatically after the associated SASDeployment custom resource is removed from Kubernetes.

Depending on cluster settings, the pod that is run by the Job might not be available after the process exits. However, if the pod remains, use the following command to generate output for its Job log:

```
kubectl \
logs \
-n name-of-SAS-Viya-namespace \
reconcile-Job-pod-name
```

Remediation

If an issue prevents the successful deployment of your software and one of the sources described above indicates the issue is associated with content in `$deploy/site-config` or the base `kustomization.yaml` file, take the following steps to address the issue before contacting SAS Technical Support:

- 1 Make corrections for the error. Debugging can include reviewing example files for formatting, file names, or path specifications. The base `kustomization.yaml` file can also be reviewed to ensure it was revised as necessary. To help with debugging, refer to the appropriate documentation, including README files.
- 2 Rebuild the SASDeployment custom resource using the instructions at [“Run the create sas-deployment-cr Command” on page 78](#).

- 3 Apply the custom resource using the instructions at [“Command and Output” on page 84](#).

Deployment Using the sas-orchestration Command

Install the Orchestration Tool

Before you can issue the command to deploy your software, you must first install the orchestration tool. Follow the instructions at [“Install the Orchestration Tool” on page 77](#).

Command

After the orchestration tool is installed, run the following command to deploy your software:

```
docker run --rm \
  -v $(pwd):mount-for-working-directory-inside-container \
  -v "mount-for-kubeconfig-file-location-inside-container" \
  -e "KUBECONFIG=assignment-of-kubeconfig-file-within-container" \
  [-e FLATTENED_IMAGE_REGISTRY=true \]
sas-orchestration \
deploy \
--namespace name-of-namespace \
--deployment-data certificates-information \
--license license-information \
--user-content location-of-deployment-files \
--cadence-name stable-or-lts \
--cadence-version cadence-version-number \
[--cadence-release cadence-release-number \]
[--image-registry mirror-registry-location \]
[--repository-warehouse repository-warehouse-location \]
```

Note: The `-e FLATTENED_IMAGE_REGISTRY=true` option should only be used if you are deploying from an image registry on Red Hat OpenShift.

Here is a description of the values to be substituted for the variables in the command:

Note: Because `$deploy/sas-bases` is restricted from modification, the orchestration tool generates a `sas-bases` folder based on the cadence information you supplied. This folder plus user-supplied content (such as the base `kustomization.yaml` file) is used for deploying or updating your software.

Note: For information about all the flags available for the `deploy` command, use the help flag:

```
docker run --rm \
  sas-orchestration \
  deploy \
  --help
```

mount-for-working-directory-inside-container

The path at which the current working directory should be mounted inside the container.

mount-for-kubeconfig-file-location-inside-container

The mounted location of the cluster's configuration file.

assignment-of-kubeconfig-file-within-container

The `KUBECONFIG` environment variable pointing to the location of the kubeconfig file within the container.

name-of-namespace

The namespace where the software is to be deployed.

certificates-information

The location of the `*-certs.zip` file. It can be a directory path, which includes the mount for the working directory, or a go-getter URL.

license-information

The location of the license, which can be a directory path, including the mount for the working directory, or a go-getter URL.

location-of-deployment-files

The location of the `$deploy` directory. This can be a directory path, including the mount for the working directory, or a go-getter URL.

stable-or-lts

Use `stable` for software in the Stable cadence, or use `lts` for the Long-Term Support cadence.

cadence-version-number

The cadence version number of the software to be deployed (for example, 2020.1.4).

[cadence-release-number] (optional)

The latest cadence release or a specific cadence release of the cadence version number of the software to be deployed.

Note: Because the orchestration tool generates an internal `sas-bases` folder based on the information in this command, you can ensure that the data is consistent by reviewing the `$deploy/sas-bases/.orchestration/cadence.yaml` file. Ensure that the `cadence-type`, `cadence-version-number`, and `cadence-release-number` flags in the command match the `name`, `version`, and `release` fields, respectively, of the `cadence.yaml` file.

[mirror-registry-location] (optional)

The URL for the docker image registry (for example, `registry.example.com`). This flag is needed if you are deploying with a mirror registry.

Note: If you are deploying on Red Hat OpenShift, the URL must be in the following format: *service-name.name-of-registry-namespace.svc:port/platform-namespace* (for example, *image-registry.openshift-image-registry.svc:5000/myviya*). Use the same value you used in the `configMapGenerator` block of the base `kustomization.yaml` file (see step 2 of “Using the `sas-orchestration` Tool on Red Hat OpenShift” on page 39).

[*repository-warehouse-location*] (optional)

The URL for the warehouse describing what should be deployed. This flag is needed if you are managing a dark environment.

Note: The files being pulled used to deploy the software must be text files. If you must use a binary file, it should be added to the custom resource by using a `go-getter` URL. For information about `go-getter` URLs, see <https://github.com/hashicorp/go-getter>.

Example

Here is an example of the `sas-orchestration deploy` command that includes the following values.

- The directory should be mounted in `/cwd/` in the container.
- The software is being deployed in the `viya1` namespace.
- The `*-certs.zip` file is located at `/cwd/SASViyaV4_69SWC4_certs.zip`.
- The license file from SAS is located at `/cwd/SASViyaV4_69SWC4_stable_2022.12_license_2022-12-08T105930.jwt`.
- The `$deploy` directory is `/cwd/deploy`.
- The software being deployed is `Stable 2022.12`.

```
docker run --rm \
-v $(pwd) :/cwd/ \
-v "/home/user/.kube/config:/kube/config" \
-e "KUBECONFIG=/kube/config" \
sas-orchestration \
deploy \
--namespace viya1 \
--deployment-data /cwd/SASViyaV4_69SWC4_certs.zip \
--license /cwd/SASViyaV4_69SWC4_stable_2022.12_license_2022-12-08T105930.jwt \
--user-content /cwd/deploy \
--cadence-name stable \
--cadence-version 2022.12
```

IMPORTANT Provider-specific code has been removed from the open source Kubernetes code base in Kubernetes 1.26. With Google Kubernetes Engine (GKE) 1.26, Google-specific artifacts are required for deployment. The `sas-orchestration deploy` command can only be used with GKE 1.26 by mounting the provider-specific artifacts to the `sas-orchestration` container.

Here is the command format for Google cloud (the provider-specific artifacts are highlighted):

```
docker run --rm \
  -v $(pwd):mount-for-working-directory-inside-container \
  --user $(id -u):$(id -g) \
  -v location-of-the-GKE-gcloud-auth-plugin:location-in-the-container \
  -v location-of-authentication-files-used-by-Google-CLI:
location-in-the-container \
  -v "mount-for-kubeconfig-file-location-inside-container" \
  -e "KUBECONFIG=assignment-of-kubeconfig-file-within-container" \
  -e "PATH=append-location-of-the-GKE-gcloud-auth-plugin-to-existing-path" \
  sas-orchestration \
  deploy \...
```

Here is an example of the command:

```
docker run --rm \
  -v $(pwd):/cwd/ \
  --user $(id -u):$(id -g) \
  -v /install/google-cloud-sdk:/usr/lib64/google-cloud-sdk \
  -v "$HOME"/.config/gcloud:/.config/gcloud \
  -v "/home/user/.kube/config:/kube/config" \
  -e "KUBECONFIG=/kube/config" \
  -e "PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:
/bin:/opt/sas/viya/home/bin:/usr/lib64/google-cloud-sdk/bin" \
  sas-orchestration \
  deploy \...
```

Deployment Using Kubernetes Commands

Note: If you have deployed the SAS Viya Platform Deployment Operator, these commands are not necessary since the operator deploys your software for you. For more information, see `$deploy/sas-bases/examples/deployment-operator/deploy/README.md` (for Markdown) or `$deploy/sas-bases/docs/sas_viya_deployment_operator.htm` (for HTML).

IMPORTANT The following `kubectl` commands require that the `kubeconfig` environment variable is set. For information about setting that variable, see [The KUBECONFIG environment variable](#). Alternatively, you can add the `--kubeconfig=namespace-kubeconfig-file` argument to each `kubectl` command for the command to work properly.

- 1 On the `kubectl` machine, create the Kubernetes manifest:

```
kustomize build -o site.yaml
```

The following message might be displayed:

Warning: 'vars' is deprecated. Please use 'replacements' instead. [EXPERIMENTAL]
Run 'kustomize edit fix' to update your Kustomization automatically.

If the message is displayed, it can safely be ignored.

2 Apply cluster-api resources to the cluster. As an administrator with cluster permissions, run

```
kubect1 apply --selector="sas.com/admin=cluster-api" --server-side --force-conflicts -f site.yaml
```

```
kubect1 wait --for condition=established --timeout=60s -l "sas.com/admin=cluster-api" crd
```

The kubect1 apply command might cause the following messages to be displayed:

error: no objects passed to apply

resource mapping not found for name: "foo" namespace: "<name-of-namespace>" from
"site.yaml": no matches for kind "bar" in version "baz"
ensure CRDs are installed first

If either message is displayed, it can safely be ignored.

3 As an administrator with cluster permissions, run

```
kubect1 apply --selector="sas.com/admin=cluster-wide" -f site.yaml
```

4 As an administrator with local cluster permissions, run

```
kubect1 apply --selector="sas.com/admin=cluster-local" -f site.yaml --prune
```

5 As an administrator with namespace permissions, run

```
kubect1 apply --selector="sas.com/admin=namespace" -f site.yaml --prune
```

The kubect1 apply command might cause the following message to be displayed:

error: error pruning nonNamespaced object

If this message is displayed, it can safely be ignored.

6 If you are performing an update, as an administrator with namespace permissions, run the following command to prune additional resources not in the default set.

```
kubect1 apply --selector="sas.com/admin=namespace" -f site.yaml --prune --prune-whitelist=autoscaling/v2/HorizontalPodAutoscaler
```

7 Wait for Kubernetes to create and start the pods. To determine whether the pods have started:

```
kubect1 -n name-of-namespace get pods
```

The output of this command looks like this:

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
d10006	annotations-66dc4479fd-qfqqr	1/1	Running	0	5s
d10006	appregistry-bbbdfb78c-tc1lv	1/1	Running	0	5s
d10006	audit-7c4ff4b8b8-zxg8k	1/1	Running	0	5s
d10006	authorization-79d4f594b9-t9sbx	1/1	Running	0	5s
d10006	cachelocator-668fcd544-hcxbs	1/1	Running	0	5s
d10006	cacheserver-7dc898d4bf-8dfgx	1/1	Running	0	5s
d10006	casaccessmanagement-64b5769d8f-mlmjf	1/1	Running	0	5s
d10006	casadministration-747746f94c-j2dm2	1/1	Running	0	5s

During startup some pods restart a number of times until other pods are ready. The value in the **Status** column is `Running` or `Completed` when the pods have either fully started or completed their expected function.

Save the \$deploy Directory

The files in the \$deploy directory are used in subsequent administration tasks with your deployment, such as updating your software and applying a new license. Therefore, you must not delete the \$deploy directory. Should you choose, you can move it to a GitOps repository or other location for later use.

IMPORTANT If you move the \$deploy directory from its original location, you must notify other potential administrators of the new location.

Readiness Service

The readiness service checks the status of the SAS Viya platform to determine whether it is ready for use. The service performs all of its checks every 30 seconds. After the software is deployed, the service should be consulted to determine whether the deployment is ready for use. The readiness service is also a useful tool for the administration of SAS Viya platform throughout its life.

Checks

The sas-endpoints-ready Check

The sas-endpoints-ready checks all Kubernetes service resources in the SAS Viya platform namespace to determine if they are ready. When all the services are ready, this check will pass. If a particular endpoint has no Subset addresses at all (if the corresponding deployment has zero replicas), it will be marked ready as well. To

disable this check for a particular service, add the `sas.com/readiness-check:disabled` label to the Kubernetes Service resource.

The sas-database-ready Check

The `sas-database-ready` check validates that the database instance required by SAS is configured properly and is available for connection. There are two parts to this check:

- 1 Ensure the database is configured properly in Consul.
- 2 Attempt to connect to the database.

If you want to skip the database connection check, set the environment variable `SAS_READINESS_CHECK_DATABASE_READY_CONNECT_ENABLED` to `false`.

The sas-oauth-provider-ready Check

The `sas-oauth-provider-ready` ensures that the OAuth provider for the deployment is available, and that the SAS Viya platform can retrieve OAuth tokens from it. There are two parts to this check:

- 1 Ensure that the `sas-logon-app` Kubernetes Endpoints resource exists and is ready to receive traffic.
- 2 Attempt to retrieve an OAuth token from the SASLogon service.

Usage

To see the results of the latest readiness check:

```
kubectl wait -n name-of-namespace \
  --for=condition=ready pod \
  --selector="app.kubernetes.io/name=sas-readiness" \
  --timeout=1800s
```

If the deployment is ready, the command has the return code value 0, and the following output is displayed:

```
pod/sas-readiness-hash condition met
```

If the deployment is not ready, the command has the return code value 1, and the following output is displayed:

```
pod/sas-readiness-hash condition met
error: timed out waiting for the condition on pod/sas-readiness-hash
```

Review the Logs

For details of the readiness check's findings, review its logs:

.....
Note: The command should be entered as a single line.
.....

```
kubectl logs -n name-of-namespace --selector="app.kubernetes.io/name=sas-readiness"
```

If the deployment is ready, the log will contain a message like the following:

```
"message": "All checks passed. Marking as ready. The first recorded failure was 30s ago."
```

In order to prevent repetitive log messages, the readiness check only records the first success message despite continuing to perform its checks. As a result, the log might appear to contain a stale success message, but no new messages will be added until the status changes to “not ready”.

The log will contain a message like the following to indicate which services are failing the readiness check:

```
"message": "The check \"sas-endpoints-ready\" failed - 6 endpoints have no available addresses: sas-catalog,sas-catalog-table-bot,sas-data-flows,sas-device-management,sas-relationships,sas-studio-app"
```

Customize the Readiness Check Period

By default, the readiness service runs its set of checks every 30 seconds. You can change this value by setting the `SAS_READINESS_CHECK_PERIOD` environment variable on the `sas-readiness` pod.

Sign In as the sasboot User

Your SAS environment is deployed with an initial administrator account that is named `sasboot`. The password for this account has expired by default, so you must reset the password before you can sign in.

To reset the password:

- 1 Get the name of the pod that contains SASLogon and search for the characters, `sasboot`:

```
kubectl -n name-of-namespace logs -c sas-logon-app $(kubectl -n name-of-namespace get pods -l "app==sas-logon-app" --no-headers --sort-by=.status.startTime -o custom-columns=:metadata.name | tail -1) | grep "/SASLogon/reset_password?"
```

Note: This command will not return a value if you set the `sas.logon.initial.password` property in a `sitedefault.yaml` file.

2 Sign in from a URL with this format:

`protocol://name-of-ingress-host:port/SASLogon/reset_password?
code=password`

protocol can be http or https, depending on how you have secured your deployment.

3 Follow the instructions on the displayed web page to reset the password.

If the URL has expired, restart the `sas-logon` pod:

```
kubectl -n name-of-namespace delete pods -l "app==sas-logon-app"
```

Then go to the log and obtain the new URL. The URL expires 24 hours after the SAS Logon service restarts. For security purposes, the URL that is specified in a browser or in a text editor also expires, even if the password is not reset.

After you reset the password, SAS Environment Manager automatically opens in your browser.

4 Click **Yes** for all of the assumable groups so that you have the permissions to perform subsequent tasks.

5 Share the following URL with any other users of your SAS Viya platform so that they can access the deployment:

`protocol://name-of-ingress-host:port/SASDrive`

protocol can be http or https, depending on how you have secured your deployment.

After the password has been reset, you should consider disabling the password reset feature. For the steps to disable that feature, see [“Disable the Password Reset Feature” in SAS Viya Platform: Identity Management](#).

Note: If you have enabled multi-tenancy in your deployment, using the `sasboot` user to reset the password opens SAS Environment Manager on the provider tenant.

Promotion and Migration

If you plan to move pre-existing content from SAS 9 or SAS Viya 3, see [System Migration and Content Migration](#).

Post-Installation Tasks

Configure Identities	98
Configure an LDAP Identity Provider	98
Configure a SCIM Identity Provider	99
Configure Multi-tenancy	99
Configure the Connection to the Mail Service	100
Configure Files Service	101
Configure Monitoring and Logging	101
Configure Guest Access	101
Obtain and Run Hadoop Tracer Script	102
Overview	102
Transfer the Hadoop Tracer Script to the Hive Node	102
Tasks Before Running the Hadoop Tracer Script	103
Runtime Options	103
Run the Hadoop Tracer Script	103
Determine If the Run Was Successful	104
How to Address Failures	104
Copy the JAR and Configuration Files to the SAS Client System	105
Configure Cloud Analytic Services (CAS)	106
Configure External Access to Amazon Web Services CAS Services	106
Configure Cloud Data Exchange	107
Configure Model Access	108
Configure Access to Models	108
Configure Access to Analytic Store Model Files	108
Configure SAS Asset and Liability Management	109
Configure SAS Business Orchestration Services	109
Configure SAS Data Engineering	110
Additional Software Associated with SAS Data Engineering	110
Configure SAS Data Quality	110
Configure the Quality Knowledge Base	110
Configure SAS Dynamic Actuarial Modeling	111

Deploy SAS Enterprise Session Monitor	111
Configure SAS Event Stream Processing	111
Configure SAS Expected Credit Loss	112
Configure SAS for Microsoft 365	112
Configure SAS Intelligent Decisioning	112
Configure SAS Model Manager	112
Configure SAS Model Risk Management	113
Configure SAS Risk Engine	113
Configure SAS Risk Modeling	113
Configure SAS Stress Testing	113
Configure SAS Visual Analytics	114
Additional Software Associated with SAS Visual Analytics	114
Configure SAS Visual Statistics	114
Additional Software Associated with SAS Visual Statistics	114
Configure SAS Viya	115
Additional Software Associated with SAS Viya	115
Configure SAS Viya Advanced	115
Additional Software Associated with SAS Viya Advanced	115
Configure SAS Viya Enterprise	116
Additional Software Associated with SAS Viya Enterprise	116
Configure SAS Viya Programming	117
Additional Software Associated with SAS Viya Programming	117
Configure SAS Viya with SingleStore	118
Set the Path to the Certificate Authority for SingleStore	118
Additional Documentation	118
Configure SAS/CONNECT Spawner	118

Configure Identities

Configure an LDAP Identity Provider

After completing the installation of the SAS Viya platform, you must configure an identity provider before your users can access SAS Environment Manager and SAS Visual Analytics. LDAP is the default identity provider. Make sure your LDAP server meets the requirements that are described in [LDAP Requirements](#).

For a full set of instructions for configuring an LDAP identity provider, see “How to Configure LDAP” in [SAS Viya Platform: Identity Management](#).

Configure a SCIM Identity Provider

After completing the installation of the SAS Viya platform, you must configure the connection to your identity provider before your users can access SAS Environment Manager and other SAS Viya platform products. To use SCIM as your source for user and group identities, you must grant Microsoft Azure access to the SAS Viya platform SCIM endpoints.

Microsoft Azure

To complete SCIM identity provider configuration, you must configure the SCIM connection to the SAS Viya platform in Microsoft Azure Active Directory. You also need to disable LDAP, which is configured by default when the deployment completes. For a full set of instructions, see [“How to Configure SCIM” in SAS Viya Platform: Identity Management](#).

Amazon Web Services

You can use [Amazon Web Services](#) to access the SAS Viya platform SCIM endpoints using SCIM as your source for user and group identities. The SAS Viya platform should support authentication and identity management with Amazon Web Services Identity and Access Management, with only the initial configuration linking the SAS Viya platform to AWS IAM. Authentication should be delegated to AWS and identity information should be synchronized into the SAS Viya platform with the SCIM identity provider configuration.

Users can change the provisioning from manual to SCIM by following the instructions to enable automatic provisioning at [Enabling Single Sign-On Between OneLogin and AWS](#).

Google Cloud Platform

You can utilize [Google Cloud Platform](#) to access the SAS Viya platform SCIM endpoints to use SCIM as your source for user and group identities. To set up Google Cloud for SCIM, follow the instructions on the web page.


Configure Multi-tenancy

After deployment, a Kubernetes administrator with elevated permissions can onboard and offboard tenants, monitor services, monitor logs, stop and start




tenants, and run multi-tenant backups and restores. SAS recommends onboarding at least one tenant into a multi-tenant environment immediately after deployment. For information about performing each of these tasks, see [SAS Viya Platform: Multi-tenancy](#).

Configure the Connection to the Mail Service

After performing a new deployment of the SAS Viya platform, you must configure the connection to your mail service. Complete these steps while you are signed in as one of the SAS Administrators.

- 1 Select the  from the side menu to open the Configuration page.
- 2 On the Configuration page, select **All Services** from the list, and then select **Mail service** from the list of services.

Note: If **Mail service** is not listed, then there is no configuration to be performed. You should skip the rest of this topic.

- 3 In the **sas.mail** section, click . In the Edit Configuration window, follow these steps:
 - a Specify a value for the following required fields: **host** and **port**. For the remaining fields, review the default values and make changes, as necessary. The default values are appropriate for most sites.
 - b Click **Save**.
- 4 (Optional) To enable the health check for the mail service, perform the following steps.
 - a Select the  from the side menu to open the Configuration page.
 - b On the Configuration page, select **All Services** from the list, and then select **Mail service** from the list of services.
 - c In the **management.health.mail** section, click .
 - d Turn the **enabled** toggle to on.
 - e Click **Save**.

When this toggle is set, health checks will be enabled after the mail service is restarted. If the mail host is not configured or is configured incorrectly, or if it cannot connect to the SMTP mail server, the mail service will indicate it is in a failed state.

Configure Files Service

After deploying the SAS Viya platform, the Files service must be configured to provide in-line responses to file contents. For the steps to perform this configuration, see [“Managing Cross-Site Scripting Risk for File Uploads” in SAS Viya Platform: Overview](#).

Configure Monitoring and Logging

Monitoring performance metrics and the logs generated by your deployment helps your SAS administrators ensure that the SAS Viya platform is running efficiently and enable them to quickly detect any problems. You can use your preferred monitoring technology to monitor your SAS Viya platform deployment. If you do not have a preferred technology, SAS provides solutions for monitoring and logging that are based on widely-used open source monitoring technologies including Prometheus, Grafana, and OpenSearch.

You can deploy these solutions from the [SAS Viya Platform Monitoring for Kubernetes](#) GitHub site.

The monitoring deployment includes these components:

- Prometheus
- Prometheus Operator
- Alert Manager
- Grafana
- Prometheus Pushgateway
- Grafana dashboards

The logging deployment includes these components:

- Fluent Bit
- OpenSearch
- OpenSearch Dashboards

Configure Guest Access

Guest access is an optional feature that provides anonymous Read-Only access to a subset of resources and functionality in participating applications. For information

about configuring your SAS Viya platform deployment for guest access, see “[Guest Access](#)” in *SAS Viya Platform: Authentication*.

Obtain and Run Hadoop Tracer Script

Note: The Hadoop tracer script is included with SAS/ACCESS Interface to Hadoop and SAS In-Database Technologies for Hadoop Cloud Services. The Hadoop tracer script might also be required by additional components that access Hadoop. If you do not need to run the Hadoop tracer script, then skip this section.

Overview

The Hadoop tracer script is a Python script that traces the system calls of various Hadoop client tools and uses the trace data to identify required client JAR files and configuration files. The script determines which specific Hadoop distribution supplied the JARs and configuration files that SAS products need for connectivity between Hadoop client machines and the Hadoop server environment. Therefore, you must run this script on the Hive node of your Hadoop cluster in order to prepare your Hadoop environment for a SAS Viya platform deployment.

The script is organized into two parts. The first part is a list of the Hadoop services that the script is tracing to a new `driver.json` file. The second part is the remainder of the code. This division allows the list of Hadoop services and additional JAR files to be modified manually without having to make Python code changes.

If the `driver.json` file is moved or removed from its default location, the Hadoop tracer script fails with an error message indicating that the file is missing.

Transfer the Hadoop Tracer Script to the Hive Node

The Hadoop tracer script runs on the Hive node of your Hadoop cluster.

- 1 On your Hadoop cluster, create a temporary directory to hold a ZIP file that you download later. For example, `/tmp/sas/hadooptracer`.
- 2 Copy and paste the following URL into a browser. Download the `hadooptracer.zip` file from that FTP site to the directory that you created in step 1.

`ftp.sas.com/techsup/download/blind/access/hadooptracer.zip`
- 3 Using a method of your choice (such as PSFTP, SFTP, SCP, or FTP), transfer the ZIP file to the Hive node on your Hadoop cluster.
- 4 Unzip the file. The `hadooptracer_py` and `driver.json` files are included in the ZIP file.

- 5 Change permissions on the `hadooptracer_py` file to include the Execute permission:

```
chmod 755 ./hadooptracer_py
```

Tasks Before Running the Hadoop Tracer Script

- Ensure that the user running the script has authorization to issue HDFS and Hive commands.
- If Hadoop is secured with Kerberos, obtain a Kerberos ticket for the user before running the script.
- Ensure that Python and the `strace` Linux library have been installed on the Hadoop cluster. Install them from the package repositories for your Linux distribution if necessary. Python 2.6 or later is required.
- If you want to pull one or more JAR files that are not included in the output of the Hadoop tracer script, modify the `driver.json` file by adding the JAR files to the `ExtraJarFiles` property at the bottom of the file. Here is an example of adding a new JAR file named `xyz-service.jar` to the `ExtraJarFiles` property:

```
"ExtraJarFiles": ["jline-*", "jruby-complete-*", "hive-warehouse-connector-assembly-*", "xyz-service-*"]
```

Runtime Options

For a list of the options available when running the Hadoop tracer script:

```
python ~/hadooptracer_py --help
```

Review these options before you run the script.

Run the Hadoop Tracer Script

Use this command to run the Hadoop tracer script:

```
python ./hadooptracer_py --filterby=latest --postprocess --jsonfile ./driver.json
```

If you want to pull the necessary JAR files without filtering, use `filterby=none`, or omit the `filterby=` option.

TIP To collect only the Hadoop configuration files and exclude JAR files:

```
python ./hadooptracer_py --svckey=Hadoop
```

TIP The `postprocess` option ensures that the `${hdp.version}` tokens are replaced. SAS strongly recommends that you run the tracer script with this option. This option is ignored for Cloudera clusters.

The Hadoop tracer script performs the following tasks:

- collects the necessary Hadoop JAR and configuration files from nodes in the cluster and copies them in the `/tmp/jars` directory and the `/tmp/sitexmls` directory, respectively.
- creates a `hadooptracer.json` file in the `/tmp` directory. If you need a custom path for the JSON output file, use this command instead:

```
python ./hadooptracer_py -f /your-path/hadooptracer.json
```

- creates a log at `/tmp/hadooptracer.log`. If you need a custom path for the log file, add this option:

```
--logfile your-path/log-file-name
```

- prints more debug entries to the log file when the `--debug` option is added.

IMPORTANT The Hadoop JAR and configuration files on the SAS client machine must be kept in sync with the Hadoop configuration. After a Hadoop cluster update, re-run the Hadoop tracer script to collect any new files, and then copy those files to the SAS client machine, replacing any previous versions of the files.

Determine If the Run Was Successful

- 1 Ensure that the required Hadoop JAR files are collected from the Hadoop cluster and placed in the `./jars` directory.

```
ls -l ./jars
```

- 2 Ensure that the required Hadoop configuration files are collected from the Hadoop cluster and placed in `./confs` directory.

```
ls -l ./confs
```

- 3 Review the `hadooptracer.log` file that is located in the default location, `/tmp`, or the custom location that you specified.

How to Address Failures

Most errors with the Hadoop tracer script stem from improper usage or an incorrect cluster configuration. If there are problems with the Hadoop cluster, they will typically show up in the stdout of the Hadoop tracer script in the form of Java traceback information.

Another common problem occurs when users try to run the Hadoop tracer script on a cluster node that doesn't have `Hadoop/Hive/HDFS/Yarn/Pig/etc` in an available `PATH`. For example,

```
2020-04-07 12:16:51,036 hadooptracer [ERROR] pig is not found in the $PATH
```

Inspect `hadooptracer.log`, located in `/tmp` by default, and use the rest of this troubleshooting section to resolve common issues. Some error messages in the console output for `hadooptracer_py` are normal and do not necessarily indicate a problem with the JAR and configuration file collection process. However, if the files are not collected as expected or if you experience problems connecting to Hadoop with the collected files, contact SAS Technical Support and include the `hadooptracer.log` and the `hadooptracer.json` files.

Copy the JAR and Configuration Files to the SAS Client System

SAS 9.4 and SAS Viya 3.5

- 1 On the SAS client machine, create two directories to hold the JAR and configuration files. For example, the `/opt/sas/hadoopfiles/lib` and `/opt/sas/hadoopfiles/conf` directories.
- 2 Using a method of your choice (such as PSFTP, SFTP, SCP, or FTP), copy the files in the `tmp/jars` and `tmp/conf`s directories on the Hadoop server to the directories on the SAS client machine that you created.

Note: If you connect to the Hadoop server with an HTTP REST API, you do not need the Hadoop JAR files on the SAS client machine.

The SAS Viya Platform in a Kubernetes Environment

For more information about configuring SAS/ACCESS Interface to Hadoop, see the README file at `$deploy/sas-bases/examples/data-access/README.md` (for Markdown) or at `$deploy/sas-bases/docs/configuring_sasaccess_and_data_connectors_for_sas_viya_4.htm` (for HTML).

Configure Cloud Analytic Services (CAS)

Configure External Access to Amazon Web Services CAS Services

To gain secure external access to CAS services running in an AWS cloud, additional steps are necessary after the load balancer setup described in [“Configure External Access to CAS” on page 58](#) has been completed and the cluster has been deployed.

All load balancer hosts must be DNS-aliased to the service names generated by setting up the load balancer. You may select an extant DNS domain/subdomain or create a new one specifically to support access to AWS CAS clusters. You may use your own DNS service or transfer DNS management of the domain to the AWS Route 53 service.

Note: AWS Route 53 offers some advantages for large scale or more complex deployment cases due to the additional routing and monitoring capabilities that it offers. See [What is Amazon Route 53?](#) for details.

Regardless of the DNS service used, you must add CNAME records to the routing table to direct references to the CAS-generated names to the load balancer FQDN. CAS-generated external service names are predetermined based upon the suffix supplied in the earlier setup. Here are examples of those names when using Data Connect Accelerators and Spark EPCS access in addition to the general ports:

```
sas-cas-server-default-bin-pisces.viya.acme.com
sas-cas-server-default-epcs-pisces.viya.acme.com
sas-cas-server-default-controller-dc-pisces.viya.acme.com
sas-cas-server-default-backup-dc-pisces.viya.acme.com
sas-cas-server-default-worker-0-dc-pisces.viya.acme.com
sas-cas-server-default-worker-1-dc-pisces.viya.acme.com
```

To find the load balancer names:

```
kubectl -n name-of-namespace get svc | grep LoadBalancer
```

Here is a partial sample of typical output, revealing the load balancer DNS names:

```

sas-cas-server-default-backup-dc LoadBalancer 10.100.178.234
a7ccc6edbca1d402c8c72ea7da3543d7-1688256446.us-east-1.elb.amazonaws.com
sas-cas-server-default-bin LoadBalancer 10.100.169.27
a5a90408edfc246bba051f830d6e58fe-2016925591.us-east-1.elb.amazonaws.com
sas-cas-server-default-controller-dc LoadBalancer 10.100.30.89
a14ca0a2c8009415891c05053ee36d83-222471472.us-east-1.elb.amazonaws.com
sas-cas-server-default-eps LoadBalancer 10.100.105.198
a8207fb65f9464e7680948e8e8d9e722-58303842.us-east-1.elb.amazonaws.com
sas-cas-server-default-http LoadBalancer 10.100.65.146
aa21afd024de44d64bffe557add1e47a-1604903159.us-east-1.elb.amazonaws.com
sas-cas-server-default-worker-0-dc LoadBalancer 10.100.4.214
a307341cce8694e639c6fdee9096cb5c-632571573.us-east-1.elb.amazonaws.com
sas-cas-server-default-worker-1-dc LoadBalancer 10.100.224.50
a04b03687b3704b79820fb72869eb5c5-1997141149.us-east-1.elb.amazonaws.com

```

With this information you can then cut your CNAME records in your DNS service of choice in this arrangement:

```

sas-cas-server-default-backup-dc-pisces.viya.acme.com =>
a7ccc6edbca1d402c8c72ea7da3543d7-1688256446.us-east-1.elb.amazonaws.com
sas-cas-server-default-bin-pisces.viya.acme.com =>
a5a90408edfc246bba051f830d6e58fe-2016925591.us-east-1.elb.amazonaws.com
sas-cas-server-default-controller-dc-pisces.viya.acme.com =>
a14ca0a2c8009415891c05053ee36d83-222471472.us-east-1.elb.amazonaws.com
...

```

Configure Cloud Data Exchange

Note: For a full description of the deployment of Cloud Data Exchange and SAS Data Agent, see [“What is the deployment process for SAS Data Agent?” in *Getting Started with SAS Viya Platform Operations*](#).

Cloud Data Exchange is included in a number of offerings from SAS:

- SAS Data Engineering Advanced
- SAS Intelligent Decisioning
- SAS Visual Analytics
- SAS Visual Forecasting
- SAS Visual Statistics
- SAS Visual Text Analytics
- SAS Viya
- SAS Viya Advanced
- SAS Viya Enterprise
- SAS Viya Programming

If you have at least one of these offerings in your software order, you must complete your deployment of Cloud Data Exchange to take advantage of its data management and connection capabilities.

After you have finished your deployment of the SAS Viya platform, you must perform a few more tasks to finish deploying Cloud Data Exchange:

- 1 Ensure that you have deployed the co-located SAS Data Agent by following the steps in the README file located at `$deploy/sas-bases/examples/sas-data-agent-server-colocated/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/configure_a_co-located_sas_data_agent.htm` (for HTML format).
- 2 Deploy the remote SAS Data Agent using the procedures described in [Remote SAS Data Agent: Deployment Guide](#).
- 3 Perform the post-installation tasks described at “Once Co-located SAS Data Agent and Remote SAS Data Agent Are Running” in [Cloud Data Exchange for the SAS Viya Platform: Administrator’s Guide](#).

Configure Model Access

Configure Access to Models

Note: SAS Model Manager is only available as part of the SAS Model Manager, SAS Viya, SAS Viya Advanced, and SAS Viya Enterprise offerings.

In order to import models into SAS Model Manager and register models from Model Studio, SAS Visual Analytics, and SAS Studio into the common model repository, as well as add a model from the common model repository into a decision flow, users must have the appropriate access permissions. For more information, see “[Access to Models](#)” in [SAS Viya Platform: Models Administration](#).

Configure Access to Analytic Store Model Files

Note: SAS Model Manager is only available as part of the SAS Model Manager, SAS Viya, SAS Viya Advanced, and SAS Viya Enterprise offerings. SAS Intelligent Decisioning is only available as part of the SAS Viya Enterprise and SAS Intelligent Decisioning offerings. If your deployment does not include any of these offerings, you should skip this section.

CAUTION

To use analytic store models, you must have completed the tasks described in the README file, “SAS Micro Analytic Service ASTORE Configuration”. The README file is located at `$deploy/sas-bases/examples/sas-microanalytic-score/astores/`

`README.md` (for Markdown) and `$deploy/sas-bases/docs/sas_micro_analytic_service_astore_configuration.htm` (for HTML).

In order to publish analytic store models from SAS Model Manager, Model Studio, and SAS Intelligent Decisioning to the SAS Micro Analytic Service destination, see “[Accessing Analytic Store Model Files](#)” in *SAS Viya Platform: Models Administration*.

Configure SAS Asset and Liability Management

To complete the configuration of this product, you must perform the post-installation tasks described in *SAS Asset and Liability Management: Administrator's Guide* (access key required).

Configure SAS Business Orchestration Services

SAS Business Orchestration Services requires a set of configuration files. When you have completed the deployment of the SAS Viya platform and SAS Business Orchestration Services, you must create an init container image that includes all of the required files, and you must add a reference to the init container to the base `kustomization.yaml` file. Instructions are provided in a README file located at `$deploy/sas-bases/examples/sas-boss/README.md` (for Markdown format) or `$deploy/sas-bases/docs/deploying_sas_business_orchestration_services.htm` (for HTML format).

The product documentation provides additional instructions for configuring and using SAS Business Orchestration Services. The documentation is available from the SAS Support website at <https://support.sas.com/en/software/business-orchestration-support.html>.

Configure SAS Data Engineering

Additional Software Associated with SAS Data Engineering

If you want to access Cloud Analytic Services (CAS) with R or Python, you must download and deploy SWAT packages for them. Those packages are available at the following locations:

- [SWAT for Python](#)
- [SWAT for R](#)

A SAS Viya 3.5 software order is associated with deploying SAS Data Engineering. The software order contains SAS Embedded Process for Hadoop and SAS Embedded Process for Teradata on Linux. The Software Order Email (SOE) specifies the location of the deployment documentation: [SAS Viya In-Database Technologies: Deployment and Administration Guide](#). You should deploy this order only if you are using SAS In-Database Technologies for Hadoop Cloud Services or SAS In-Database Technologies for Teradata.

A second SAS Viya platform order is also associated with deploying SAS Data Engineering. It contains SAS Embedded Process for Spark.

Note: The SOEs for the associated orders do not refer to SAS Data Engineering explicitly.

Configure SAS Data Quality

Configure the Quality Knowledge Base

For an overview of SAS Data Quality, see “[Overview](#)” in [SAS Viya Platform: QKB Management](#). SAS Data Quality relies on a collection of rules and reference data called a SAS Quality Knowledge Base (QKB).

Your deployment includes the latest version of the QKB for Contact Information. It is configured to use the `English, United States` locale by default. To deploy a custom QKB into CAS, you should first create a QKB Archive and import it into CAS using SAS Environment Manager. For information about these tasks, see “[Create a QKB Archive \(QARC\) File](#)” in [SAS Viya Platform: QKB Management](#).

If you are upgrading from SAS Viya 3.5, you must manually import any older or custom QKBs into this version. For instructions to import the QKB into SAS Studio, see “[How To \(Kubernetes\)](#)” in *SAS Viya Platform: QKB Management* and “[Set the Default QKB and the Default Locale](#)” in *SAS Viya Platform: QKB Management*. To import a QKB into CAS, create a QARC archive and use the SAS Environment Manager; for details, see “[Create a QKB Archive \(QARC\) File](#)” in *SAS Viya Platform: QKB Management*.

Configure SAS Dynamic Actuarial Modeling

Additional tasks are required to deploy this product. For details, see [SAS Dynamic Actuarial Modeling: Administrator's Guide](#) (access key required).

Deploy SAS Enterprise Session Monitor

SAS Enterprise Session Monitor extends the monitoring and scheduling capabilities of your SAS Viya platform deployment by providing users with easy access to metrics that lead to resource optimization and enhanced throughput. However, deploying SAS Enterprise Session Monitor requires a separate download and deployment process from the rest of the SAS Viya platform.

If you have ordered SAS Enterprise Session Monitor, you will have received an email, separate from your Software Order Email, describing where and how to download the required software. The email describes the location of the deployment, administration, and usage documentation. Follow the instructions in that email to download and deploy SAS Enterprise Session Monitor.

Configure SAS Event Stream Processing

In order to start using SAS Event Stream Processing, you must first use SAS Event Stream Processing Studio to launch a project. This causes the ESP Operator to start an ESP Server pod, which then runs the project.

For more information, see “[Running Event Stream Processing Projects](#)” in *SAS Event Stream Processing: Overview*.

Configure SAS Expected Credit Loss

Additional tasks are required to deploy this product. For details, see [SAS Expected Credit Loss: Administrator's Guide](#) (access key required).

Configure SAS for Microsoft 365

For the information to complete the configuration of this product, including required security configuration, see “Steps for SAS Administrator” in [SAS for Microsoft 365: User's Guide](#).

Configure SAS Intelligent Decisioning

.....
Note: SAS Intelligent Decisioning is only available as part of the SAS Viya Enterprise and SAS Intelligent Decisioning offerings. If your deployment does not include any of these offerings, you should skip this section.
.....

You must perform post-installation tasks such as configuring access to analytic store model files and setting configuration properties for SAS Intelligent Decisioning. For more information, see [SAS Intelligent Decisioning: Administrator's Guide](#).

Configure SAS Model Manager

.....
Note: SAS Model Manager is only available as part of the SAS Model Manager, SAS Viya, SAS Viya Advanced, and SAS Viya Enterprise offerings. If your deployment does not include any of these offerings, you should skip this section.
.....

You must perform post-installation tasks such as configuring access to analytic store model files and a workflow client user account for use with SAS Model Manager. For more information, see [SAS Model Manager: Administrator's Guide](#).

Configure SAS Model Risk Management

To complete the configuration of this product, you must perform post-installation tasks to configure additional access control. For more information, see [“Designing and Implementing Access Control”](#) in *SAS Risk Cirrus: Administrator’s Guide*.

Configure SAS Risk Engine

To complete the configuration of SAS Risk Engine, you must perform post-installation tasks such as verifying the deployment and granting users access to objects in the user interface. For more information, see [SAS Risk Engine: Administration](#) (access key required).

Configure SAS Risk Modeling

To complete the configuration of this product, you must perform post-installation tasks such as working with user groups and creating a destination for publishing a model. For more information, see [SAS Risk Modeling: Administrator’s Guide](#) (access key required).

Configure SAS Stress Testing

Additional tasks are required to deploy this product. For details, see [SAS Stress Testing: Administrator’s Guide](#) (access key required).

Configure SAS Visual Analytics

Additional Software Associated with SAS Visual Analytics

If you want to access Cloud Analytic Services (CAS) with R or Python, you must download and deploy SWAT packages for them. Those packages are available at the following locations:

- [SWAT for Python](#)
- [SWAT for R](#)

A SAS 9.4 order is associated with deploying SAS Visual Analytics. The software order contains SAS Visual Analytics Add-In for Office. The SOE specifies the location of the deployment documentation.

Note: The SOE for the associated order does not refer to SAS Visual Analytics explicitly.

Configure SAS Visual Statistics

Additional Software Associated with SAS Visual Statistics

If you want to access Cloud Analytic Services (CAS) with R or Python, you must download and deploy SWAT packages for them. Those packages are available at the following locations:

- [SWAT for Python](#)
- [SWAT for R](#)

A SAS 9.4 order is associated with deploying SAS Visual Statistics. The software order contains SAS Visual Analytics Add-In for Office. The SOE specifies the location of the deployment documentation.

Note: The SOE for the associated order does not refer to SAS Visual Statistics explicitly.

Configure SAS Viya

Additional Software Associated with SAS Viya

Note: Before the 2023.01 version of the SAS Viya platform, this offering was named SAS Visual Machine Learning.

If you want to access Cloud Analytic Services (CAS) with R or Python, you must download and deploy SWAT packages for them. Those packages are available at the following locations:

- [SWAT for Python](#)
- [SWAT for R](#)

A SAS 9.4 order is associated with deploying SAS Viya. The software order contains SAS Visual Analytics Add-In for Office. The SOE specifies the location of the deployment documentation.

Note: The SOE for the associated order does not refer to SAS Viya explicitly.

Configure SAS Viya Advanced

Additional Software Associated with SAS Viya Advanced

Note: Before the 2023.01 version of the SAS Viya platform, this offering was named SAS Visual Data Science.

If you want to access Cloud Analytic Services (CAS) with R or Python, you must download and deploy SWAT packages for them. Those packages are available at the following locations:

- [SWAT for Python](#)

- [SWAT for R](#)

A SAS 9.4 order is associated with deploying SAS Viya Advanced. The software order contains SAS Visual Analytics Add-In for Office. The SOE specifies the location of the deployment documentation.

Note: The SOE for the associated order does not refer to SAS Viya Advanced explicitly.

Configure SAS Viya Enterprise

Additional Software Associated with SAS Viya Enterprise

Note: Before the 2023.01 version of the SAS Viya platform, this offering was named SAS Visual Data Science Decisioning.

If you want to access Cloud Analytic Services (CAS) with R or Python, you must download and deploy SWAT packages for them. Those packages are available at the following locations:

- [SWAT for Python](#)
- [SWAT for R](#)

A SAS Viya 3.5 software order is associated with deploying SAS Viya Enterprise. The software order contains SAS Embedded Process for Hadoop and SAS Embedded Process for Teradata on Linux. The Software Order Email (SOE) specifies the location of the deployment documentation: [SAS Viya In-Database Technologies: Deployment and Administration Guide](#). You should deploy this order only if you are using SAS In-Database Technologies for Hadoop Cloud Services or SAS In-Database Technologies for Teradata.

A SAS 9.4 order is associated with deploying SAS Viya Enterprise. The software order contains SAS Visual Analytics Add-In for Office. The SOE specifies the location of the deployment documentation.

A second SAS Viya platform order is also associated with deploying SAS Viya Enterprise. It contains SAS Embedded Process for Spark.

Note: The SOEs for the associated orders do not refer to SAS Viya Enterprise explicitly.

Configure SAS Viya Programming

Additional Software Associated with SAS Viya Programming

Note: Before the 2023.01 version of the SAS Viya platform, this offering was named SAS Data Science Programming.

If you want to access Cloud Analytic Services (CAS) with R or Python, you must download and deploy SWAT packages for them. Those packages are available at the following locations:

- [SWAT for Python](#)
- [SWAT for R](#)

A SAS Viya 3.5 software order is associated with deploying SAS Viya Programming. The software order contains SAS Embedded Process for Hadoop and SAS Embedded Process for Teradata on Linux. The Software Order Email (SOE) specifies the location of the deployment documentation: [SAS Viya In-Database Technologies: Deployment and Administration Guide](#). You should deploy this order only if you are using SAS In-Database Technologies for Hadoop Cloud Services or SAS In-Database Technologies for Teradata.

A SAS 9.4 order is associated with deploying SAS Viya Programming. The software order contains SAS Visual Analytics Add-In for Office. The SOE specifies the location of the deployment documentation.

A second SAS Viya platform order is also associated with deploying SAS Viya Programming. It contains SAS Embedded Process for Spark.

Note: The SOEs for the associated orders do not refer to SAS Viya Programming explicitly.

Configure SAS Viya with SingleStore

Set the Path to the Certificate Authority for SingleStore

Encryption between CAS and SingleStore is supported under the same TLS mode that you set for SAS Viya. However, you can specify the path to a different certificate authority file by setting the `ssl_ca=` option in the SingleStore data connector. For more information about the `ssl_ca=` option, see [“Security” in SAS Viya Platform with SingleStore: Administration and Configuration Guide](#).

Additional Documentation

For more information about configuring and administering SAS Viya with SingleStore, see [SAS Viya Platform with SingleStore: Administration and Configuration Guide](#).

Configure SAS/CONNECT Spawner

If you are performing a new deployment and added a reference to the `enable-spawned-servers.yaml` example file to your base `kustomization.yaml` in order to spawn SAS/CONNECT servers in the spawner pod, you must make a change to the configuration in SAS Environment Manager.

- 1 Go to SAS Environment Manager.
- 2 Select **Configurations**, then **SAS/CONNECT Spawner**.
- 3 In the contents field for **sas.connect.spawner: startup_commands**, remove `- nolocallaunch` from the `USERMODS` line.
- 4 Click **Save**.

For more information about this task, see the “Allow the Ability to Spawn Servers within the Spawner Pod” section of the “Configure SAS/CONNECT Spawner in SAS Viya” README file, located at `$deploy/sas-bases/examples/sas-connect-spawner/README.md` (for Markdown format) and at `$deploy/sas-bases/docs/configure_sasconnect_spawner_in_the_sas_viya_platform.htm` (for HTML format).

Validating the Deployment

<i>The SAS Viya Platform: Deployment Validation</i>	119
<i>The SAS Viya Platform and the SAS Operational Quality Tool</i>	119

The SAS Viya Platform: Deployment Validation

Validating your SAS Viya platform deployment is not a Kubernetes-based task and requires a SAS administrator. The procedures for validating products in your deployment are described in [SAS Viya Platform: Deployment Validation](#).

The SAS Viya Platform and the SAS Operational Quality Tool

SAS Operational Qualification Tool (SAS OQ) helps qualify the use of SAS software in regulated industries. SAS OQ supports the qualification aspect of the essential migration, integration, and verification processes that customers must perform to validate the SAS deployment. SAS OQ helps to demonstrate that SAS is operational. For more information about SAS OQ, see [SAS Operational Qualification Tool: User's Guide](#).

Modifying the Deployment Configuration

<i>Modify Existing Customizations in a Deployment</i>	121
---	-----

Modify Existing Customizations in a Deployment

After your software has been deployed, you might want to modify some of the customizations that you made in your initial deployment. For example, you might want to move from an SMP deployment of CAS to an MPP deployment. Or you might want to move from using TLS in front-door mode to using full-stack TLS.

Modifying your configuration does not require new software. To modify the customizations, simply repeat some of the steps that you performed in the initial deployment of your software:

IMPORTANT Performing these steps causes an outage while the software is re-deployed with the new configuration settings. Ensure that you plan for an outage before continuing with the steps..

- 1 Perform the modifications. Refer to the topic in [“Common Customizations” on page 36](#) or the README file that describes the changes that should be made. The README indexes, located at `$deploy/sas-bases/README.md` for README files in Markdown language or `$deploy/sas-bases/docs/index.htm` for README files in HTML, can help direct you to the appropriate README.
- 2 Redeploy the software.
 - If you are deploying with the SAS Viya Platform Deployment Operator, first modify the SASDeployment custom resource with the command provided at

[“Run the create sas-deployment-cr Command” on page 78](#). Then run the deployment command described at [“Command and Output” on page 84](#).

- If you are deploying with the Orchestration Tool, perform the command described at [“Command” on page 87](#).
- If you are deploying with Kubernetes commands, run the series of commands that are described at [“Deployment Using Kubernetes Commands” on page 90](#).

Uninstalling

<i>Uninstall with the SAS Viya Platform Deployment Operator</i>	123
Uninstall a SAS Viya Platform Deployment from a Namespace Only	123
Remove All SAS Viya Platform Deployments in the Cluster and the SAS Viya Platform Deployment Operator	126
<i>Remove the SAS Viya Platform Deployment Operator Only</i>	128
<i>Uninstall Deployments That Do Not Use the Deployment Operator</i>	129
Remove the SAS Prepull and SAS Workload Orchestrator ClusterRoles and ClusterRoleBindings	129
Remove Service Account Links and SCCs	129
Remove Internal PostgreSQL Resources	130
Remove the SAS Viya Platform	131
Remove CRDs	131

Uninstall with the SAS Viya Platform Deployment Operator

The following steps assume that the SAS Viya platform is running and functional. If the software is not running, start it. If the software is not functioning, contact SAS Technical Support.

Uninstall a SAS Viya Platform Deployment from a Namespace Only

The steps in this section remove the SAS Viya platform from a single namespace, but leave the SAS Viya Platform Deployment Operator running in the cluster. As long as any namespace in the cluster is running and managed by the operator, do not remove the operator. If you want to remove all of the deployments from all of the namespaces in a cluster and also remove the operator from that cluster, see

“Remove All SAS Viya Platform Deployments in the Cluster and the SAS Viya Platform Deployment Operator” on page 126.

Remove the SAS Prepull and SAS Workload Orchestrator ClusterRoles and ClusterRoleBindings

Before you begin to uninstall your SAS Viya platform software, perform the following tasks to account for some special cluster-wide permissions.

- 1 If you enabled the Image Staging Node List Option in your deployment, follow the instructions in the “Disable the Node List Option” section of the “SAS Image Staging Configuration Options” README located at `$deploy/sas-bases/examples/sas-prepull/README.md` (for Markdown format) and `$deploy/sas-bases/docs/sas_image_staging_configuration_options.htm` (for HTML format).
- 2 If your deployment includes Workload Management and Cluster Role was enabled for it, follow the instructions in the “Disable the Cluster Role” section of the “Cluster Privileges for SAS Workload Orchestrator Service” README located at `$deploy/sas-bases/examples/sas-workload-orchestrator/configure/README.md` (for Markdown format) and `$deploy/sas-bases/docs/configuration_settings_for_sas_workload_orchestrator_service.htm` (for HTML format).

Remove SAS Viya Platform from a Namespace

When the operator deploys the SAS Viya platform, it assigns an ownerReference from the custom resource to every namespace-scoped resource deployed by the operator. When the custom resource is deleted, Kubernetes deletes anything that was owned by that resource, thereby uninstalling the SAS Viya platform. However, some resources in a SAS Viya platform deployment are cluster-scoped (such as CRDs, ClusterRoles, and ClusterRoleBindings). These are shared by all SAS Viya platform deployments in the cluster, so they are not considered “owned” by any one custom resource. Therefore, they are not deleted when the custom resource is removed for a namespace.

- 1 Remove the resources for an internal instance of PostgreSQL from the namespace:

Note: If your deployment does not include an internal instance of PostgreSQL, skip this step.

```
kubectl -n name-of-namespace delete postgresclusters --selector="sas.com/
deployment=sas-viya"
```

Note: If your Crunchy Data PVs do not have a reclaimPolicy of "retain", performing this command will delete their associated PVCs, resulting in the data they contain being lost. Therefore, ensure that you have backed up that data or are otherwise unconcerned about losing it.

If your Crunchy Data PVs do have a reclaimPolicy of "retain", their PVCs will not be deleted when the cluster is deleted. If you no longer need them, you can delete them manually by using the `kubectl get pvc` command to get the list of PVCs and the `kubectl delete` command to delete them.

- 2 Remove the SASDeployment custom resource from the namespace. If you kept your `$deploy-sasdeployment.yaml` file, run the following command:

```
kubectl -n name-of-namespace delete -f $deploy-sasdeployment.yaml
```

If you did not keep the `$deploy-sasdeployment.yaml` file, run the following command:

```
kubectl -n name-of-namespace delete sasdeployments SASdeployment-custom-resource-name
```

Remove Service Account Links

Note: If your deployment is not on Red Hat OpenShift, skip this section.

For Red Hat OpenShift systems, the namespace's service accounts must be removed from the assigned SCCs. To find the namespace's service accounts that are linked to SCCs, run the following code to generate the commands to review and execute:

```
SAS_NS=name-of-namespace
for oc_scc in $(oc get clusterrolebinding | grep
"system:openshift:scc:" | awk '{ print $1 }'); do
    crb_name=$(oc get clusterrolebinding $oc_scc -o json | jq -r
'.roleRef.name')
    sa_names=$(oc get clusterrolebinding $oc_scc -o json | jq -r
'.subjects[] | select(.namespace == \"'$SAS_NS'\") | .name')
    for sa_name in $sa_names; do
        if [[ ! -z $sa_name ]]; then
            echo "oc -n $SAS_NS adm policy remove-scc-from-user $
{crb_name##*:.} -z $sa_name"
        fi
    done
done
```

Run each command that is produced as output from the code.

Remove All SAS Viya Platform Deployments in the Cluster and the SAS Viya Platform Deployment Operator

Remove the SAS Prepull and SAS Workload Orchestrator ClusterRoles and ClusterRoleBindings

Before you begin to uninstall your SAS Viya platform software, perform the following tasks to account for some special cluster-wide permissions.

- 1 If you enabled the Image Staging Node List Option in your deployment, follow the instructions in the “Disable the Node List Option” section of the “SAS Image Staging Configuration Options” README located at `$deploy/sas-bases/examples/sas-prepull/README.md` (for Markdown format) and `$deploy/sas-bases/docs/sas_image_staging_configuration_options.htm` (for HTML format).
- 2 If your deployment includes Workload Management and Cluster Role was enabled for it, follow the instructions in the “Disable the Cluster Role” section of the “Cluster Privileges for SAS Workload Orchestrator Service” README located at `$deploy/sas-bases/examples/sas-workload-orchestrator/configure/README.md` (for Markdown format) and `$deploy/sas-bases/docs/configuration_settings_for_sas_workload_orchestrator_service.htm` (for HTML format).

Remove the Deployments and the Operator

To uninstall all the SAS Viya platform deployments in a cluster, repeat the instructions in [“Remove SAS Viya Platform from a Namespace” on page 124](#) for each namespace in the cluster that contains a SAS Viya platform deployment. After all the SAS Viya platform deployments have been removed, run the following command from the `$operator-deploy` directory to remove the SAS Viya Platform Deployment Operator:

```
kustomize build . | kubectl -n name-of-deployment-operator-namespace delete -f -
```

Remove the Operator ClusterRoles

After all of the SAS Viya Platform Deployment Operator instances are removed from the cluster, some shared cluster resources must be removed.

```
kubectl delete clusterroles \
  --selector "app.kubernetes.io/name=sas-deployment-operator"
```

Remove the Operator CustomResourceDefinition

After all of the SASDeployment custom resources have been removed from the cluster, the operator CustomResourceDefinition must be removed.

```
kubectl delete crds \
  --selector "app.kubernetes.io/name=sas-deployment-operator"
```

Additional Clean-up of Cluster Resources

Remove cluster-wide resources that have been left behind by the previous commands.

```
kubectl delete crd --selector "sas.com/admin=cluster-api"
kubectl delete crd --selector "sas.com/admin=cluster-wide"
kubectl delete clusterrole --selector "sas.com/admin=cluster-wide"
```

Remove SCCs

Note: If your deployment is not on Red Hat Openshift, skip this section.

If you have removed the last namespace running the SAS Viya platform in the cluster, then the SAS-specific SCCs can be removed. The following can be run to generate the commands to review and execute:

Note: Removing SCCs is a cluster-wide operation. Ensure that no namespaces with SAS Viya platform software exist in the cluster before removing SCCs. Otherwise, any remaining SAS Viya platform deployments in the cluster may cease to work as expected.

```
for sas_scc in $(oc get scc | grep sas | awk '{ print$1 }'); do echo
"oc delete scc $sas_scc"; done
for sas_scc in $(oc get scc | grep pgo | awk '{ print$1 }'); do echo
"oc delete scc $sas_scc"; done
```

Run each command that is produced as output from the code.

Validate the Removal

- 1 After you have removed the last namespace with SAS Viya platform software from the cluster, run the following commands:

```
kubectl get crd --selector "sas.com/admin=cluster-wide"
kubectl get crd --selector "sas.com/admin=cluster-api"
kubectl get clusterrole --selector "sas.com/admin=cluster-wide"
```

If you removed the cluster-wide CRDs, each command should return zero results. If you do get results, repeat the steps in [“Additional Clean-up of Cluster Resources” on page 127](#).

- 2 If you are removing a deployment from OpenShift, run the following commands:

```
kubectl get scc | grep sas
kubectl get scc | grep pgo
kubectl get clusterrole | grep "system:openshift:scc:sas-"
```

If you removed the cluster-wide SCC settings, each command should return zero results. If you do get results, repeat the steps in [“Remove SCCs” on page 127](#).

Remove the SAS Viya Platform Deployment Operator Only

To remove SAS Viya Platform Deployment Operator from the cluster but leave the instances of the SAS Viya platform running in their namespaces, run the following commands:

```
kubectl delete serviceaccounts,roles,rolebindings,deployments,secrets \
  -n name-of-deployment-operator-namespace \
  --selector "app.kubernetes.io/name=sas-deployment-operator"

kubectl delete clusterrolebindings \
  $(kubectl get clusterrolebindings \
    --selector "app.kubernetes.io/name=sas-deployment-operator" \
    -o jsonpath='{.items[?(@.subjects[0].namespace == "name-of-deployment-operator-namespace")].metadata.name}')
```

After the commands are completed, see [“Remove the Operator ClusterRoles” on page 127](#) for information to remove some remaining resources.

Uninstall Deployments That Do Not Use the Deployment Operator

The instructions in this section should be used to uninstall a deployment that was created using either Kubernetes commands (referred to as a "manual deployment" in earlier versions of this document) or with the `sas-orchestration` command. The steps assume that the SAS Viya platform is running and functional. If the software is not running, start it. If the software is not functioning, contact SAS Technical Support.

Remove the SAS Prepull and SAS Workload Orchestrator ClusterRoles and ClusterRoleBindings

Before you begin to uninstall your SAS Viya platform software, whether from a single namespace or from an entire cluster, perform the following tasks to account for some special cluster-wide permissions.

- 1 If you enabled the Image Staging Node List Option in your deployment, follow the instructions in the "Disable the Node List Option" section of the "SAS Image Staging Configuration Options" README located at `$deploy/sas-bases/examples/sas-prepull/README.md` (for Markdown format) and `$deploy/sas-bases/docs/sas_image_staging_configuration_options.htm` (for HTML format).
- 2 If your deployment includes Workload Management and Cluster Role was enabled for it, follow the instructions in the "Disable the Cluster Role" section of the "Cluster Privileges for SAS Workload Orchestrator Service" README located at `$deploy/sas-bases/examples/sas-workload-orchestrator/configure/README.md` (for Markdown format) and `$deploy/sas-bases/docs/configuration_settings_for_sas_workload_orchestrator_service.htm` (for HTML format).

Remove Service Account Links and SCCs

Note: If your deployment is not on Red Hat OpenShift, skip this section.

For Red Hat OpenShift systems, the namespace's service accounts must be removed from the assigned SCCs.

- 1 To find the namespace's service accounts that are linked to SCCs, run the following to generate the commands to review and run:

```

SAS_NS=name-of-namespace
for oc_scc in $(oc get clusterrolebinding | grep
"system:openshift:scc:" | awk '{ print $1 }'); do
    crb_name=$(oc get clusterrolebinding $oc_scc -o json | jq -r
'.roleRef.name')
    sa_names=$(oc get clusterrolebinding $oc_scc -o json | jq -r
'.subjects[] | select(.namespace == \"'$SAS_NS'\") | .name')
    for sa_name in $sa_names; do
        if [[ ! -z $sa_name ]]; then
            echo "oc -n $SAS_NS adm policy remove-scc-from-user $
{crb_name##*:} -z $sa_name"
        fi
    done
done

```

Run each command that is produced as output from the code.

- 2 If you have removed the last namespace running the SAS Viya platform in the cluster, then the SAS-specific SCCs can be removed. The following can be run to generate the commands to review and run:

Note: Removing SCCs is a cluster-wide operation. Ensure that no namespaces with SAS Viya platform software exist in the cluster before removing SCCs. Otherwise, any remaining SAS Viya platform deployments in the cluster may cease to work as expected.

```

for sas_scc in $(oc get scc | grep sas | awk '{ print$1 }'); do echo
"oc delete scc $sas_scc"; done
for sas_scc in $(oc get scc | grep pgo | awk '{ print$1 }'); do echo
"oc delete scc $sas_scc"; done

```

Run each command that is produced as output from the code.

- 3 Run the following commands:

```

kubectl get scc | grep sas
kubectl get scc | grep pgo
kubectl get clusterrole | grep "system:openshift:scc:sas-"

```

If you removed the cluster-wide SCC settings, each command should return zero results. If you do get results, repeat the steps above to remove the SCCs or ClusterRoles.

Remove Internal PostgreSQL Resources

Note: If your deployment does not include an internal instance of PostgreSQL, skip this topic.

Remove the resources for an internal instance of PostgreSQL from the namespace:

```

kubectl -n name-of-namespace delete postgresclusters --selector="sas.com/
deployment=sas-viya"

```

Note: If your Crunchy Data PVs do not have a reclaimPolicy of "retain", performing this command will delete their associated PVCs, resulting in the data they contain being lost. Therefore, ensure that you have backed up that data or are otherwise unconcerned about losing it.

If your Crunchy Data PVs do have a reclaimPolicy of "retain", their PVCs will not be deleted when the cluster is deleted. If you no longer need them, you can delete them manually by using the `kubectl get pvc` command to get the list of PVCs and the `kubectl delete` command to delete them.

Remove the SAS Viya Platform

Because your SAS Viya platform software is required to be deployed in a dedicated namespace, uninstall your deployment by deleting that namespace. As an administrator with cluster permissions, run the following command:

```
kubectl delete namespace name-of-namespace
```

When the namespace is deleted, your SAS Viya platform deployment has been uninstalled.

Note: If you have deployed the SAS Viya platform on multiple namespaces, repeat the command on each namespace to remove the SAS Viya platform from the cluster.

Remove CRDs

If you are removing the last namespace with a SAS Viya platform deployment from the cluster, the SAS custom resource definitions (CRDs) must be removed as well.

Note: Removing CRDs is a cluster-wide operation. Ensure that no namespaces with SAS Viya platform software exist in the cluster before removing CRDs. Otherwise, any remaining SAS Viya platform deployments in the cluster may cease to work as expected.

- 1 As a user with rights to list CRDs, run the following two commands to determine if any SAS CRDs still exist:

```
kubectl get crd --selector "sas.com/admin=cluster-wide"
```

```
kubectl get crd --selector "sas.com/admin=cluster-api"
```

- 2 If either of those commands return objects, run the following two commands as a user with rights to remove CRDs:

```
kubectl delete crd --selector "sas.com/admin=cluster-wide"
```

```
kubectl delete crd --selector "sas.com/admin=cluster-api"
```

- 3 Run the following command as a user with rights to remove clusterroles:

```
kubectl delete clusterrole --selector "sas.com/admin=cluster-wide"
```

4 Run the following commands:

```
kubectl get crd --selector "sas.com/admin=cluster-wide"
```

```
kubectl get crd --selector "sas.com/admin=cluster-api"
```

```
kubectl get clusterrole --selector "sas.com/admin=cluster-wide"
```

If you removed the cluster-wide CRDs, each command should return zero results. If you do get results, repeat the steps above to remove the CRDs.

Appendix 1

(Optional) Using a Mirror Registry

Create a Mirror Registry	133
About SAS Mirror Manager	134
System Requirements	134
Download Required Files	135
Create a Basic Mirror Registry	137
Required Flags	138
Optional Flags	138
Using a Configuration File	140
Deploy the Software from the Registry	140
Create and Populate a Mirror Registry in Microsoft Azure	141
Additional Requirements	141
Configure Authentication for Azure Container Registry	141
Create and Populate a Mirror Registry in Amazon ECR	142
Additional Requirements	142
Using SAS Mirror Manager on AWS	143
Create and Populate a Mirror Registry in Google Cloud Platform	144
Additional Requirements	144
Create the Registry	144
Grant Permissions to the SAS Mirror Manager User	144
Using SAS Mirror Manager on GCP	145
Create and Populate a Mirror Registry in Red Hat OpenShift	145
Create and Populate a Mirror Registry in JFrog Artifactory	147
Create a Mirror Registry at a Dark Site	147

Create a Mirror Registry

SAS Mirror Manager is a command-line utility that pulls SAS container images from the SAS Container Registry and pushes the images to your container registry, which can be either on-premises or in the cloud.

A mirror registry is optional, but it supports specific use cases, such as sites with limited internet access or those with a requirement to perform pre-deployment security scanning of SAS packages. Setting up the mirror requires access to the internet. However, after mirroring has been completed, the deployment can proceed at sites that lack internet access.

About SAS Mirror Manager

In addition to enabling deployment from a pre-populated mirror registry, SAS Mirror Manager supports other use cases. For example, you can use it to get a detailed list of the available SAS Viya platform stable and cadence releases in the SAS Container Registry. SAS Mirror Manager can help you determine whether differences exist between a SAS Viya platform repository and the current contents of the SAS Container Registry.

Using the SAS Mirror Manager list functionality, you can run checks before creating a mirror registry for your organization. Here are some examples of those checks:

To display a list of all the available software cadences and versions for a SAS software order:

```
mirrormgr list remote cadences --deployment-data SASViyaV4_order-
number_certs.zip
```

To display a list of all the available software versions and releases:

```
mirrormgr list remote cadence releases --deployment-data
SASViyaV4_order-number_certs.zip
```

To determine the required disk space for the available images:

```
mirrormgr list remote repos size --latest --deployment-data
SASViyaV4_order-number_certs.zip
```

System Requirements

The number of containers and the size of the images that you obtain from SAS depend on the products in your software order. SAS Mirror Manager downloads each image locally and then uploads it to the destination registry. The images that are downloaded require approximately 30–50 GB of disk space on the machine on which SAS Mirror Manager is running. You can use the `--remove-after-upload` option with the `mirrormgr mirror registry` command to remove container images from the local cache after they have been uploaded to the destination registry. However, make sure that you have enough space to store all of the concurrent downloads.

You can run SAS Mirror Manager on 64-bit Linux, Windows, or macOS. By default, SAS Mirror Manager performs the number of concurrent downloads that corresponds to the number of available CPUs. You can override this setting with the `--workers` option and a value that defines the number of concurrent threads to be used for downloads. For more information, see [“Required Flags” on page 138](#).

If you plan to use the SAS Viya Platform Deployment Operator with SAS Mirror Manager, perform the steps to download SAS Mirror Manager and create the mirror before deploying the operator. When you deploy the software for the operator,

download the certificates again. For more information, see [“Deploy the SAS Viya Platform Deployment Operator” on page 10](#).

SAS Mirror Manager does not support mirroring of Harbor versions 1.9.x and 1.10.x because of a Harbor limitation. Harbor versions earlier than 1.9.0 are supported, and Harbor version 2.0 and later are supported.

Before you run SAS Mirror Manager, perform the following steps:

- Configure your firewall to enable outbound connections to `ses.sas.download`, `ses.sas.com`, `crbwp.sas.com`, `bwp2.ses.sas.download`, and `cr.sas.com` on port 443.

SAS Mirror Manager uses mutual TLS for authentication. If you use a proxy server, configure it to allow access to these SAS sites to bypass the proxy.

- Verify network bandwidth.

SAS Mirror Manager requires a high-speed network environment. A lack of sufficient bandwidth increases the amount of time that is required to create the mirror registry and can lead to time-out errors when communicating with the destination registry.

- Determine the authentication mechanism for the target registry.

If you are using a cloud registry provider, use the corresponding CLI to access the user name and password for your cloud registry. The user account that is specified in the authentication must have permission to push images to the cloud registry.

See the following documentation for more information about target registries that require additional configuration:

- Azure Container Registry: [ACR documentation](#)
- Amazon Elastic Container Registry (Amazon ECR): [Amazon ECR documentation](#)
- Google Cloud Platform Container Registry: [GCR documentation](#)
- Red Hat OpenShift Container Registry: [OpenShift Container Platform registry documentation](#)
- JFrog Artifactory Registry: [Artifactory Package Management documentation](#)

In a multi-tenant deployment, all tenant onboarding and offboarding YAML files must include the transformer file `site-config/mirror.yaml`. For more information, see the README file at `$deploy/sas-bases/examples/sas-tenant-job/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/onboard_or_offboard_tenants_using_the_tenant_job.htm` (for HTML format).

Download Required Files

To create a mirror registry with SAS Mirror Manager, start by obtaining the software.

When you order SAS software, SAS sends a Software Order Email (SOE) to your business or organization. Your SOE includes information about the software order, including a link to `my.sas.com`.

- 1 Create a directory on your kubectl machine or on a machine that can be reached by your kubectl machine:

```
mkdir directory-name
```

SAS recommends that you name the directory `deploy`, but you should use a name that is meaningful to you. The directory is referred to as `$deploy` in this guide.

- 2 Click the **Get Started** link that is provided in your Software Order Email (SOE). The `my.sas.com` page opens in a web browser.
- 3 Click **My Orders** near the top of the page. The full list of your SAS software orders displays.
- 4 Click the **Deployment Tools** button, which is located above the list of orders. A popup window displays. It provides access to the SAS Viya Order API, SAS Container Manager, and SAS Mirror Manager.
- 5 Below the SAS Mirror Manager section, click **Learn More** in order to access the SAS Mirror Manager page. The page opens in a separate browser tab.
- 6 On the SAS Mirror Manager for SAS Viya page, click the download link in order to download the appropriate SAS Mirror Manager package for the machine where you want to create your mirror registry.

Note: If you have a version of SAS Mirror Manager that you used with an earlier version of SAS Viya (such as SAS Viya 3.5), it is not compatible with the current architecture of the SAS Container Registry.

- 7 Return to the MySAS - Order List page in its separate browser tab.
- 8 Click to select the order that you want to deploy.
- 9 In the Viya Orders > Order *Order ID* pane that opens, examine the order information. The description indicates the release cadence, the version of SAS Viya platform that you are about to deploy, and whether it has been downloaded previously.
- 10 Click the **Downloads** tab near the top of the page. The line above the **Asset Type** table indicates the release cadence and the version of SAS Viya platform software that you will deploy.
- 11 (Optional) If you want to deploy a different version, select the cadence and version from the menus. Or scroll down to see previous versions for this order.
- 12 Review the deployment documentation by clicking the **How to Deploy** link.
- 13 If you are performing a manual deployment, skip this series of steps:
 - a Create a new directory parallel to the `$deploy` directory in which to store your license and `*-certs.zip` files.

```
mkdir directory-name
```

You should use a name that is meaningful to you. The directory is referred to as `$license` in this guide. Replace `$license` with the directory name that you prefer.

- b** Select **Deployment Assets** from the **Asset Type** list, and then click the **Download** button to download the files that are required to deploy your software.

After unzipping the downloaded file, place the *-certs.zip file and the license file, identified by the .jwt extension, in the \$license directory.

Note: If you plan to have multiple SAS Viya platform deployments in your cluster, you should organize the \$license directory as you see fit. Whatever strategy you use to organize, ensure that you can easily differentiate the license and certs.zip files by order.

Save the TGZ file (the deployment assets) to the directory that you created in step 1.

- 14** Extract the files from the TGZ file in the same directory:

```
tar xvfz file-name.tgz
```

The result is a directory structure that looks like this:

```
$deploy/
└─ sas-bases/
    ├── base/
    ├── components/
    ├── docs/
    ├── examples/
    └─ overlays/
```

For a description of each subdirectory, see “\$deploy/sas-bases Directory” on page 27.

- 15** Uncompress the SAS Mirror Manager file, *mirrormgr-operating-system.tgz*, in your \$deploy directory.
- 16** (Optional) SAS recommends saving the SOE in the same directory.

Create a Basic Mirror Registry

The following command creates and populates a mirror registry:

```
mirrormgr mirror registry \
  --destination myregistry.mydomain.com \
  --username myregistryuser \
  --password myregistrypassword \
  --deployment-data ~path-to-certs-zip-file
```

The `mirror registry` command performs the following tasks:

- Pulls images from the SAS Container Registry.
- Pushes the images to the specified private registry.

The optional `--deployment-assets` flag is recommended. Without it, SAS Mirror Manager picks up the latest release of all supported cadence versions, and additional disk space is required to accommodate the software that is downloaded. For more information, see “Optional Flags” on page 138.

If the target registry for the images requires an `imagePullSecret` for access to those images, you will need to create one. For more information, see [“Use ImagePullSecrets to Access the Mirror Registry” on page 39](#).

Additional steps are required in order to create a mirror in each of the private cloud environments that the SAS Viya platform supports. Follow one of the documentation links below to set up the target registry:

- [“Create and Populate a Mirror Registry in Microsoft Azure” on page 141](#)
- [“Create and Populate a Mirror Registry in Amazon ECR” on page 142](#)
- [“Create and Populate a Mirror Registry in Google Cloud Platform” on page 144](#)
- [“Create and Populate a Mirror Registry in Red Hat OpenShift” on page 145](#)
- [“Create and Populate a Mirror Registry in JFrog Artifactory” on page 147](#)

Required Flags

Here is the basic command syntax, which includes required flags:

```
mirrormgr mirror registry required-flags
```

--deployment-data *path-to-certs-zip-file*

Specifies the full path to the `certs.zip` file that was downloaded to the `$deploy` directory. You downloaded the ZIP file from `my.sas.com`.

--destination *registry-location*

Specifies the host name and repository path of the registry to which images are replicated. Here is an example:

```
registry.example.com/my-project
```

--password *destination-registry-password*

Specifies the password that grants you access to the destination registry.

--username *destination-registry-user-name*

Specifies the user name that grants you access to the destination registry.

Optional Flags

--cacert *path-to-CA-certificate*

Specifies the CA certificate for communicating with the upstream repository server. The default is `$HOME/.config/mirrormgr/ca.pem`.

--cadence *cadence-version*

Specifies the cadence (or frequency) at which you deploy new SAS software. It can be used with the `--release` flag. Here is an example of specifying a release with the *stable* cadence, version 2020.1.1:

```
--cadence stable-2020.1.1
```

Note: The `--cadence` flag is not required. However, it enables you to mirror the exact container image versions from which the accompanying deployment orchestration assets were generated. By including the cadence and release, you

ensure that the mirror provides all the images that are required for a SAS Viya platform deployment. If you have downloaded deployment assets from my.sas.com, use the `--deployment-assets` flag instead of the `--cadence` flag.

--cert *path-to-auth-certificate*

Specifies the certificate for authenticating with the upstream repository server. The default is `$HOME/.config/mirrormgr/cert.pem`.

--config *config-file-location*

Enables you to automate the process of creating and populating a mirror registry. The default value for the location is `$HOME/.config/mirrormgr/mirrormgr.yaml`. For more information and an example, see [“Using a Configuration File” on page 140](#).

-d, --debug

Enables debug-level logging.

--deployment-assets

Uses data from the Deployment Assets archive that you downloaded with your software order to determine the cadence name, version, and release values. Replaces the `--cadence` and `--release` flags.

--latest

Mirrors only the latest versions of SAS content.

--log-file *path-to-log-file*

Specifies the log file location. The default is `$HOME/.local/share/mirrormgr/mirrormgr.log`.

-p, --path *path-to-repositories*

Specifies the path where repositories are downloaded. The default is `$HOME/sas_repos`.

Note: SAS recommends that you include this flag. Otherwise, the disk backing the default home directory might fill up.

--platform *platform-filter*

Specifies a filter for identifying the repositories to mirror, by platform. Here is an example of a platform filter:

```
--platform x64-oci-linux-2
```

--push-only

Instructs SAS Mirror Manager to mirror the local copy to the specified registry without connecting to the SAS Entitlement Service or to the SAS Container Registry. This flag is only used along with the `--destination` flag. For example of the syntax, see [“Create a Mirror Registry at a Dark Site” on page 147](#).

--release *release-cadence-timestamp-ID*

Specifies the software release that matches the deployment orchestration tools that are used in your deployment assets file (in TGZ format). The release includes the timestamp when the specified cadence version was last updated.

If the `--release` flag is used, the `--cadence` flag must also be used. If the `--release` flag is not used, SAS Mirror Manager mirrors the latest release of a cadence version.

Note: You can derive the cadence, version, and release from the file name of the deployment assets in the TGZ file that you download from my.sas.com. Here is an example file name:

```
SASViyaV4_09S468_stable-2020.1.1_20201016.1594946240211_deploymentAssets_2020-07-17T022620.tgz
```

In this example, `stable-2020.1.1` corresponds to the cadence and version, and the release includes a release year, month, and day plus an ID, such as `20201016.1594946240211`. However, if you have downloaded deployment assets, you can use the `--deployment-assets` flag to pass in the cadence and release automatically.

--remove-after-upload

Removes container images from the local cache after they have been uploaded to a remote registry.

-r, --repo

Specifies a list of repositories to mirror. If not specified, all entitled repositories are mirrored. To obtain a full list of repositories, run this command:

```
mirrormgr list remote repos
```

--url URL

Specifies the base URL of upstream repositories. The default is `https://ses.sas.download/ses`.

--workers number

Specifies the number of download threads to use. By default, the value is the same as the number of cores on the local machine.

A complete list of flags used in the command is available in the Help for SAS Mirror Manager. Access the Help using `-h` or `--help`.

Using a Configuration File

Passing in the location of a configuration file with the `--config` flag lets you largely automate the process of mirror creation and use a repeatable process. By default, the location of the configuration file is assumed to be `$HOME/.config/mirrormgr/mirrormgr.yaml`.

You can add virtually any of the available command flags to the configuration file by removing the leading hyphens (`--`). The file must be in YAML format. Enclose numeric string values in quotation marks (`" "`).

Here is an example of a valid configuration file:

```
deployment-data: ./certs.zip
cadence: "stable-2023.02"
release: "20230217.1667851938227"
path: ./sas_repos
log-file: ./mm.log
debug: true
```

Deploy the Software from the Registry

When you have verified that the SAS images to which you are entitled have been added to the specified registry, you are ready to deploy the SAS Viya platform.

- 1 Some updates to your `kustomization.yaml` file are required in order to configure the deployment to use the mirror registry. For more information, see [“Using a Mirror Registry” on page 36](#).
- 2 Your next steps depend on your deployment method:
 - If you are performing a deployment using the SAS Viya Platform Deployment Operator, follow the steps in [“Deploy the SAS Viya Platform Deployment Operator” on page 10](#) in order to prepare your environment.
 - If you are performing a manual deployment, follow the steps in [“Retrieve Required Files” on page 18](#) to download additional files that are required for the deployment. Then you can proceed to the [Installation on page 31](#) steps.

Create and Populate a Mirror Registry in Microsoft Azure

SAS Mirror Manager requires some additional configuration when the target container registry is hosted on Microsoft Azure.

Additional Requirements

Based on the average size of a SAS container deployment, a single deployment can be performed from a Standard-sized Azure container registry. However, if you intend to have a long-running registry that consumes SAS software updates over time, you should select a Premium size. For more information, see the [Microsoft Azure sizing documentation](#).

Configure Authentication for Azure Container Registry

Several methods of authenticating to the Azure Container Registry are available. You can see a full list in the Azure container authentication docs. You can use the Azure CLI to do the following:

- 1 Log in to your Microsoft Azure account in order to authenticate your client:

```
az login
```

- 2 Get an authorization token:

```
az acr login --name registry-name --expose-token
```

For *registry-name*, substitute the name of the destination container registry.

Note: The URL of any Azure container registry must be specified in all lowercase letters.

The output that is returned includes the access token and the login server, as in the following example:

```
docker login loginServer -u 00000000-0000-0000-0000-000000000000 -p accessToken
{
  "accessToken": "token-contents",
  "loginServer": "mirrormgr.azurecr.io"
}
```

- 3 (Optional) The token might contain a large amount of data. It is a best practice to save it as an environment variable:

```
TOKEN=$(az acr login --name registry-name --expose-token --output tsv
--query accessToken)
```

- 4 Run SAS Mirror Manager with your credentials and append the token environment variable:

```
mirrormgr mirror registry --deployment-data SASViyaV4_order-
number_certs.zip
--latest --destination registry-name.mydomain.com --username
00000000-0000-0000-0000-000000000000 --password $TOKEN
```

For *registry-name.mydomain.com*, specify the `loginServer` that was returned from the previous command as the `--destination`. The URL must be specified in all lowercase letters.

For *access-token*, specify the Azure token that was returned previously.

Note: The user name must be 00000000-0000-0000-0000-000000000000.

Return to [“Required Flags” on page 138](#) for more information about the SAS Mirror Manager commands. Consider whether you want to use any of the [“Optional Flags” on page 138](#). Then proceed to [“Deploy the Software from the Registry” on page 140](#) for additional deployment instructions.

Create and Populate a Mirror Registry in Amazon ECR

SAS Mirror Manager requires additional configuration when the target container registry is hosted on AWS.

Additional Requirements

In order to support Amazon ECR as the mirror destination, make sure you have fulfilled some additional requirements:

- The AWS CLI must be installed.
- The AWS CLI must be authenticated.
- The user account must have permission to create new Amazon ECR repositories. This permission is specified in the Identity and Access Management (IAM) section of the AWS console.

Using SAS Mirror Manager on AWS

The image repositories must be created in Amazon ECR before they are pushed. Take the following steps to use SAS Mirror Manager in an AWS environment:

- 1 Create or configure an AWS private registry with Amazon ECR.
Your AWS account has an associated default private Amazon ECR registry. For more information, see [Amazon ECR private registries](#).
- 2 Follow the steps in “Download Required Files” on page 135 to obtain the SAS Mirror Manager software.

- 3 Create any required repositories:

```
for repo in $(mirrormgr list target docker repos --deployment-data
SASViyaV4_order-number_certs.zip
--destination namespace) ; do
    aws ecr describe-repositories --repository-names $repo --region
region || aws ecr create-repository --repository-name $repo --region
region
done
```

Note: The call to `describe-repositories` runs in order to determine whether the repositories exist. If they do not exist, an error message is returned to alert you about this condition. This error may be safely ignored.

For *namespace*, supply the name of the target namespace. The name of the repository and namespace should conform to the repository name requirements that are described in the [Amazon ECR User Guide](#).

For *region*, supply the name of the AWS region. Be sure to use the same region in all commands.

- 4 Create and populate the mirror based on the associated software order:

```
mirrormgr mirror registry -p ./sas_repos --deployment-data
SASViyaV4_order-number_certs.zip
--latest --destination registry-URL/namespace --username 'AWS'
--password $(aws ecr get-login-password --region region)
```

The same options that are described in “Optional Flags” on page 138 are also applicable to Amazon ECR.

If you encounter errors that report too many requests, you can reduce the number of worker threads. By default, the number of threads that is used is the same as the number of CPU cores on the machine. Use the `--workers` option to set the number of threads.

- 5 Return to “[Required Flags](#)” on page 138 for information about the SAS Mirror Manager commands. Consider whether you want to use any of the “[Optional Flags](#)” on page 138. Then proceed to “[Deploy the Software from the Registry](#)” on page 140 for additional deployment instructions.

Create and Populate a Mirror Registry in Google Cloud Platform

SAS Mirror Manager requires additional configuration when the target container registry is hosted in GCP.

Additional Requirements

In order to support Google Cloud Platform as the mirror destination, make sure you have fulfilled the following additional requirements:

- To create a mirror registry in GCP, the GCP SDK is required. For instructions to download and install the GCP SDK, see <https://cloud.google.com/sdk/docs/quickstart-linux>.
- In order to create a mirror registry in GCP using SAS Mirror Manager on a Windows system, you must have Microsoft PowerShell.

Create the Registry

Create a container registry in your GCP account before you run SAS Mirror Manager. Follow these instructions: https://cloud.google.com/container-registry/docs/quickstart?hl=en_US.

Grant Permissions to the SAS Mirror Manager User

You must explicitly grant two storage-related permissions that enable SAS Mirror Manager to access your GCP environment.

To change required storage permissions:

- 1 Log in to your Google Cloud Platform account, and select **Cloud Storage** from the side panel.
- 2 Expand **artifacts.\$PROJECT.appspot.com** and select **permissions**.

- 3 In the **Add Members** field, enter the user name that you want to use to run SAS Mirror Manager commands. This account can be your user account or a service account.
- 4 Grant the following permissions to the account that you specified:
 - Storage Object Creator
 - Storage Object Viewer

Using SAS Mirror Manager on GCP

The command to run SAS Mirror Manager is slightly different in GCP environments. Use the `gcloud` command to fetch your access token. Here is an example:

```
mirror registry --deployment-data SASViyaV4_order-number_certs.zip --
latest --destination gcr.io/$GCP_PROJECT/$NAMESPACE --username
oauth2accesstoken --password $(gcloud auth print-access-token)
```

Note: The value for `username` is always `oauth2accesstoken`.

The value for `$NAMESPACE` must start with a letter and can contain only lowercase letters, numbers, hyphens, underscores, and forward slashes. Using a double hyphen, an underscore, or a forward slash is not supported.

Return to [“Required Flags” on page 138](#) for information about the SAS Mirror Manager commands. [“Optional Flags” on page 138](#). Consider whether you want to use any of the [“Optional Flags” on page 138](#). Then proceed to [“Deploy the Software from the Registry” on page 140](#) for additional deployment instructions.

Create and Populate a Mirror Registry in Red Hat OpenShift

When the target container registry is hosted in Red Hat OpenShift, SAS Mirror Manager creates a mirror of the SAS Viya platform software from your software order in the OpenShift Container Platform internal image registry.

Note: Before the SAS mirror registry can be created, a cluster administrator must configure the image registry for VMware vSphere and verify that a storage location is available. For more information, see [Configuring the registry for vSphere](#).

Take the following steps to create a mirror registry for an OpenShift deployment:

- 1 Create a namespace on the OpenShift cluster. You must specify the namespace where you will later deploy the SAS Viya platform:

```
oc create namespace name-of-namespace
```

- 2 A password is required to upload content to the image registry. Use one of the following methods to retrieve the password:

- Using the `oc whoami` command

- 1 Issue the following command:

```
oc whoami -t
```

- 2 Store the access token as an environment variable:

```
export password=token
```

For *token*, substitute the token value that was returned in the previous step.

- Extracting the password from an existing secret

- 1 When the namespace is created, several secrets should get created automatically:

```
oc -n name-of-namespace get secrets | grep builder-dockercfg
```

The command output should resemble the following:

```
builder-dockercfg-#####    kubernetes.io/dockercfg
```

- 2 Run the following command to display the connection information in a specified secret:

```
kubectl get secret -n namespace builder-dockercfg-##### --  
output=jsonpath={.data.\.dockercfg}" | base64 --decode
```

For *builder-dockercfg-#####*, specify the name of the secret from the previous step.

- 3 Store the password token as an environment variable:

```
export password=password
```

For *password*, use the long string value that was found for *password* in the output from the previous step.

- 3 Create the registry using SAS Mirror Manager. Run the command and append the token variable. The command should resemble the following example:

```
./mirrormgr mirror registry --destination path-to-openshift-image-  
registry/project-name --flatten --deployment-data SASViyaV4_order-  
ID_certs.zip --latest --username user-name --password $password --  
path /path-to-local-image-cache
```

If you generated an access token, specify your login ID as the *user-name*.

If you extracted the password from the secret, specify the user name that was provided in the output. For example, specify *serviceaccount* for the *user-name*.

Note: The `--flatten` flag is required for a deployment into OpenShift.

Additional flags are optional. For example, you could add the `--deployment-assets` or `--cadence` flags in order to mirror only the content that you require, rather than mirroring the latest versions of all supported cadences. For more information, see [“Optional Flags” on page 138](#).

A transformer that is applicable only to "flattened" registries has been provided in your deployment assets. You must add it to your base `kustomization.yaml` file just

after `overlays/required/transformers.yaml`. Proceed to [“Deploy the Software from the Registry” on page 140](#) for instructions.

If you are using the SAS Viya Platform Deployment Operator to manage your SAS Viya platform deployment, verify that the data from the local mirror copy is accessible by the operator.

Create and Populate a Mirror Registry in JFrog Artifactory

Artifactory is a supported option for a private or internal Docker registry. SAS has tested using the `local docker registry` option from a self-hosted Artifactory. Authentication to the registry can use either a user ID and password combination or authentication tokens.

- 1 Generate an API token through the JFrog Platform Access Token UI.

- 2 Export the token as an environment variable:

```
export accessToken=token-string
```

- 3 Run SAS Mirror Manager to create the mirror registry:

```
mirrormgr mirror registry --remove-after-upload --platform platform-  
filter --destination registry-name.mydomain.com --username user-name --  
password $accessToken --deployment-data path-to-certs-zip-file --  
deployment-assets SASViyaV4_order-  
number_cadence_release_deploymentAssets_order-ID.tgz --path /path-to-  
repositories
```

- 4 Return to [“Required Flags” on page 138](#) for information about the SAS Mirror Manager commands. Consider whether you want to use any of the [“Optional Flags” on page 138](#). Then proceed to [“Deploy the Software from the Registry” on page 140](#) for additional deployment instructions.

Create a Mirror Registry at a Dark Site

You can use SAS Mirror Manager to enable deployment of the SAS Viya platform to an air-gapped network or to a network where systems are not allowed to communicate with hosts on the internet. SAS Mirror Manager enables you to download content to a local directory, from which you can copy the content to the isolated network. From there, you can upload SAS Viya platform content to an internal registry.

To deploy the SAS Viya platform to a dark site:

- 1 Follow the steps in [“Download Required Files” on page 135](#) to obtain the SAS Viya platform software and entitlements and to download SAS Mirror Manager.

2 Mirror the software to a local directory.

```
mirrormgr mirror registry --path ./sas_repos --deployment-data
SASViyaV4_order-number_certs.zip
```

If you do not provide a destination to the `mirrormgr mirror registry` command, all files are downloaded to the directory that is specified by the `--path` flag. If you do not specify a path, `~/sas_repos` is used by default.

TIP Be sure to include any cadence version filters that you want to use in order to limit the software that is downloaded. For more information, see [“Optional Flags” on page 138](#).

3 Mirror the local directory to an internal registry.

```
mirrormgr mirror registry --deployment-data SASViyaV4_order-
number_certs.zip --path ./sas_repos --destination registry.example.com
--username registry-user-name --password registry-password --push-only
```

Specifying the `--destination` flag along with the `--push-only` flag mirrors the local copy to the specified registry without connecting to the SAS Entitlement Service or to the SAS Container Registry.

Return to [“Required Flags” on page 138](#) for more information about the SAS Mirror Manager commands. Consider whether you want to use any of the [“Optional Flags” on page 138](#). Then proceed to [“Deploy the Software from the Registry” on page 140](#) for additional deployment instructions.

Appendix 2

SAS Viya Deployment Operator Fields and Messages

Overview	149
Fields in the SASDeployment Custom Resource	150
Primary Fields	150
User File Options	151
Secret Key Selector Fields	151
Repository Warehouse Options	152
Deployment User Content Options	152
Manage Updates	152
Environment Variables for the Operator Pod	153
Communications from the Operator	156
Overview	156
Status Fields	157
Deployment Condition Fields	157
Deployment Release Fields	158
Deployment Release Details Fields	158

Overview

This appendix contains lists of fields and variables that can be included in the files to configure the SAS Viya Platform Platform Deployment Operator or that are included in its status messages.

Fields in the SASDeployment Custom Resource

Primary Fields

The following fields can be added to the SASDeployment custom resource in the spec (specification) block.

Note: Apply the indentation rules from the preceding examples to new resources that you create

Table A2.1 Fields in the spec Block

Field	Description
caCertificate	Optional. Information about the certificate authority. See “User File Options” on page 151 for more information about the format.
clientCertificate	Optional. Information about the client certificate. See “User File Options” on page 151 for more information about the format.
cadenceName	Required. The name of the cadence to deploy.
cadenceRelease	Optional. The release of the cadence to deploy.
cadenceVersion	Required. The version of the cadence to deploy.
deploymentType	Optional. The type of deployment to use.
imageRegistry	Optional. The image registry to deploy from.
license	Optional. The license to apply to the deployment. For more information about the format, see “User File Options” on page 151
repositoryWarehouse	Optional. Information about the repository warehouse. See “Repository Warehouse Options” on page 152 for more information about the options.

Field	Description
userContent	Required. The user customizations. See “Deployment User Content Options” on page 152 for information about the options.

User File Options

A field that requires a user file can have one of three values:

Table A2.2 User File Options

Field	Description
content	A nested map.
secretKeyRef	A secret in the namespace of this SASDeployment. See “Secret Key Selector Fields” on page 151 for more information.
url	A <i>go-getter</i> URL from which to fetch content. See GitHub for more information about go-getters.

Secret Key Selector Fields

A field that requires a secret key selector has two parts:

Table A2.3 Secret Key Selector Fields

Field	Description
name	Required. The name of the secret.
key	Required. The name of the secret key.

Here is an example of the format:

```
spec:
  ...
  caCertificate:
    secretKeyRef:
      name: secret-name
      key: key-name
```

...

Repository Warehouse Options

The repository warehouse field has two options. For an example of the format, see [“Manage Updates” on page 152](#).

Table A2.4 Repository Warehouse Options

Field	Description
updatePolicy	Optional. The policy by which the operator automatically applies content changes to the repository warehouse. See “Manage Updates” on page 152 for details of the possible values.
url	Optional. The repository warehouse URL from which to deploy your SAS Viya platform. The default is <code>https://ses.sas.download/ses</code> .

Deployment User Content Options

Deployment user content must use one of these two options:

Table A2.5 Deployment User Content Options

Field	Description
files	A nested map.
url	A go-getter URL from which to fetch content. See GitHub for more information about go-getters.

Manage Updates

The SAS Viya Platform Deployment Operator can automatically download and apply updates to your SAS Viya platform with the addition of the `updatePolicy` field in the custom resource.

...

```
spec:
  cadenceName: "LTS"
  cadenceVersion: "2020.1"
  cadenceRelease: "20200817.1597696671121"
  repositoryWarehouse:
  updatePolicy: {{ VALUE }}
```

Replace the `{{ VALUE }}` variable with one of the following values:

Never

The operator performs no updates. This is the default. Alternatively, if you do not want the operator to perform updates, you can omit the `updatePolicy` field entirely.

Releases

The operator automatically applies any `cadenceRelease` updates within the selected `cadenceName` and `cadenceVersion`. This setting restricts the updates to software with two decimal places in their version numbers. For example, if you choose "Releases" with the version/name/release in the example above, then the operator would update the deployment for each stable release in 2020.1 (2021.1.1, 2021.1.2, and so on) and any patch releases of these versions. It would not update to 2020.2.

Updates are checked once every 24 hours by a cronjob that is created in the SAS Viya platform namespace at deployment time. The exact start time is randomized with each deployment in an attempt to keep all of the deployments in a given time zone from updating at the same time. The schedule can be modified for all SAS Viya platform deployments managed by an operator by setting the `AUTOUPDATE_SCHEDULE` environment variable on the SAS Deployment Operator's Deployment resource. The schedule can also be modified for a specific deployment by setting the `environment.orchestration.sas.com/AUTOUPDATE_SCHEDULE` annotation on the `SASDeployment` custom resource for that deployment.

To prevent updates from occurring when building and applying changes to the `SASDeployment` custom resource, set the custom resource's cadence variables to their current values. When you are ready to perform an update or re-enable automatic updates, modify and apply the custom resource again to remove these settings. An example of when this strategy is useful is making permanent changes to the base `kustomization.yaml` file after having made other, ad hoc, changes.

Environment Variables for the Operator Pod

The following variable is required in the pod specification of the file that deploys the operator.

SERVICE_ACCOUNT_NAME

Specifies the name of the service account under which the operator is running. This value is typically provided by the Downward API in Kubernetes.

The following variables are optional in the pod specification of the file that deploys the operator.

AUTOUPDATE_SCHEDULE

Specifies the schedule for the automatic update job, in CRON format. The default is daily. SAS recommends not setting this schedule to more often than once a day.

DEBUG_LOG_LOCATION

Enables debug logging to a file at the specified path. The directory that houses the file must exist and be writable. The special values `stdout` and `stderr` enable debug logging to standard output and standard error, respectively.

DISABLE_APPLY

Disables the application of the results of generation when the value is `true`. When combined with `DISABLE_CLEANUP`, this variable enables the user to test the operator results without making any changes to their cluster.

DISABLE_AUTO_RBAC

Disables the creation of the service accounts and associated RBAC resources that the operator normally creates for the `autoupdate` and `reconcile` workloads. When the variable is set to `true`, the Kubernetes administrator is responsible for ensuring that the service accounts for these two workloads have been created and have the appropriate permissions. This variable is usually used with `SERVICE_ACCOUNT_AUTOUPDATE_NAME` and `SERVICE_ACCOUNT_RECONCILE_NAME` to define the names of the service accounts the administrator created.

DISABLE_CLEANUP

Disables cleanup of the work directory when the variable is set to `true`. The operator writes temporary files to the work directory. Users can mount persistent storage to this location and set this variable to `true` to test operator results.

DISABLE_LEADER_ELECTION

Disables leader election when the variable is set to `true`. The user must ensure that no more than one instance of the operator is managing the same namespace.

DISABLE_LIVENESS_PROBE

Disables the liveness endpoint when the variable is set to `true`. Otherwise, the operator starts a liveness endpoint and creates a Service resource in Kubernetes for that endpoint.

DISABLE_METRICS_SERVICE

Disables the metrics service when the variable is set to `true`. Otherwise, the operator starts a metrics endpoint and creates a Service resource in Kubernetes for that endpoint.

FLATTENED_IMAGE_REGISTRY

Excludes the deployable unit's unit path in the calculation of the image path when the variable is set to `true`.

HTTP_PROXY

When a forward proxy server is in use, the `HTTP_PROXY` environment variable defines the proxy URL for HTTP requests.

HTTPS_PROXY

When a forward proxy server is in use, the `HTTPS_PROXY` environment variable defines the proxy URL for HTTPS requests.

KUBECONFIG

When the operator is running outside of a Kubernetes cluster, this variable can be set to identify and provide access to a Kubernetes cluster. By default, the `~/.kube/config` file is used.

MINIMUM_UPDATE_INTERVAL

This variable is deprecated with version Stable 2020.1.1. While this variable is still supported, you should use `AUTOUPDATE_SCHEDULE` instead. See [“AUTOUPDATE_SCHEDULE” on page 154](#) for more information.

NO_PROXY

When a forward proxy server is in use, the `NO_PROXY` environment variable defines requests that the proxy should ignore.

SAS_UPDATE_CHECK_INTERVAL

This variable is deprecated with version Stable 2020.1.1. While this variable is still supported, you should use `AUTOUPDATE_SCHEDULE` instead. See [“AUTOUPDATE_SCHEDULE” on page 154](#) for more information.

SERVICE_ACCOUNT_AUTOUPDATE_NAME

Defines the name of the service account that is used by the `autoupdate` job launched by the operator.

SERVICE_ACCOUNT_NAMESPACE

When the operator is running outside of a Kubernetes cluster, this variable can be set to specify the namespace of the service account under which the operator is running.

SERVICE_ACCOUNT_RECONCILE_NAME

Defines the name of the service account that is used by the `reconcile` job launched by the operator.

WATCH_NAMESPACE

Specifies the namespace to manage. If this variable is not set, or is set to an empty string, then the operator functions across the cluster and watches all namespaces.

WORK_DIRECTORY

Specifies an override for the work directory that is used by the operator.

To revise these settings, you should modify the deployment operator manifest file. The manifest file is located at `$operator-deploy/operator-base/deployment.yaml`. The variables should be listed in the `env` block using the indentation already present in the file. Here is an example using the proxy variables:

```
containers:
  - args:
    - reconcile
    command: []
    env:
      - name: HTTP_PROXY
        value: "http://webproxy.example.com:5000"
      - name: HTTPS_PROXY
        value: "http://webproxy.example.com:5000"
      - name: NO_PROXY
        value: "localhost,noproxy.example.com,apiserver.example.com"
    ...
```

Communications from the Operator

Overview

When the operator performs tasks, descriptions of actions and results are added to the status block in the SASDeployment custom resource. The following tables describe the format and content of the status block.

Here is an example of a status block for reference:

```

apiVersion: orchestration.sas.com/v1alpha1
kind: SASDeployment
...
status:
  conditions:
  - lastTransitionTime: "2020-10-06T13:19:24Z"
    message: Reconcile succeeded
    reason: ReconcileSucceeded
    status: "True"
    type: Ready
  displayName: Stable
  latest:
    releaseForCadence:
      cadenceName: stable
      cadenceRelease: "20201006.1601989747737"
      cadenceVersion: "2020"
      displayName: Stable
      matchesDeployedRelease: true
      supportEnd: "2030-03-15T04:00:00.000Z"
      supportLevel: SUPPORTED
    releaseForVersion:
      cadenceName: stable
      cadenceRelease: "20201006.1601989747737"
      cadenceVersion: "2020"
      displayName: Stable
      matchesDeployedRelease: true
      supportEnd: "2030-03-15T04:00:00.000Z"
      supportLevel: SUPPORTED
  messages:
  - ""
  - ""
  - |-
    role.rbac.authorization.k8s.io/sas-deployment-operator-mickey-test-
    a-sas-deployment created
  ...

```

Status Fields

The following fields can be added to the status block of the SASDeployment custom resource (referred to as *resource* in the table)..

Table A2.6 *Deployment Status Fields*

Field	Description
conditions	The current state of this SASDeployment. See “Deployment Condition Fields” on page 157 for more information.
displayName	The display name of the cadence for this SASDeployment.
messages	The list of messages from the last operation that was performed by the operator on this SASDeployment.
state	The state of this SASDeployment. Valid values are PENDING, RECONCILING, SUCCEEDED, or FAILED.
supportEnd	The date on which support for this SASDeployment ends.
supportLevel	The support level for this SASDeployment.

Deployment Condition Fields

The following fields are used to describe the condition of the SASDeployment custom resource.

Table A2.7 *Deployment Condition Fields*

Field	Description
type	The type of this condition.
status	The status of this condition.
lastTransitionTime	The timestamp that corresponds to the last status change of this condition.
reason	The reason for this condition's last transition.

Field	Description
message	The description of the last transition.

Deployment Release Fields

The following fields are used to describe the cadence release and version.

Table A2.8 *Deployment Release Fields*

Field	Description
releaseForVersion	The latest available release for the version of this SASDeployment.
releaseForCadence	The latest available release for the cadence of this SASDeployment.

For more information about either of these fields, see [“Deployment Release Details Fields” on page 158](#).

Deployment Release Details Fields

The following fields are used to provide details about the cadence release and version.

Table A2.9 *Deployment Release Details Fields*

Field	Description
cadenceName	The name of this available release.
cadenceRelease	The release of this available release.
cadenceVersion	The version of this available release.
displayName	The displayName of this available release.
matchesDeployedRelease	The value is <code>True</code> if this available release is currently deployed by this SASDeployment.
supportEnd	The date on which support for this available release ends.

Field	Description
supportLevel	The support level of this available release.

Appendix 3

Change Deployment Methods

<i>Kubernetes Commands to SAS Viya Platform Deployment Operator</i>	161
Deploy the SAS Viya Platform Deployment Operator	161
Create the \$license Directory	162
Create the SASDeployment Custom Resource	162
Deploy the Software	162
<i>Kubernetes Commands to sas-orchestration Command</i>	163
<i>SAS Viya Platform Deployment Operator to Any Other Method</i>	163

Kubernetes Commands to SAS Viya Platform Deployment Operator

.....
Note: Changing your deployment method requires an outage for your SAS Viya platform software.
.....

Deploy the SAS Viya Platform Deployment Operator

If you have not already deployed the SAS Viya Platform Deployment Operator, follow the instructions in [“Deploy the SAS Viya Platform Deployment Operator”](#) on [page 10](#).

Create the \$license Directory

The SAS Viya Platform Deployment Operator needs an explicit location for the license file and the zip file that contains the certificates for your software order. To create this directory, see step 6 in [“Retrieve Required Files” on page 18](#).

Create the SASDeployment Custom Resource

To create the custom resources required by the SAS Viya Platform Deployment Operator, follow the instructions in [“Create the SASDeployment Custom Resource” on page 77](#).

When you reach the step to issue the `create sas-deployment-cr` command, you must enter the correct values for the `--cadence-*` fields. To find those values, run the following command:

```
kubectl -n name-of-namespace get cm -o yaml | grep ' SAS_CADENCE '
```

Use the output of this command to determine the correct values for the `create sas-deployment-cr` command.

Note: If you have manually modified any of the resources since the SAS Viya platform was last deployed, those changes are overwritten and replaced with their original values when this command is run.

Deploy the Software

To avoid updating your SAS Viya platform software while changing the deployment method, ensure that the `updatePolicy` field in the custom resource is set to “never” or that the cadence variables are set to their current values. For more information, see [“Manage Updates” on page 152](#).

Follow the instructions in [“Deployment Using the SAS Viya Platform Deployment Operator” on page 84](#) to deploy the software using the SAS Viya Platform Deployment Operator.

If you did not update your software while changing deployment methods, after you have confirmed that the operator is working as you expect, follow the instructions at [“Update to a New Version Using the Deployment Operator” in SAS Viya Platform Operations: Updating Software](#) to update your software.

Kubernetes Commands to sas-orchestration Command

Note: Changing your deployment method requires an outage for your SAS Viya platform software.

Follow the instructions at [“Deployment Using the sas-orchestration Command” on page 87](#) to deploy or update your software using the sas-orchestration command. Be sure to install the orchestration tool, described at that link, as part of the process of changing deployment methods.

SAS Viya Platform Deployment Operator to Any Other Method

To change deployment methods from the SAS Viya Platform Deployment Operator to anything else requires dissociating the SAS Viya deployment from the operator.

- 1 Dissociate the deployment from the operator by following the steps in the README located at `$deploy/sas-bases/examples/deployment-operator/disassociate/README.md` (for Markdown format) or at `$deploy/sas-bases/docs/disassociate_a_sas_viya_platform_deployment_from_the_sas_viya_platform_deployment_operator.htm` (for HTML format).
- 2 (Optional) If all the deployments deployed by the operator have been dissociated from the operator, you can remove the operator. For the steps to remove the operator, see [“Remove the SAS Viya Platform Deployment Operator Only” on page 128](#).
- 3 After the dissociation, administer and update your deployment according to the procedures described for your chosen deployment method.

Appendix 4

Troubleshooting

<i>Deployment Operator Custom Resource Is Too Large</i>	165
<i>Error Checking Database Connection Error Message</i>	166
<i>Certframe Container Fails at init Stage</i>	166
<i>Issues with External PostgreSQL Server Size and Connections</i>	167
<i>SAS Viya Platform Deployment Operator Reconcile Failures</i>	168
<i>Performing Updates Produces Many Error Messages</i>	168
<i>Failure to Create Symbolic Link</i>	169
<i>Long Delay for SAS/CONNECT Server Sign-on from SAS Studio</i>	169
<i>"Operation not permitted" Error in SAS Configurator for Open Source Logs</i> ...	170
<i>Uninstall Leaves the Deployment Namespace in a Terminating State</i>	170

Deployment Operator Custom Resource Is Too Large

Error

When you are performing a deployment with the SAS Viya Deployment Operator and you attempt to apply the SASDeployment custom resource, either of the following error messages is displayed:

```
Error from server (RequestEntityTooLarge)
```

```
The SASDeployment "sas-viya" is invalid: metadata.annotations: Too long: must have at most 262144 bytes
```

Explanation

Applying the custom resource requires that it not exceed a size limit. The RequestEntityTooLarge and Too long errors indicate that the custom resource exceeds that limit.

Resolution

When the custom resource is created, it consumes the content of the \$deploy directory. To minimize the size of the custom resource, ensure that only the files that are needed are included in the \$deploy directory. For example, site.yaml files from earlier manual deployments and old or duplicate kustomization.yaml files should be moved from the \$deploy directory or deleted.

Error Checking Database Connection Error Message

Error

The following error message is displayed:

```
Error checking database connection: x509: certificate signed by unknown authority
```

Explanation

The SAS Viya platform cannot connect securely with a database during deployment.

Resolution

For instructions to provide CA certificates to the SAS Viya platform, see the "Incorporating Additional CA Certificates into the SAS Viya Deployment" section of the README file at \$deploy/sas-bases/examples/security/README.md (for Markdown format) or at \$deploy/sas-bases/docs/configure_network_security_and_encryption_using_sas_security_certificate_framework.htm (for HTML format).

Certframe Container Fails at init Stage

Error

The certframe container on one or more pods fails with the following log entries:

```
[ERROR] - Certificate failed to show a status of 'Ready' with a new resource version after retrying for 300 seconds
```

or

```
[ERROR] - Could not find a CertificateRequest in 'Ready=False' state associated with Certificate
```

Explanation

A pod has not completed initialization, and the certframe init container is moving into a failed state.

Resolution

- 1 If a pod has not completed pod initialization due to the certframe container, search the logs of the certframe container for an error message such as the following:

```
[ERROR] - Could not find a CertificateRequest in 'Ready=False' state associated with
```

Certificate sas-risk-management-8d2d37cz42-p7rj9

- 2 If you find the error message, check the `certificaterequests` resource to determine whether there are multiple requests.

```
kubectl -n name-of-namespace --kubeconfig kubeconfig get certificaterquests
| grep search-term
```

For *search-term*, use part of the pod name. From the example error message, `risk` or `management` is a good search term. Any part of the pod name can be used, but the more uncommon the search term is, the fewer pod names there are to search through. Here is an example command that uses the sample error message in a namespace named `test1`:

```
kubectl -n test1 --kubeconfig kubeconfig get certificaterquests | grep risk
```

Here is sample output from the command:

<code>sas-risk-management-app-8d2d37cz42-p7rj9-5xpr7</code>	<code>True</code>	<code>123m</code>
<code>sas-risk-management-app-8d2d37cz42-p7rj9-j6g5c</code>	<code>True</code>	<code>126m</code>

- 3 If you find multiple `certificaterquests`, delete the oldest ones until only one is left.

```
kubectl -n name-of-namespace --kubeconfig kubeconfig delete
certificaterequest full-pod-name
```

Here is an example that uses output from the preceding steps.

```
kubectl -n test1 --kubeconfig kubeconfig delete certificaterquest sas-
risk-management-app-8d2d37cz42-p7rj9-j6g5c
```

- 4 Check the pod status.

Issues with External PostgreSQL Server Size and Connections

Note: If your deployment uses an internal instance of PostgreSQL, skip this topic.

Error

Deployments with an external instance of PostgreSQL have any of the following messages in the logs.

- Error messages associated with `sas-transfer` that include the following:

```
GenericJDBCException: Unable to acquire JDBC Connection
```

- Error messages associated with `sas-backupjob`:

```
The backup agent encountered an error for the "postgres" data source.
```

```
The "backup" job for the "postgres" data source finished with the status "Failed".
```

- The `sas-readiness` service shows the `sas-database-ready` check going in and out of the ready state with the following message:

Connection failed: pq: remaining connection slots are reserved for non-replication superuser connections

Explanation

These error messages indicate that the PostgreSQL database, which is required to support SAS Infrastructure Data Server, has exhausted its pool of available connections.

Resolution

Verify that you have met all the requirements for an external PostgreSQL database in [“External PostgreSQL Requirements” in System Requirements for the SAS Viya Platform](#).

An external PostgreSQL server should support max_connections and max_prepared_transactions of at least 1024. However, in some environments, particularly multi-tenant SAS Viya platform deployments, this value might not be sufficient. For more information, see [“PostgreSQL Requirements for a Multi-Tenant Deployment” in System Requirements for the SAS Viya Platform](#).

SAS Viya Platform Deployment Operator Reconcile Failures

Error

After you apply the SASDeployment custom resource, checking the status results in a FAILED state.

Explanation

The state of the content modified for deployment, such as the kustomization.yaml file or any of the example YAML files used for configuration, is not what the SAS Viya Platform Deployment Operator expects. The result is a conflict that causes the deployment to fail.

Resolution

See [“Initial Troubleshooting” on page 85](#) for information about determining the cause of failure. The `.status.messages` field contains the actual messages from Kustomize and may prove the most useful in diagnosing problems.

Performing Updates Produces Many Error Messages

Error

During the update process, you see many errors messages in the logs.

Explanation

A number of error messages come from timing issues that resolve themselves in the course of the update. Some might come from third-party software and are

more like warnings or information messages. But since they originate with third-party software, SAS cannot change the content of the messages.

Resolution

The error messages should begin to diminish after 30 minutes. If more errors continue to appear in the logs, then you need to investigate.

Failure to Create Symbolic Link

SAS Configurator for Open Source requires a PVC for Python builds. If you want to use Microsoft Azure File Services (azurefile) as a storage class for the PVC, `msyslinks` must be supplied as a mount option to the azurefile storage class in order to enable the update functionality. Here is an example:

```
kubectl describe storageclass azurefile -n production
Name: azurefile
IsDefaultClass: No
Annotations: none
Provisioner: kubernetes.io/azure-file
Parameters: skuName=Standard_LRS
AllowVolumeExpansion: True
MountOptions:
msyslinks
actimeo=30 ReclaimPolicy: Delete
VolumeBindingMode: Immediate
Events: none
```

If the option is not supplied, a non-fatal message is generated in the SAS Configurator for Open Source pod log that reports a failure to create a symbolic link on the file system. The message provides the target of the link to be referenced. The target can be used to reference the Python install.

Long Delay for SAS/CONNECT Server Sign-on from SAS Studio

You might see a long (2-minute) delay during sign-on to a SAS/CONNECT server from an internal SAS Viya platform client, such as SAS Studio. You can work around this issue by setting `TCPNOIPADDR` to 1.

This problem has been observed on Microsoft Azure, AWS, and GCP. The delay is caused by an errant connection back to the SAS/CONNECT spawner. The spawner hangs until time-out, at which point a successful connection is made to the launched SAS/CONNECT server.

A fix is planned. In the meantime, run your programs with `TCPNOIPADDR=1` to work around the issue that causes the delay.

"Operation not permitted" Error in SAS Configurator for Open Source Logs

Error

When a pod ends up in an error state, searching the logs reveals that the following fatal error is included for SAS Configurator for Open Source:

```
install: cannot change permissions of '/opt/sas/viya/home/sas-pyconfig/...':
Operation not permitted
```

Explanation

SAS Configurator for Open Source performs a number of operations in the PVC. To use Microsoft Azure File Services (azurefile) as a storage class for the PVC, add the `noperm` mount option to the azurefile storage class. The `noperm` mount option allows client-side permissions checks to pass if options such as `file_mode` or `dir_mode` are set on the PVC mount in the pod.

Resolution

Add the `noperm` mount option to the azurefile storage class:

```
Name: azurefile
IsDefaultClass: No
Annotations: none
Provisioner: kubernetes.io/azure-file
Parameters: skuName=Standard_LRS
AllowVolumeExpansion: True
MountOptions:
noperm
actimeo=30 ReclaimPolicy: Delete
VolumeBindingMode: Immediate
Events: none
```

Uninstall Leaves the Deployment Namespace in a Terminating State

Note: If you have uninstalled your software using the SAS Viya Platform Deployment Operator, you should skip this note. Also, if your deployment did not have an internal instance of PostgreSQL, you should skip this note.

Error

When you perform an uninstall, the namespace in which your software was deployed is left in an unending `Terminating` state.

Explanation

When you uninstalled your software, the PostgreSQL resources were not removed prior to uninstalling your deployment.

Resolution

- 1 Create a list of the PostgresClusters in your deployment:

```
kubectl -n name-of-namespace get postgresclusters
```

- 2 Run the following command for each PostgresCluster returned in step 1, replacing *PostgresCluster-name* with the name of the PostgresCluster.

```
kubectl -n name-of-namespace patch postgrescluster PostgresCluster-name -p  
'{"metadata":{"finalizers":[]}}' --type=merge
```

